

How to Program a Calculator

Numerical analysis of common functions

Jake Darby

Abstract

This document will discuss and analyse various numerical methods for computing functions commonly found on calculators. The aim of this paper is to compare and contrast several algorithms, for each function, in regards to their efficiency and accuracy.

Contents

1	Introduction	3
1.1	Code and Computers used	3
2	General Definitions and Theorems	5
2.1	Methods	5
2.1.1	Newton-Raphson Method	5
2.1.2	Taylor Series Expansion	6
2.2	Errors	7
2.3	Convergence	7
2.4	Efficiency Metrics	10
3	Root Functions	10
3.1	Digit by Digit Method	10
3.2	Bisection Method	15
3.3	Newton's Method for Square Roots	19
3.4	Newton's Inverse Square Root Method	25
3.5	Comparison of Methods	27
4	Trigonometric Functions	29
4.1	Trigonometric Identities	29
4.2	Calculating π	30
4.3	Geometric Method	32
4.4	Taylor Series	38
4.5	CORDIC	44
4.6	Comparison of Methods	52
5	Logarithms and Exponentials	54
5.1	Calculating x^a	55
5.2	Calculating x^b	56
5.3	Naive Method	57
5.4	Taylor Series Method	58
5.5	Hyperbolic Series Method	63
5.6	Continued fractions	64
5.7	Comparison of Methods	72
6	Conclusion	76
A	Code	78
A.1	General Code	78
A.2	Square Root Code	84
A.3	Trigonometric Code	103
A.4	Exponential and Logarithm Code	134

1 Introduction

For many thousands of years all calculations that a person might want performing had to be done by hand. For simple calculations such as addition, subtraction and multiplication this was not such an issue, but as society evolved we wanted to know the answer to increasingly hard questions. The Greeks' sought to find a value for π , and ended up with the bounds that $\frac{223}{71} < \pi < \frac{22}{7}$ [14][19, p. 106], which while sufficient for their needs is not sufficient for ours in the present.

At the same time many functions were being studied to find solutions, often arising from practical concerns. For instance finding the square root of any arbitrary number has been important to architects since the time of the ancient Babylonian mathematics[25]. Similarly relevant have been the periodic trigonometric functions due to their relation to triangles, and exponential functions due to their use in finance for example to find interest on loans.

The difficulty of these methods is that there is typically no simple way of getting an exact answer, if in fact one is available. Over time methods were developed that would allow a person to calculate an approximate answer to their problem, given enough time and patience. Such methods tended to be long and tedious work, which even lead to the profession of a human computer from the early 17th century until the 20th century; who would be hired for that purpose.

By the time of the Renaissance period people had started to build early mechanical calculators to help in these endeavours. Such calculators were typically capable of only addition and subtraction, which could be used to implement multiplication and division if one so wished. Later these machines became more elaborate, capable of multiple simple functions, or designed to perform one more complicated function. A famous example is Charles Babbage's difference engine[16] which was a large mechanical calculator that would tabulate polynomial functions developed in the early 1800s.

Eventually in the 20th century electronic computers were created and soon replaced both mechanical and human calculators. Such electronic machines had many benefits over both their human and mechanical counterparts, and soon it became common place to use electronic computers to perform mathematical computations. Today computers have become faster and smaller, and the average person's phone outstrips the entire computing power of NASA during the Apollo missions.

However despite the speed of the calculations these modern computers still need to be instructed in how to evaluate the functions asked of it. This document will take some common functions that any calculator will answer in the blink of an eye accurate to around 10 significant digits, and explore how they may be computed. In particular this document will be comparing the speed at which these computations can be performed versus the accuracy of their results.

1.1 Code and Computers used

During this project I will be discussing the implementation of various algorithms. I will be implementing these algorithms in the C programming language, using the C11 standard.

I chose the C programming language to implement my algorithms in, because once it compiles to binary machine code, the programs produced tend to be very efficient. This is partly due to the low-level of C programming, having relatively close control over direct CPU actions; however this does come at the cost of losing higher functionality that many other programming languages offer. A second reason for the efficiency is due to C's long history, originally being developed in 1969-1970, which has lead to several very efficient compilers being developed.

I will be implementing most programs using C's built in primitive types, typically `int`, `unsigned int` and `double`. On a computer an `int` is an integer that can represent both positive and negative bits using twos complement, this gives an `int` using n bits a minimum value of -2^n and a maximum value of $2^n - 1$. Typically a computer will store an `int` as 32 bits, though some computers may use more or less bits. An `unsigned int` is very similar to an `int`, but does not represent negative values, and thus an `unsigned int` of n bits has a minimum value of 0 and a maximum value of $2^{n+1} - 1$.

If an integer of a specific number of bits is needed then the header `stdint.h` may be used which defines `int_N` and `uint_N` which respectively represent `int` of N bits and `unsigned int` of N bits; The typical values of N are 8, 16, 32 and 64.

In C a `double` is a floating point representation of a real value, that typically follows the IEEE 754 standard[4] for double-precision binary floating points. This standard has:

- 1 bit for the sign of the number, s
- 11 bits for the exponent, e
- 52 bits for the significand, $b = b_0b_1b_2 \dots b_{51}$
- A value that is understood to be:

$$(-1)^s \left(1 + \sum_{i=1}^{52} b_{52-i} 2^{-i} \right) \times 2^{e-1023}$$

This gives a `double` value a precision of around 15-17 significant decimal digits. While this is good for most applications, there are some applications where we may desire even more precision than this. To solve this I will be implementing certain algorithms using the GNU Multiple Precision Arithmetic Library[24] (referred to as GMP) as well as GNU MPFR Library[23] (referred to as MPFR), which was built upon GMP to correct and optimise the original. These libraries allow the use of arbitrary precision real values, given enough memory space, as well as integers longer than C's standard integer types can hold.

An important point to note that will be useful later on is that due to the storage structure of C's `double`s and the MPFR `mpfr_t`s which also use a floating point representation. In the storage of the significand both data types work such that the value of b is in the range $[\frac{1}{2}, 1)$. This is useful as it means that if we have a stored value x , then it is very easy to extract $\alpha \in [\frac{1}{2}, 1), \beta \in \mathbb{Z}$ such that $x = \alpha \cdot 2^\beta$; an operation that would usually be equivalent to calculating the non-trivial $\log_2(x)$. The value of this observation will be in restricting the

range over which functions need to be evaluated later in the document.

I will be compiling and testing all of my code on a benchmark machine running a light version of Ubuntu 14.04, using the GNU C Compiler. The specifications of the machine, that may impact performance are:

- An Intel i5-4690K processor running at 4GHz.
 - This processor uses a 64 bit architecture.
- 8Gb of DDR3 RAM
- A modern chipset on the motherboard

2 General Definitions and Theorems

This section will list some general definitions and theorems which will be used throughout the document. This will not be an exhaustive or in depth view of such concepts but merely an overview to allow easier reading of the material moving forwards.

2.1 Methods

In this document we will look at various functions, such as root functions and trigonometric functions, among others. Despite the variety of functions being analysed there are several methods that are useful for more than one function, or are worth analysing before their use.

2.1.1 Newton-Raphson Method

The Newton-Raphson Method is named after Sir Isaac Newton and Joseph Raphson[7, p. 84]. It is a method that takes a continuously differentiable function f and it's derivative f' , as well as an initial guess x_0 , to create successively more accurate solutions to x where $f(x) = 0$.

The motivation of the method can be seen in figure 2.1.1, where we take an initial guess x_0 of the root x^* . The tangent to the curve above x_0 is then found, and has the equation $y = f'(x_0)(x - x_0) + f(x_0)$, by setting $y = 0$ and solving for x we find x_1 . By repeating this process and starting from a good enough x_0 we hope to find successively closer approximations to x^* .


The specific definition of the Newton-Raphson method that I will be using in this document is below:

Definition 2.1.1.1. Given $f \in C^\infty(\mathbb{R})$, f' being the derivative of f , and $x_0 \in \mathbb{R}$; then we define:

$$x_{n+1} := x_n - \frac{f(x_n)}{f'(x_n)} \quad \forall n \in \mathbb{N}$$

The Newton Raphson method is not suitable for all problems and there are in fact many cases in which it behaves poorly. One such case is when $f'(x_n) \approx 0$ as the value of x_{n+1} will be very close to x_n and thus $f'(x_{n+1}) \approx 0$. Further bad choices of x_0 can lead to the method diverging or entering cycles between two points indefinitely, however we will see that we do not need to be concerned with these issues for our uses of the method.

Figure 2.1.1: Demonstration of Newton-Raphson Method



2.1.2 Taylor Series Expansion

The Taylor Series formulation was created by Brook Taylor in 1715[22], based off of the work of Scottish mathematician James Gregory. The Taylor Series describes a method of representing any infinitely differentiable function as an infinite power series.

Definition 2.1.2.1. Given $f : \mathbb{R} \rightarrow \mathbb{R}$ which is infinitely differentiable on an open interval \mathcal{I} centred at $a \in \mathbb{R}$, we define the Taylor Series of f on \mathcal{I} to be:

$$\sum_{n=0}^{\infty} \frac{f^{(n)}(a)}{n!} (x - a)^n$$

It was shown that on the open interval \mathcal{I} from the above definition we have that $f(x) = \sum_{n=0}^{\infty} \frac{f^{(n)}(a)}{n!} (x - a)^n$, i.e. a function is equal to its Taylor polynomial on the interval for which it is defined. We can then use this fact to define a polynomial that will approximate our function f at $x \in \mathcal{I} \subset \mathbb{R}$

Definition 2.1.2.2. Given $f : \mathbb{R} \rightarrow \mathbb{R}$ which has a Taylor Series of $\sum_{n=0}^{\infty} c_n x^n$, we define the Taylor Polynomial of degree $N \in \mathbb{N}$ to be

$$p_N(x) := \sum_{n=0}^N c_n x^n = c_0 + c_1 x + c_2 x^2 + \cdots + c_N x^N$$

A commonly used type of Taylor series is the Maclaurin series which is a Taylor series in an interval around $a = 0$. Thus a Maclaurin series has the form:

$$\sum_{n=0}^N \frac{f^{(n)}(0)}{n!} x^n$$

Some examples of simple Maclaurin Series are:

$$\frac{1}{1-x} = \sum_{n=0}^{\infty} x^n \quad \forall x \in (-1, 1) \quad (2.1.1)$$

$$(1+x)^k = \sum_{n=0}^{\infty} \binom{k}{n} x^n \quad \forall x \in (-1, 1), k \in \mathbb{N} \quad (2.1.2)$$

2.2 Errors

The error of an approximation \tilde{v} for some v is a measure of how much \tilde{v} differs from v . We will use the error of approximations to discuss the convergence of methods as well as describing their accuracy.

There are several ways of evaluating the error of an approximation which each have their own uses. The error measures that we will use in this document are detailed below:

Definition 2.2.1. If we have a value v and it's approximation \tilde{v} , then the absolute error is

$$\epsilon := |v - \tilde{v}|$$

The absolute error is useful in guaranteeing a certain level of accuracy that a given implementation of a method will give; for instance if $\epsilon < 10^{-3}$ then the approximation is accurate to at least 3 decimal places. Uses of absolute error in the document will use ϵ as their absolute error variable.

As the absolute error of an approximation is hard or impossible to accurately calculate during program execution, we need a way to estimate it. Typically our computations will produce a sequence of approximations x_0, x_1, x_2, \dots , and thus we define the following:

Definition 2.2.2. If we have the sequence $(x_n : n \in \mathbb{N})$, then the iteration error is defined as:

$$\delta_n := |x_n - x_{n-1}|$$

While it is often impossible to calculate ϵ_n it is very easy to calculate δ_n from the generated approximations. This estimate is best used when we know that the convergence is rapid, as in these cases δ_n is a good approximation of ϵ_n .

2.3 Convergence

As our methods of approximating functions will typically generate a sequence of values x_0, x_1, x_2, \dots then we want to ensure that the approximations are approaching the correct value. We consider here what it means for a sequence to converge to a limit value, and some useful results for later chapters.

Definition 2.3.1. A sequence $(x_n \in \mathbb{R} : n \in \mathbb{N})$ converges to x uniformly if

$$\forall \tau \in \mathbb{R}_0^+ \exists N \in \mathbb{N} \text{ s.t. } \epsilon_n := |x - x_n| < \tau \quad \forall n \in [N, \infty) \cap \mathbb{Z}$$

Remark 2.3.1.1. We will typically use the notation that $\lim_{n \rightarrow \infty} |x_n - x| = 0$, to denote that $(x_n : n \in \mathbb{N})$ converges to x .

Theorem 2.3.1. $(x_n \in \mathbb{R} : n \in \mathbb{N})$ converges to x uniformly if and only if

$$\forall \tau \in \mathbb{R}_0^+ \exists N \in \mathbb{N} \text{ s.t. } |x_n - x_m| < \tau \quad \forall m, n \in [N, \infty) \cap \mathbb{Z}$$

Proof. For \implies :

Suppose that $(x_n : n \in \mathbb{N})$ converges to x uniformly. Then

$$\forall \tau \in \mathbb{R}_0^+ \exists N \in \mathbb{N} \text{ s.t. } |x_n - x| < \tau \quad \forall n \in [N, \infty) \cap \mathbb{Z}$$

Thus suppose $N \in \mathbb{N}$ is such that $|x_n - x| < \frac{\tau}{2} \quad \forall n \in [N, \infty) \cap \mathbb{Z}$.

Then if $n, m \geq N$ we see that

$$|x_n - x_m| \leq |x_n - x| + |x_m - x| \leq \tau$$

For \Leftarrow :

Omitted for brevity. □

We have shown now what it means for a value to converge to a limit, but not all sequences that approach a limit do so at the same pace. For example if we consider the sequences $x_n := 2^{-n}$ and $y_n := 10^{-n}$, then it is obvious that the limit of both sequences is 0, but y_n approaches the limit faster. This leads to the following definition of the rate of convergence.

Definition 2.3.2. If $(x_n \in \mathbb{R} : n \in \mathbb{N})$ is a sequence that converges to x , then it is said to converge:

- Linearly if $\lambda \in \mathbb{R}^+$ and

$$\lim_{n \rightarrow \infty} \frac{|x_{n+1} - x|}{|x_n - x|} = \lambda$$

- Quadratically if $\lambda \in \mathbb{R}^+$ and

$$\lim_{n \rightarrow \infty} \frac{|x_{n+1} - x|}{|x_n - x|^2} = \lambda$$

- Order $\alpha \in \mathbb{R}_0^+$ if $\lambda \in \mathbb{R}^+$ and

$$\lim_{n \rightarrow \infty} \frac{|x_{n+1} - x|}{|x_n - x|^\alpha} = \lambda$$

The higher the order of convergence of a sequence the faster it approaches its limit, therefore we are looking for algorithms with high orders of convergence. Many regular series have linear convergence and quadratic convergence is typically very rapid, while orders above quadratic are hard to construct for useful functions.

A useful result is that, under the correct circumstances, the Newton-Raphson method can be shown to have quadratic convergence. The following proof assumes that $\epsilon_n := |x^* - x_n|$:

Theorem 2.3.2. Let f be a twice differentiable function, x^* be a solution to $f(x) = 0$ and $(x_n : n \in \mathbb{N})$ be a sequence produced by the Newton-Raphson Method from some initial point x_0 . If the following are satisfied, then $(x_n : n \in \mathbb{N}_0)$ converges quadratically to x^* :

NR₁: $f'(x) \neq 0 \quad \forall x \in I := [x^* - r, x^* + r]$, where $r \in [|x^* - x_0|, \infty)$

NR₂: $f''(x)$ is continuous $\forall x \in I$

NR₃: $M|\epsilon_0| < 1$ where $M := \sup \left\{ \left| \frac{f''(x)}{f'(x)} \right| : x \in I \right\}$

Proof. By Taylor's Theorem with Lagrange Remainders[22, p. 80] we have that

$$0 = f(x^*) = f(x_n) + (x^* - x_n)f'(x_n) + \frac{1}{2}(x^* - x_n)^2 f''(y_n)$$

where $0 < |x^* - y_n| < |x^* - x_n|$.

Then we get the following derivation:

$$\begin{aligned} f(x_n) + (x^* - x_n)f'(x_n) &= -\frac{1}{2}(x^* - x_n)^2 f''(y_n) \\ \implies \left(\frac{f(x_n)}{f'(x_n)} - x_n \right) + x^* &= -\frac{1}{2} \frac{f''(y_n)}{f'(x_n)} (x^* - x_n)^2 \quad \text{as NR}_3 \implies f'(x_n) \neq 0 \\ \implies x^* - x_{n+1} &= -\frac{1}{2} \frac{f''(y_n)}{f'(x_n)} (x^* - x_n)^2 \\ \implies \epsilon_{n+1} &= \frac{1}{2} \left| \frac{f''(y_n)}{f'(x_n)} \right| \epsilon_n^2 \quad \text{by taking absolute values} \end{aligned}$$

As NR₂ holds then M exists and is positive, and therefore we have:

$$\epsilon_n \leq M\epsilon_{n-1}^2 \leq M^{2^n-1}\epsilon_0^{2^n}$$

We now aim to show that we have convergence, i.e. $\lim_{n \rightarrow \infty} x_n = x^*$; to do this it suffices to show that $\lim_{n \rightarrow \infty} \epsilon_n = 0$.

Consider the sequence $(z_n := M^{2^n-1}\epsilon_0^{2^n} : n \in \mathbb{N}_0)$. We know that $0 \leq \epsilon_n \leq z_n \forall n \in \mathbb{N}_0$, so it then follows that if $\lim_{n \rightarrow \infty} z_n = 0$, then $\lim_{n \rightarrow \infty} \epsilon_n = 0$ by the Squeeze Theorem[20, p. 909].

Now as $M\epsilon_0 < 1$ by NR₃, then we see that:

$$\begin{aligned} \lim_{n \rightarrow \infty} z_n &= \lim_{n \rightarrow \infty} (M\epsilon_0)^{2^n-1} \epsilon_0 \\ &= \epsilon_0 \lim_{n \rightarrow \infty} (M\epsilon_0)^{2^n-1} \\ &= \epsilon_0 \cdot 0 \quad \text{because } M|\epsilon_0| < 1 \\ &= 0 \end{aligned}$$

Now to show that this sequence converges quadratically we see that $\epsilon_{n+1} = \frac{1}{2} \left| \frac{f''(y_n)}{f'(x_n)} \right| \epsilon_n^2$, and therefore $\frac{\epsilon_{n+1}}{\epsilon_n^2} = \frac{1}{2} \left| \frac{f''(y_n)}{f'(x_n)} \right|$.

Because $|x^* - y_n| < |x^* - x_n|$ and $\lim_{n \rightarrow \infty} x_n = x^*$, then it follows that $\lim_{n \rightarrow \infty} y_n = x^*$. Therefore we see that

$$\lim_{n \rightarrow \infty} \frac{\epsilon_{n+1}}{\epsilon_n} = \frac{1}{2} \left| \frac{f''(x^*)}{f'(x^*)} \right| \in \mathbb{R}^+$$

Hence as the above limit exists and is positive then the sequence is quadratically convergent. \square

2.4 Efficiency Metrics

Now that we have discussed how to measure the accuracy of our results by their errors, we wish to consider the efficiency method. There is typically a trade-off between accuracy and efficiency in that to gain a more accurate result, more calculations are required thus taking up more resources. In general however, we will be using efficiency metrics to compare how efficient two different algorithms are at getting the same result.

There are two main ways in which we will measure the efficiency of an algorithm. The first of these methods is the theoretical complexity of the algorithm, which represents the number of steps/operations an algorithm needs to achieve it's goal. The complexity of an algorithm is denoted by the big O notation, which represents the order of the complexity, i.e. the highest order term in the number of operations required.

Typically the execution of an algorithm depends on the size of the input and so if we consider that an input has size n we can discuss different complexities. The first consideration is that if one algorithm takes $2n$ operations while another takes $20n$ operations, then both algorithms have a complexity of $\mathcal{O}(n)$.

A complexity of $\mathcal{O}(n)$ is not a bad complexity for an algorithm as the number of operations needed rises linearly with the size of the input. Complexities of $\mathcal{O}(n^2)$, $\mathcal{O}(2^n)$ and $\mathcal{O}(n!)$ are all poor complexities for an algorithm[8] with the latter two becoming infeasible for larger n . On the other hand complexities better than $\mathcal{O}(n)$ include $\mathcal{O}(\log(n))$ and $\mathcal{O}(1)$, the latter of these is particularly significant as it means that the algorithm takes the same number of steps regardless of the size of the input.

The second method of assessing efficiency consists of timing of functions during execution. This method directly observes how long it takes a computer to perform the calculations for a given algorithm and can be used to empirically test the speed of two algorithms. One remark is that due to the speed of modern computers it is infeasible to time the execution of a single function, and one typically times the same algorithm with the same input being calculated multiple times to get accurate and measurable timings.

3 Root Functions

Root functions are a vital part of mathematics and have been used for millennia, originally studied for their useful relation to architecture; root functions also have many modern day applications. The majority of this section will be dealing with the commonly used square root function \sqrt{N} , which always gives an irrational answer if N is not a square number.

We will consider several methods for approximating root functions, but for our purposes here we are only going to consider roots of $N \in \mathbb{R}_0^+$, this is because if $N \in \mathbb{R}^-$ then it follows that $\sqrt{N} = i\sqrt{|N|}$.

3.1 Digit by Digit Method

The first method we will examine is an old method, which was used to accurately generate the square root of numbers one digit at a time. This method differs from others discussed as

it generates each digit of the root with perfect accuracy, one at a time, thus in a theoretical sense this algorithm is the most accurate of the methods we will view; we will see however that this method is slow.

Now suppose we are looking for \sqrt{N} , then we know that $\sqrt{N} = a_0 10^n + a_1 10^{n-1} + a_2 10^{n-2} + \dots$ for some $n \in \mathbb{Z}$; it then follows that $N = (a_0 10^n + a_1 10^{n-1} + a_2 10^{n-2} + \dots)^2$. By expanding the quadratic value we get that

$$N = a_0^2 10^{2n} + (20a_0 + a_1)a_1 10^{2n-2} + (20(a_0 10 + a_1) + a_2)a_2 10^{2n-4} + \dots \\ + \left(20 \sum_{i=0}^{k-1} a_i 10^{k-i-1} + a_k \right) a_k 10^{2n-2k}$$

An observation should be made regarding the value of n that we use for the theorem. We could of course try different values of n , in some structured procedure, that will find the largest n such that $10^n \leq N$. However we can note that $\log_{10}(\sqrt{N}) = \frac{1}{2} \log_{10}(N)$, thus $10^{\frac{1}{2} \log_{10}(N)} = \sqrt{N}$. Using this information, and the fact that $n \in \mathbb{Z}$, we can have $n := \lfloor \frac{1}{2} \log_{10}(N) \rfloor$.

This allows us to get successive approximations of N where $N_0 = a_0^2 10^{2n}$, $N_1 = N_0 + (20a_0 + a_1)a_1 10^{2n-2}$, $N_2 = N_1 + (20(a_0 10 + a_1) + a_2)a_2 10^{2n-4}$. This will allow us to create an algorithm that will give successive approximations of $\sqrt{N} = a_0 10^n + a_1 10^{n-1} + \dots$, more importantly each approximation will give us the exact next digit in the decimal representation of \sqrt{N} .

Thus we can have an iterative method to solve the problem, where at each stage we are trying to find the largest digit which satisfies the inequality $(20 \sum_{i=0}^{k-1} a_i 10^{k-i-1} + a_k) a_k 10^{2n-2k} \leq N - N_{k-1}$. Thus we get the following pseudo-code, which outputs two sequences, one indicating the digits before the decimal point and one afterwards. I will use set notation to indicate the sequences, but in this case order is important and repetition is allowed.

Algorithm 3.1.1: Exact Digit by Digits Square Root

```

1  exactRootDigits( $N \in \mathbb{R}_0^+, d \in \mathbb{N}$ ):
2       $Digits_a := \emptyset$ 
3       $Digits_b := \emptyset$ 
4       $k := 0$ 
5       $n := \lfloor \frac{1}{2} \log_{10}(N) \rfloor$ 
6      while  $k < d$ :
7           $a_k := \max \left\{ t \in [0, 9] \cap \mathbb{Z} : \left( 20 \sum_{i=0}^{k-1} a_i 10^{k-i-1} + t \right) t 10^{2n-2k} \leq N \right\}$ 
8           $N \mapsto N - \left( 20 \sum_{i=0}^{k-1} a_i 10^{k-i-1} + a_k \right) a_k 10^{2n-2k}$ 
9          if  $n - k < 0$ :
10              $Digits_b \mapsto Digits_b \cup \{a_k\}$ 
11          else:
12              $Digits_a \mapsto Digits_a \cup \{a_k\}$ 
13              $k \mapsto k + 1$ 
14          if  $Digits_a = \emptyset$ :
15              $Digits_a := \{0\}$ 
16          if  $Digits_b = \emptyset$ :
```

```

17       $Digits_b := \{0\}$ 
18      return  $(Digits_a, Digits_b)$ 

```

This method has a computational complexity of $\mathcal{O}(d^2)$, as each loop requires the operations of summing k elements, and the loop is repeated for $k \in [0, d] \cap \mathbb{Z}$. We will see that by considering some changes to the algorithm we can change the complexity class to be $\mathcal{O}(d)$.

First we will note that line 5 is not an issue, as if we only care about the first significant digit of $\frac{1}{2}\log_{10}(N)$, then this is $\mathcal{O}(|\log(N)|)$. This can be seen as if we start from $n = 0$ we can either count up or down until we find 10^{2n} at most or at least N , respectively. This obviously takes at most $|\log_{10}(N)|$ steps, giving us our stated complexity. We will also assume that $\mathcal{O}(|\log(N)|) \leq \mathcal{O}(d)$, as we have already seen that we can manipulate our input N to be within a reasonable range.

Second we note that on line 7 we calculate $\sum_{i=0}^{k-1} a_i 10^{k-i-1}$ for each value of t ; we can reduce the complexity of this line by pre-calculating this value. However we can do even better if we consider that at step $k + 1$ we are calculating $\sum_{i=0}^k a_i 10^{k-i} = a_k + 10 \sum_{i=0}^{k-1} a_i 10^{k-i-1}$. Thus if we introduce $P_0 := 0$, and for each k we calculate $P_{k+1} := 10P_k + a_k$, then we can reduce the complexity from $\mathcal{O}(k)$ to $\mathcal{O}(1)$.

This calculation of P_k , then carries over to reduce the complexity of line 8 to be $\mathcal{O}(1)$ instead of $\mathcal{O}(k)$. Combining this we can create the modified algorithm below:

Algorithm 3.1.2: Exact Digit by Digits Square Root version 2

```

1  exactRootDigits_v2 ( $N \in \mathbb{R}_0^+, d \in \mathbb{N}$ ):
2       $Digits_a := \emptyset$ 
3       $Digits_b := \emptyset$ 
4       $k := 0$ 
5       $n := \lfloor \frac{1}{2}\log_{10}(N) \rfloor$ 
6       $P_0 := 0$ 
7      while  $k < d$ :
8           $a_k := \max \{t \in [0, 9] \cap \mathbb{Z} : (20P_k + t) 10^{2n-2k} \leq N\}$ 
9           $N \mapsto N - (20P_k + a_k) a_k 10^{2n-2k}$ 
10          $P_{k+1} := 10P_k + a_k$ 
11         if  $n - k < 0$ :
12              $Digits_b \mapsto Digits_b \cup \{a_k\}$ 
13         else:
14              $Digits_a \mapsto Digits_a \cup \{a_k\}$ 
15          $k \mapsto k + 1$ 
16     if  $Digits_a = \emptyset$ :
17          $Digits_a := \{0\}$ 
18     if  $Digits_b = \emptyset$ :
19          $Digits_b := \{0\}$ 
20     return  $(Digits_a, Digits_b)$ 

```

This method is useful, but can be difficult to implement as it requires high precision for the representation of the real value of N . In my implementation using C, I utilised the MPFR library to utilise high precision integers, but still encountered issues regarding loss of precision.

As an example the table below shows the number of digits of accuracy I was able to calculate for $\sqrt{2}$ using the above algorithm, compared to the number of bits of precision used in the calculations.

Bits of Precision	Maximum Accuracy
8	2
16	5
32	9
64	18
128	39
256	77
512	154
1024	308
2048	615
4096	1234
8192	2466

This data is highly structured and so we can hope to create a simple function that would allow us to calculate how much precision would be needed for a given number of digits of accuracy, at least for single digit inputs for N . We can see that the average ratio of Precision to Accuracy is 3.41259..., which ranges from 3.31928... to 4.0. From this we can draw a general trend that Digits of Accuracy $\approx 3.4 \times$ Bits of Precision; thus if we take the more generous assumption that Digits of Accuracy $4 \times$ Bits of Precision, we can use this to pre-determine the accuracy needed.

It should be noted that to ensure accuracy we should over-estimate the required precision, however if we overestimate the precision, then our calculations will be performed using unnecessarily large data structures and thus computation time will increase.

One particular use of this technique is to find an approximation of a square root to it's integer part, calculated in base 2. This algorithm is of note as we will see that it has a computation time of $\mathcal{O}(1)$.

The algorithm uses the same basis as the base 10 version, for it's calculations, but due to the nature of being in binary several changes can be made for computational efficiency. To do this we will view the problem as follows: if we know some $r \in \mathbb{Z}_0^+$ which is our current approximation of our root, we are looking for some $e \in \mathbb{Z}_0^+$ such that $(r + e)^2 \leq N$. Expanding this out we get $r^2 + 2re + e^2 \leq N$, and if we keep track of $M = N - r^2$, we can test if $2re + e^2 \leq M$.

Now we can consider our choice of e , the most practical method is to test successive $e_m := 2^m$, where m is descending starting with $m = \max m \in \mathbb{Z}_0^+ : 4^m \leq N$. We can use an iterative formula to build up the integer square root, where we start with $r = 0, M = N$ and have $rr + e_m$ whenever $2re_m + e_m^2 \leq M$, stopping when $m < 0$. This is then implemented as follows:

Algorithm 3.1.3: Integer Square Root Algorithm

```

1  integerSquareRoot( $N \in \mathbb{Z}_0^+$ ):
2       $M := N$ 
```

```

3       $m := \max m \in \mathbb{Z}_0^+ : 4^m \leq M$ 
4       $r := 0$ 
5      while  $m \geq 0$ :
6          if  $2r(2^m) + 4^m \leq M$ :
7               $M \mapsto M - 2r(2^m) + 4^m$ 
8               $r \mapsto r + 2^m$ 
9               $m \mapsto m - 1$ 
10     return  $r$ 

```

If we now consider an implementation of the above algorithm using an unsigned integer system with K bits, where $2|K$. We will use `res` to represent $2re_m$, which means at the start of the algorithm we will have `res = 0`; similarly we can use `bit` to represent e_m^2 . As we know that K bits are used and $2|K$, it then follows that the largest power of 4 less than the maximum representable value ($2^K - 1$ is 2^{K-2} , which can be calculated as `bit = 1 << (K - 2)` using bit shift operations. Finally we will use `num` to represent M .

Now that we have discussed the set-up we can consider how to implement some of the steps above. First to implement line 3 we can simply keep dividing `bit` by 4 while `bit > num`, which can be efficiently implemented as `bit >> 2` by using bit shifts in place of division by powers of 2. The same technique can be used in place of line 9, which leads us to re-evaluating our usage of line 5. As we are using bit shifting and a bit shift that would take a number past 0 instead results in 0, we also know that $2|K$ and so eventually we will reach `bit == 1`, which represents $m = 0$; therefore we can use `bit > 0` as our stopping criteria on line 5.

Line 6 is easy to convert, given our definitions of `res`, `bit` and `num`, as is line 7. All that remains is to consider how to update `res`, which has two different ways of being updated depending on whether `res + bit <= num`. If it is false that `res + bit <= num`, then we wish for `res` to represent $2re_{m-1}$; this is easily achieved if we consider that $2re_{m-1} = \frac{1}{2}(2re_m)$, which prompts the update `res = res >> 1`. For the second case, when `res + bit <= num` is true, we want `res` to represent $2(r + e_m)e_{m-1}$; to implement this we consider the following derivation:

$$\begin{aligned}
 2(r + e_m)e_{m-1} &= \frac{1}{2} \cdot 2(r + e_m)e_m \\
 &= \frac{1}{2} \cdot 2(re_m + e_m^2) \\
 &= \frac{1}{2}(2re_m) + e_m^2
 \end{aligned}$$

Using this above derivation we see that we can calculate this as `res = (res >> 1) + bit`. Below is a simple implementation of this in C using the unsigned 32 bit integer type `uint32_t`. A more commented and slightly modified version can be found in Appendix .

```

1  uint32_t int_sqrt(uint32_t num)
2  {
3      uint32_t res = 0, bit = (1 << 30);
4
5      while (bit > num)
6          bit = bit >> 2;
7
8      while (bit > 0)
9      {

```

```

10         if (res + bit <= num)
11         {
12             num = num - (res + bit);
13             res = (res >> 1) + bit;
14         }
15         else
16             res = res >> 1;
17
18         bit = bit >> 2;
19     }
20
21     return res;
22 }

```

We should consider the final step of the loop, when `bit == 1`. In this case when `res` is updated we have `res` represent either $2(r+e_0)e_{-1} = r+e_0$, or $2re_{-1} = r$; thus the algorithm exits with the correct value.

Now that the algorithm is correctly constructed using simple unsigned integer addition, subtraction and bit shifting (which we can assume all have computational time of $\mathcal{O}(1)$), we can look at the worst case complexity of the algorithm:

- The complexity of the set up of variables is constant time.
- The worst case complexity would be to to have `bit <= num` at the start.
- The loop would execute 16 times for our 32 bit integers, and contains a single operation which is $\mathcal{O}(1)$ complexity.
 - The worst case within the loop is to have `res + bit <= num` for each iteration.
 - Within the first `if` branch there are a constant 4 operations.
 - Each loop has an additional operation operation to update `bit`.
 - This makes 5 operations per loop, giving $\mathcal{O}(1)$ complexity within the loops.

Therefore we see that the algorithm has $\mathcal{O}(1)$ time complexity, and even has the same in storage complexity. In particular our 32 bit example requires 163 operations, including assignments, comparisons and calculations. This means that the integer square root of any number up to 4294967295 can be calculated extremely quickly.

3.2 Bisection Method

The Bisection Method is a general method for approximating the zero, α , of a function, f , on a bounded interval, $I := [a, b]$, where f has the property $f(x)f(y) < 0 \forall (x, y) \in [a, \alpha) \times (\alpha, b]$; we may assume, without loss of generality, that $f(x) < 0 \forall x \in [a, \alpha]$.

The bisection method starts with initial bounds $a_0 = a, b_0 = b$, where the initial approximation for the root is $x_0 = \frac{1}{2}(a + b)$. We will consider pseudo-code of the iteration process, that uses $b_n - a_n < \tau$ or $f(x_n) = 0$ as exit criteria. Here τ is a tolerance threshold, and if the exit criteria is met it means that $|x_n - \alpha| \leq \frac{\tau}{2}$, while the other exit criteria means we have reached an exact solution.

Algorithm 3.2.1: General Bisection Method

```

1  bisectionMethod ( $a \in \mathbb{R}, b \in (a, \infty), f \in \mathcal{C}[a, b], \tau \in \mathbb{R}^+$ )
2       $a_0 := a$ 
3       $b_0 := b$ 
4       $x_0 := \frac{1}{2}(a + b)$ 
5       $n := 0$ 
6      while  $f(x_n) \neq 0$  AND  $b_n - a_n > \tau$ :
7          if  $f(x_n) < 0$ :
8               $a_{n+1} := x_n$ 
9               $b_{n+1} := b_n$ 
10         else:
11              $a_{n+1} := a_n$ 
12              $b_{n+1} := x_n$ 
13          $n \mapsto n + 1$ 
14          $x_n := \frac{1}{2}(a_n + b_n)$ 
15     return  $x_n$ 

```

For our purposes we are trying to find the zero of $f(x) = x^2 - N$, which is a strictly increasing function on \mathbb{R}_0^+ . If $N \geq 1$, then $\sqrt{N} \in [0, N]$, while $N < 1 \implies \sqrt{N} \in [0, 1]$. It is obvious that our function has the required property, and thus we get the following method for finding the square root of N :

Algorithm 3.2.2: Bisection Method for Square Roots

```

1  bisectionSquareRoot ( $N \in \mathbb{R}_0^+, \tau \in \mathbb{R}^+$ )
2       $a_0 := 0$ 
3       $b_0 := \max 1, N$ 
4       $x_0 := \frac{1}{2}(a_0 + b_0)$ 
5       $n := 0$ 
6      while  $x_n^2 - N \neq 0$  AND  $b_n - a_n > \tau$ :
7          if  $x_n^2 - N < 0$ :
8               $a_{n+1} := x_n$ 
9               $b_{n+1} := b_n$ 
10         else:
11              $a_{n+1} := a_n$ 
12              $b_{n+1} := x_n$ 
13          $n \mapsto n + 1$ 
14          $x_n := \frac{1}{2}(a_n + b_n)$ 
15     return  $x_n$ 

```

The implementation of this method is efficiently achieved in C using only addition, subtraction and multiplication by a constant. Before this method is implemented, however, we must first consider if and or when it converges to the correct answer. From an intuitive standpoint we would assume that if there is only one root in the interval, it would follow that we would converge to the root.

Proposition 3.2.1. $\lim_{n \rightarrow \infty} x_n = \sqrt{N}$ for Algorithm 3.2.2

Proof. To prove this statement it suffices to prove that $\sqrt{N} \in [a_n, b_n] \forall n \in \mathbb{N}$ and $\lim_{n \rightarrow \infty} |x_n - \sqrt{N}| = 0$.

Claim 1: $\sqrt{N} \in [a_n, b_n] \forall n \in \mathbb{N}$

Proof. $a_0 := 0 \implies a_0 \leq \sqrt{N}$

$b_0 := \max\{1, N\} \implies b_0 \geq \sqrt{N}$

Therefore it is obvious that $\sqrt{N} \in [a_0, b_0]$

Now suppose $\sqrt{N} \in [a_n, b_n]$ for some $n \in \mathbb{N}$

It should be noted that $a_n, b_n, x_n \in \mathbb{R}_0^+ \forall n \in \mathbb{N}$ as $a_0, b_0 \in \mathbb{R}_0^+$ and all the subsequent values are derived from these using only addition and multiplication by positive factors.

We then see that $x_n := \frac{1}{2}(a_n + b_n)$, and we consider the two cases that $x_n^2 - N \leq 0$ or $x_n^2 - N \geq 0$.

Case $x_n^2 - N \leq 0$:

$a_{n+1} := x_n, b_{n+1} := b_n$

It is therefore obvious that $\sqrt{N} \leq b_{n+1}$.

Now we see that $x_n^2 - N \leq 0 \implies x_n^2 \leq N \implies x_n \leq \sqrt{N}$ as all the values are non-negative.

Thus $\sqrt{N} \in [a_{n+1}, b_{n+1}]$.

Case $x_n^2 - N \geq 0$:

$a_{n+1} := a_n, b_{n+1} := x_n$

It is therefore obvious that $\sqrt{N} \geq a_{n+1}$.

Now we see that $x_n^2 - N \geq 0 \implies x_n^2 \geq N \implies x_n \geq \sqrt{N}$ as all the values are non-negative.

Thus $\sqrt{N} \in [a_{n+1}, b_{n+1}]$.

Hence $\sqrt{N} \in [a_n, b_n] \implies \sqrt{N} \in [a_{n+1}, b_{n+1}] \forall n \in \mathbb{N}$

As $\sqrt{N} \in [a_0, b_0]$ then we see that $\sqrt{N} \in [a_n, b_n] \forall n \in \mathbb{N}$ ■

Claim 2: $\lim_{n \rightarrow \infty} |x_n - \sqrt{N}| = 0$

Proof. Let $n \in \mathbb{N}$ be arbitrary.

As $x_n := \frac{1}{2}(a_n + b_n)$ then we see that $|a_n - x_n| = |b_n - x_n| = \frac{1}{2}(b_n - a_n)$.

Now as $\sqrt{N} \in [a_n, b_n]$ it follows that $|\sqrt{N} - x_n| \leq \frac{1}{2}(b_n - a_n)$.

As the modulus function is a mapping from \mathbb{R} to \mathbb{R}_0^+ , it is clear that $|\sqrt{N} - x_n|$ is bounded below by 0.

Now as for each $n \in \mathbb{N}$, either $a_{n+1} = x_n$ or $b_{n+1} = x_n$, we see that $b_{n+1} - a_{n+1} = \frac{1}{2}(b_n - a_n)$.

Further we can see that $b_n - a_n \geq 0 \forall n \in \mathbb{N}$ because $b_n \geq a_n$.

Therefore the sequence of $\frac{1}{2}(b_n - a_n)$ is a strictly decreasing sequence that is bounded below, by 0. Thus $\lim_{n \rightarrow \infty} \frac{1}{2}(b_n - a_n) = 0$

Therefore $\lim_{n \rightarrow \infty} |x_n - \sqrt{N}| = \lim_{n \rightarrow \infty} \frac{1}{2}(b_n - a_n) = 0$ ■

By using our two claims above we see that $\lim_{n \rightarrow \infty} x_n = \sqrt{N}$. □

The algorithm can be generalised to search for $\sqrt[k]{N}$, where $k \in [2, \infty) \cap \mathbb{Z}$. We can do this by using or implementing an integer power function, `intPow`, to use in place of x_n^2 . This gives the following algorithm:

Algorithm 3.2.3: Bisection Method for General Roots

```

1  kRootBisectionMethod( $N \in \mathbb{R}_0^+, k \in [2, \infty) \cap \mathbb{Z}, \tau \in \mathbb{R}^+$ )
2       $a_0 := 0$ 
3       $b_0 := \max 1, N$ 
4       $x_0 := \frac{1}{2}(a_0 + b_0)$ 
5       $n := 0$ 
6      while  $\text{intPow}(x_n, k) - N \neq 0$  AND  $b_n - a_n > \tau$ :
7          if  $\text{intPow}(x_n, k) - N < 0$ :
8               $a_{n+1} := x_n$ 
9               $b_{n+1} := b_n$ 
10         else:
11              $a_{n+1} := a_n$ 
12              $b_{n+1} := x_n$ 
13          $n \mapsto n + 1$ 
14          $x_n := \frac{1}{2}(a_n + b_n)$ 
15     return  $x_n$ 

```

The proof that method converges to the correct root is very similar to the proof of convergence for algorithm 3.2.2; and as such will be omitted here.

We can now consider the accuracy that can be achieved by our algorithm, for our purposes we will be considering \sqrt{N} , though the same applies for $\sqrt[k]{N}$. We know that $\sqrt{N} \in [a_n, b_n] \forall n \in \mathbb{N}$, and in particular we know that either $\sqrt{N} \in [a_n, x_n]$ or $\sqrt{N} \in [x_n, b_n] \forall n \in \mathbb{N}$; therefore we know that $\epsilon_n := |x_n - \sqrt{N}| \leq \frac{1}{2}(b_n - a_n) \forall n \in \mathbb{N}$. Then as we know that $b_{n+1} - a_{n+1} = \frac{1}{2}(b_n - a_n)$, we know that $\epsilon_n \leq \frac{1}{2^n}(b_0 - a_0)$.

We can consider that $\forall N \in \mathbb{R}_0^+ \exists (r, k) \in [\frac{1}{4}, 1) \times \mathbb{Z} : N = r \cdot 2^{2k}$; using this we know that $\sqrt{N} = \sqrt{r} \cdot 2^k$. As we have the fixed initial bounds of $a_0 = 0$ and $b_0 = 1$, then if we are finding \sqrt{r} we know that $\epsilon_n \leq \frac{1}{2^n} \forall n \in \mathbb{N}$. Hence we can calculate the precision of our current estimate beforehand for any $n \in \mathbb{N}$, and thus we can guarantee d significant digits of accuracy for $r \in [\frac{1}{4}, 1)$.

To get this accuracy must find $n \in \mathbb{N}$ such that $\epsilon_n \leq 10^{-d}$, to achieve this we must find $n \in \mathbb{N}$ such that $2^n \geq 10^d$. For example the following table indicates the minimum required n , required for certain significant digits of accuracy.

d	$n : 2^n \geq 10^d$
1	0
5	15
10	30
20	64
50	163
100	329

Now usually finding r and k as above would be as hard as calculating the logarithm of N ; however due to the way that C stores real numbers as either `double` or in the MPFR library, finding these values is actually fairly trivial. Both provide a functionality to find $(a, b) \in [\frac{1}{2}, 1) \times \mathbb{Z} : N = a \cdot 2^b$, and from this we merely require a simple comparison and division by 2

if b is not even. This leads to the following algorithm, which has the above maximum number of iterations for a required accuracy:

Algorithm 3.2.4: Bisection Method for Square Roots with fixed bounds

```

1  bisectionSquareRoot ( $N \in \mathbb{R}_0^+, \tau \in \mathbb{R}^+$ )
2      Let  $(r, e) \in [\frac{1}{2}, 1) : N = r \cdot 2^e$ 
3      if  $2 \nmid e$ :
4           $r \mapsto \frac{r}{2}$ 
5           $e \mapsto e - 1$ 
6       $a_0 := 0$ 
7       $b_0 := 1$ 
8       $x_0 := \frac{1}{2}(a_0 + b_0)$ 
9       $n := 0$ 
10     while  $x_n^2 - N \neq 0$  AND  $b_n - a_n > \tau$ :
11         if  $x_n^2 - N < 0$ :
12              $a_{n+1} := x_n$ 
13              $b_{n+1} := b_n$ 
14         else:
15              $a_{n+1} := a_n$ 
16              $b_{n+1} := x_n$ 
17          $n \mapsto n + 1$ 
18          $x_n := \frac{1}{2}(a_n + b_n)$ 
19     return  $x_n \cdot 2^{\frac{e}{2}}$ 

```

3.3 Newton's Method for Square Roots

If we consider our equation $f(x) = x^2 - N$, then we can see that it is differentiable on $x \in \mathbb{R}^+$ with $f'(x) = 2x$; we can therefore hope to use the Newton-Raphson method to approximate $x^* \in \mathbb{R}^+ : f(x^*) = 0$. Now it is obvious that if $f(x^*) = 0$ then $x^* = \sqrt{N}$ and so the Newton-Raphson method should converge to the \sqrt{N} provided we start at a suitable x_0 .

The iterative step of Newton's method for square roots is $x_{n+1} = x_n - \frac{x_n^2 - N}{2x_n}$ which when implemented in C, requires the calculation of `x = x - (x*x - N) / (2*x)` each iteration, which requires 5 operations. However if we re-arrange our equation, we instead get $x_{n+1} = \frac{1}{2}x_n + \frac{N}{x_n}$, which when implements is `x = 0.5 * (x + N/x)`, which now uses only 3 operations.

We can then use the following pseudo-code as the basis of our implementations of the Newton-Raphson Method for Square Roots:

Algorithm 3.3.1: Basic Newton Method for Square Root

```

1  NewtonSquareRoot ( $N \in \mathbb{R}, x_0 \in \mathbb{R}, \tau \in (0, 1)$ ):
2       $n := 0$ 
3      loop:
4           $x_{n+1} := \frac{1}{2}(x_n + \frac{N}{x_n})$ 
5           $\delta_n := |x_{n+1} - x_n|$ 
6          if  $\delta_n \leq \tau$ :
7              return  $x_{n+1}$ 

```

Next we want to consider our initial estimate x_0 ; it is prudent to first consider when our initial estimate will converge to the correct root. By looking at a graph of the function, and in particular the tangents to the curve, it would seem reasonable to wonder if $\lim_{n \rightarrow \infty} x_n = \sqrt{N}$.

Proposition 3.3.1. *If $x_0 \in \sqrt{N}, \infty$ and $\{x_n : n \in \mathbb{N}\}$ is a sequence of approximations of \sqrt{N} found via the Newton-Raphson Method, as detailed above, then:*

$$\lim_{n \rightarrow \infty} x_n = \sqrt{N}$$

Proof. Suppose $x_n > \sqrt{N}$, then

$$\begin{aligned} x_{n+1} &= \frac{1}{2} \left(x_n + \frac{N}{x_n} \right) \\ &< \frac{1}{2} \left(x_n + \frac{N}{\sqrt{N}} \right) && \text{as } \sqrt{N} < x_n \implies \frac{1}{x_n} < \frac{1}{\sqrt{N}} \\ &= \frac{1}{2} (x_n + \sqrt{N}) \\ &< \frac{1}{2} (2x_n) \\ &= x_n \end{aligned}$$

Therefore we see that $\{x_k : k \in [n, \infty) \cap \mathbb{Z}\}$ is a strictly decreasing sequence.

Now suppose that $x_n \geq \sqrt{N}$ and then, for a contradiction, assume that $x_{n+1} < \sqrt{N}$. We then see that:

$$\begin{aligned} \frac{1}{2} \left(x_n + \frac{N}{x_n} \right) &< \sqrt{N} \\ \implies x_n + \frac{N}{x_n} &< 2\sqrt{N} \\ \implies x_n^2 + N &< 2\sqrt{N}x_n \\ \implies x_n^2 - 2\sqrt{N}x_n + N &< 0 \\ \implies (x_n - \sqrt{N})^2 &< 0 \end{aligned}$$

This is a contradiction as $x_n, \sqrt{N} \in \mathbb{R} \implies (x_n - \sqrt{N})^2 \geq 0$.

Therefore $x_n \geq \sqrt{N} \implies x_{n+1} \geq \sqrt{N}$.

Hence if $x_0 > \sqrt{N}$, then it follows that $\{x_n : n \in \mathbb{N}\}$ is a strictly decreasing sequence that is bounded below. Therefore by an elementary result from limit theory, we see that $\lim_{n \rightarrow \infty} x_n = \inf\{x_n : n \in \mathbb{N}\}$. \square

The most obvious choice for x_0 would be N , but we see that $N \in (0, 1)$, then $N < \sqrt{N}$. In this case, we could choose $x_0 = 1$ for the case that $N \in (0, 1)$. Therefore we can choose

$$x_0 := \begin{cases} N & : N \in (1, \infty) \\ 1 & : N \in (0, 1) \end{cases}$$

In our choice of x_0 , we have so far left out the cases where $N \in \{0, 1\}$. In both of these cases we already know the correct answer, namely $\sqrt{N} = N$ provided $N \in \{0, 1\}$. Therefore we can exclude them from our calculations, as we can pre-asses the value of N , simply returning the correct answer if one of these cases is encountered.

This then leads to an updated version of the above pseudo-code:

Algorithm 3.3.2: Basic Newton Method for Square Root

```

1  NewtonSquareRoot ( $N \in \mathbb{R}_0^+, \tau \in (0, 1)$ ):
2      if  $N \in \{0, 1\}$ :
3          return  $N$ 
4      if  $N > 1$ :
5           $x_0 := N$ 
6      else:
7           $x_0 := 1$ 
8       $n := 0$ 
9      loop:
10          $x_{n+1} := \frac{1}{2}(x_n + \frac{N}{x_n})$ 
11          $\delta_n := |x_{n+1} - x_n|$ 
12         if  $\delta_n \leq \tau$ :
13             return  $x_{n+1}$ 
14          $n \mapsto n + 1$ 

```

An alternative would be to use the integer square root method discussed in Section 3.1 to improve our initial choice of x_0 . We will start by showing, that for intervals $I \subset \mathbb{R}^+$, the first two criteria for quadratic convergence of the Newton Raphson method are met.

Proposition 3.3.2. *If $I \subset \mathbb{R}^+$ then NR_1 and NR_2 are satisfied for $f(x) = x^2 - N$*

Proof. $f(x) = x^2 - N \implies f'(x) = 2x \implies f''(x) = 2$

Now as $x \in \mathbb{R}^+ \forall x \in I$, then it is obvious that $f'(x) > 0$

Therefore $f'(x) \neq 0 \forall x \in I$, and so NR_1 is satisfied.

As $f''(x)$ is a constant function, then it is continuous on all of \mathbb{R} .

Hence $f''(x)$ is continuous $\forall x \in I$ and so NR_2 is satisfied. □

Now the integer square root function will always produce a root that is at most a distance of 1 from \sqrt{N} ; therefore we can consider $I = [\sqrt{N} - 1, \sqrt{N} + 1]$. Now if $N \leq 1$, then $I \not\subset \mathbb{R}^+$ and so we cannot guarantee the satisfaction of NR_1 . Therefore we can proceed with our analysis of the case that $N > 1$.

If $N > 1$ we need to find when we can satisfy NR_3 . First, we remember that $M := \sup \left\{ \left| \frac{f''(x)}{f'(x)} \right| : x \in I \right\}$ and $\epsilon_0 := |x_0 - \sqrt{N}|$. Then to satisfy NR_3 , we must have that $M\epsilon_0 < 1$.

We can guarantee that $\epsilon_0 \leq 1$ because $x_0 \in I$ from the integer square root algorithm; therefore it suffices to find the situation where $M < 1$. As both f' and f'' are continuous and non-zero

on I it follows that $M = \sup\{x^{-1} : x \in I\} = (\sqrt{N} - 1)^{-1}$. We then see that:

$$\begin{aligned} M < 1 &\iff \sqrt{N} - 1 > 1 \\ &\iff \sqrt{N} > 2 \\ &\iff N > 4 \end{aligned}$$

Therefore we can get the following new choice for x_0 , and thus new pseudo-code:

$$x_0 := \begin{cases} 1 & : N \in (0, 1) \\ N & : N \in (1, 4] \\ \text{intSqrt}(N) & : N \in (4, \infty) \end{cases}$$

Algorithm 3.3.3: Basic Newton Method for Square Root

```

1  NewtonSquareRoot( $N \in \mathbb{R}_0^+, \tau \in (0, 1)$ ):
2    if  $N \in \{0, 1\}$ :
3      return  $N$ 
4    if  $N < 1$ :
5       $x_0 := 1$ 
6    else:
7      if  $N \leq 4$ :
8         $x_0 := N$ 
9      else:
10        $x_0 := \text{IntSqrt}(N)$ 
11    $n := 0$ 
12   loop:
13      $x_{n+1} := \frac{1}{2}(x_n + \frac{N}{x_n})$ 
14      $\delta_n := |x_{n+1} - x_n|$ 
15     if  $\delta_n \leq \tau$ :
16       return  $x_{n+1}$ 
17    $n \mapsto n + 1$ 
```

If we consider any $N \in \mathbb{R}_0^+$, then $\exists a \in [\frac{1}{2}, 1), b \in \mathbb{Z} : N = a \times 2^b$. Finding this value would be a hard as finding the logarithm of N base 2, but due to the representation of numbers within C, both standard C and MPFR have functions that allow us to extract these two values with minimal computational expenditure.

This helps as we can then narrow our problem, to only finding $\sqrt{a} : a \in [\frac{1}{2}, 1)$, and then calculating

$$\sqrt{N} = \sqrt{a} \times 2^{\lfloor \frac{b}{2} \rfloor} \times \alpha \text{ where } \alpha = \begin{cases} 1 & : b \in 2\mathbb{Z} \\ \sqrt{2} & : b \in \mathbb{Z}^+ \setminus 2\mathbb{Z} \\ \frac{1}{\sqrt{2}} & : b \in \mathbb{Z}^- \setminus 2\mathbb{Z} \end{cases}$$

Where $\lfloor \cdot \rfloor : \mathbb{R} \rightarrow \mathbb{Z}$ is the nearest integer function. We can then implement the above observations in the following algorithm:

Algorithm 3.3.4: Newton Method for Square Root v3

```

1  NewtonSquareRoot( $N \in \mathbb{R}_0^+, \tau \in (0, 1)$ ):
2    Let  $(a, b) \in [\frac{1}{2}, 1) \times \mathbb{Z}$  s.t.  $N = a \cdot 2^b$ 
3     $x_0 := 1$ 
```

```

4      if  $b \equiv 0 \pmod{2}$ :
5           $\alpha := 1$ 
6      else:
7          if  $b > 0$ :
8               $\alpha := \sqrt{2}$ 
9          else:
10              $\alpha := \frac{1}{\sqrt{2}}$ 
11      $n := 0$ 
12     loop:
13          $x_{n+1} := \frac{1}{2}(x_n + \frac{a}{x_n})$ 
14          $\delta_n := |x_{n+1} - x_n|$ 
15         if  $\delta_n \leq \tau$ :
16             return  $\alpha \cdot x_{n+1} \cdot 2^{\lfloor \frac{b}{2} \rfloor}$ 
17      $n \mapsto n + 1$ 

```

We must first consider the fact that the algorithm requires the pre-calculation of both $\sqrt{2}$ and $\frac{1}{\sqrt{2}}$, to be able to calculate all values. However, it is the case that we can use the algorithm itself to generate these values as $2 = \frac{1}{2} \cdot 2^2$, and as the exponent of 2 is even then the algorithm does not require $\sqrt{2}$ for this computation. Similarly $\frac{1}{2} = \frac{1}{2} \cdot 2^0$, which again is an even exponent. We can thus run our algorithm to find an arbitrarily accurate values for $\sqrt{2}$ and $\frac{1}{\sqrt{2}}$ to allow us to run the algorithm for other values.

With this observation can then consider $N \in [\frac{1}{2}, 1)$. As this is a small range and, as per our previous algorithm, we use an initial guess of $x_0 = 1$, we can then prove that our algorithm will converge quadratically to \sqrt{N} .

Proposition 3.3.3. *Algorithm 3.3.4, satisfies the criteria of Theorem 2.3.2, and thus has quadratic convergence to \sqrt{N} .*

Proof. To fulfil the criteria of Theorem 2.3.2, we must find an interval $I := [\sqrt{N} - r, \sqrt{N} + r]$ for some $r \geq \epsilon_0$.

Consider $\epsilon_0 = |\sqrt{N} - x_0| = 1 - \sqrt{N}$. We see that as $N \geq \frac{1}{2}$ then $\sqrt{N} \geq \sqrt{2}^{-1}$, and thus $\epsilon_0 \leq 1 - \sqrt{2}^{-1}$. Let us have $r := 1 - \frac{1}{\sqrt{2}}$, and I as defined above.

If we look at the lower bound of I , then we see that:

$$\begin{aligned}
 \sqrt{N} - r &\geq \frac{1}{\sqrt{2}} - (1 - \frac{1}{\sqrt{2}}) \\
 &= \frac{2}{\sqrt{2}} - 1 \\
 &= \sqrt{2} - 1 \\
 &> 0
 \end{aligned}$$

Therefore we see that $I \subset \mathbb{R}^+$, and so by Proposition 3.3.2 we get that NR_1 and NR_2 are satisfied. It then remains to show that NR_3 is satisfied on I .

Now by the definition in Theorem 2.3.2, we have that $M = \sup \left\{ \frac{1}{2} \left| \frac{f''(x)}{f'(y)} \right| : x, y \in I \right\}$. We know that I is bounded, $f''(x) = 2$ and $f'(x) = 2x$ meaning that $\frac{1}{2} \left| \frac{f''(x)}{f'(y)} \right| = \frac{1}{f'(x)}$ as $x \in \mathbb{R}^+$.

Therefore our problem is reduced to finding $\max \left\{ \frac{1}{2x} : x \in I \right\}$, which is equivalent to finding $\min \{x : x \in I\} = \sqrt{N} - r$. Therefore by passing this information back up the chain we get that

$$M = \frac{1}{2(\sqrt{N} - r)}$$

Then we see that:

$$\begin{aligned} M_{\epsilon_0} &= \frac{1 - \sqrt{N}}{2(\sqrt{N} - r)} \\ &\leq \frac{1 - \frac{1}{\sqrt{2}}}{2(\sqrt{N} - r)} \quad \text{as } \sqrt{N} \geq \frac{1}{\sqrt{2}} \\ &\leq \frac{1 - \frac{1}{\sqrt{2}}}{2(\frac{1}{\sqrt{2}} - r)} \quad \text{as } \sqrt{N} \geq \frac{1}{\sqrt{2}} \\ &= \frac{1 - \frac{1}{\sqrt{2}}}{2(\frac{2}{\sqrt{2}} - 1)} \\ &= \frac{1 - \frac{1}{\sqrt{2}}}{2\sqrt{2}(1 - \frac{1}{\sqrt{2}})} \\ &= \frac{1}{2\sqrt{2}} \\ &< 1 \quad \text{as } 2\sqrt{2} > 1 \end{aligned}$$

As we have confirmed that $M_{\epsilon_0} < 1$, then we have confirmed that NR_3 is satisfied on I , and so the algorithm converges quadratically to the desired root. \square

Using the previous proposition we can, similar to our previous methods, consider how many iterations would be needed to reach a required tolerance. To start we consider that, as mentioned in the proof of Theorem 2.3.2, that $\epsilon_n \leq (M_{\epsilon_0})^{2^{n-1}} \epsilon_0$.

We know that $M_{\epsilon_0} \leq \frac{1}{2\sqrt{2}}$ and that $\epsilon_0 \leq 1 - \frac{1}{\sqrt{2}}$, giving:

$$\epsilon_n \leq \left(\frac{1}{2\sqrt{2}} \right)^{2^{n-1}} \left(1 - \frac{1}{\sqrt{2}} \right)$$

Thus if we want to achieve a tolerance of $\epsilon_n \leq \tau$, then it suffices to find $n \in \mathbb{N}_0$ such that:

$$\left(\frac{1}{2\sqrt{2}} \right)^{2^{n-1}} \leq \tau$$

Then,

$$(2^n - 1) \log \left(\frac{1}{2\sqrt{2}} \right) \leq \log \left(\frac{\tau}{1 - \frac{1}{\sqrt{2}}} \right)$$

By noting that $\log(\frac{1}{a}) = -\log(a)$, then we get

$$(1 - 2^n) \log(2\sqrt{2}) \leq \log\left(\frac{\tau}{1 - \frac{1}{\sqrt{2}}}\right)$$

Once this is rearranged we get the following inequality:

$$2^n \geq \frac{\log\left(\frac{2(\sqrt{2}-1)}{\tau}\right)}{\log(2\sqrt{2})}$$

By taking logarithms again and re-arranging we get that

$$n \geq \frac{\log\left(\frac{\log\left(\frac{2(\sqrt{2}-1)}{\tau}\right)}{\log(2\sqrt{2})}\right)}{\log(2)} = \log_2\left(\log_{2\sqrt{2}}\left(2\frac{\sqrt{2}-1}{\tau}\right)\right)$$

Now for an example, suppose we want to know how many iterations we need to perform to find \sqrt{N} to within 10 decimal places, i.e. $\tau = 10^{-10} = 0.0000000001$. We remember that $\sqrt{N} \in [\frac{1}{2}, 1)$, and then we will apply transformations to this value afterwards, therefore this is equivalent to finding 10 significant digits of accuracy for our square root (ignoring any loss of accuracy that may arise from multiplications afterwards).

Now in this case we want to find $n \in \mathbb{N}$ such that $n \geq \log_2(\log_{2\sqrt{2}}(2 \cdot 10^{10}(\sqrt{2} - 1)))$. Calculating this value we find that we need $n \geq 4.457144\dots$ and so we can take $n = 5$. This means that we could modify our algorithm and implementation to do 5 fixed iterations of Newton's Method to guarantee at least 10 decimal places of accuracy.

In terms of efficiency versus accuracy trade-off, modifying the problem thusly would improve it's efficiency by removing now unnecessary calculation and comparison of δ_n at each stage. However this does need a fixed guaranteed accuracy, and therefore such a program would no longer be suitable if we needed to calculate a square root accurate to 15 decimal places.

Below is a table that lists the minimum $n \in \mathbb{N}$ such that n satisfies our inequality, where our tolerance is 10^k for some $k \in \mathbb{N}$. This will give us the maximum number of iterations that must be performed for the required accuracy.

$k : \tau = 10^k$	n
5	4
10	5
100	8
1,000	12
1,000,000	22

3.4 Newton's Inverse Square Root Method

While the Newton's method discussed in the previous section is acceptable, it has a small issue when it comes to performance, namely that division is slow for a computer to perform compared to multiplication. With this knowledge in mind we would like to find a way of utilising

Newton's method without having to perform any division operations.

If we consider $f(x) = N - \frac{1}{x^2}$ then if x^* is a solution to $f(x) = 0$ we see that $x^* = \frac{1}{\sqrt{N}}$. As $f'(x) = \frac{2}{x^3}$, then the Newton's Method, will give

$$x_{n+1} = x_n - \frac{N - \frac{1}{x_n^2}}{\frac{2}{x_n^3}} = x_n \left(\frac{3}{2} - \frac{N}{2} x_n^2 \right)$$

where x_0 is a given initial guess. As can be seen this algorithm requires no division if we multiply by real constants rather than the division implied above.

We can then consider that, similar to Algorithm 3.3.4, any N can be represented as $a \cdot 2^b$ where $a \in [\frac{1}{2}, 1)$. This will, again allow us to narrow our problem to a known range of values, by using the following transformations.

$$\begin{aligned} N = a \cdot 2^b &\implies \frac{1}{N} = \frac{1}{a} \cdot 2^{-b} \\ &\implies \frac{1}{\sqrt{N}} = \frac{1}{a} \cdot 2^{\lfloor \frac{-b}{2} \rfloor} \cdot \alpha \quad \alpha := \begin{cases} 1 & : b \equiv 0 \pmod{2} \\ \sqrt{2} & : b \equiv 1 \pmod{2}, b \in \mathbb{Z}^- \\ \frac{1}{\sqrt{2}} & : b \equiv 1 \pmod{2}, b \in \mathbb{Z}^+ \end{cases} \\ &\implies \sqrt{N} = N \cdot \frac{1}{\sqrt{a}} \cdot 2^{\lfloor \frac{-b}{2} \rfloor} \cdot \alpha \end{aligned}$$

Therefore we only need to calculate inverse square roots for values of N in the range $[\frac{1}{2}, 1)$. Thus giving us the following algorithm:

Algorithm 3.4.1: Newton Inverse Square Root Method

```

1  NewtonInvSquareRoot( $N \in \mathbb{R}_0^+, \tau \in (0, 1)$ ):
2    Let  $(a, b) \in [\frac{1}{2}, 1) \times \mathbb{Z}$  s.t.  $N = a \cdot 2^b$ 
3     $x_0 := 1$ 
4    if  $b \equiv 0 \pmod{2}$ :
5       $\alpha := 1$ 
6    else:
7      if  $b > 0$ :
8         $\alpha := \frac{1}{\sqrt{2}}$ 
9      else:
10        $\alpha := \sqrt{2}$ 
11     $n := 0$ 
12    loop:
13       $x_{n+1} := x_n(\frac{3}{2} + \frac{a}{2}x_n^2)$ 
14       $\delta_n := |x_{n+1} - x_n|$ 
15      if  $\delta_n \leq \tau$ :
16        return  $N \cdot \alpha \cdot x_{n+1} \cdot 2^{\lfloor \frac{-b}{2} \rfloor}$ 
17       $n \mapsto n + 1$ 
```

With this method we can once again consider it's convergence properties, in particular does it satisfy the criteria for quadratic convergence in Theorem 2.3.2.

Proposition 3.4.1. *Algorithm 3.4.1 satisfies the criteria of Theorem 2.3.2, and thus has quadratic convergence to \sqrt{N} .*

Proof. We know that we only need to consider $N \in [\frac{1}{2}, 1)$, and therefore $\sqrt{N^{-1}} \in (1, \sqrt{2}]$. Also $x_0 = 1$ and so we see that

$$\epsilon_0 = \left| x_0 - \sqrt{N^{-1}} \right| = \sqrt{N^{-1}} - x_0 \leq \sqrt{2} - 1$$

Now let $r := \epsilon_0 = \sqrt{N} - 1$ and $I := [\sqrt{N^{-1}} - r, \sqrt{N^{-1}}]$. If we consider the lower bound of I we see that $\sqrt{N^{-1}} - (\sqrt{N^{-1}} - 1) = 1$, and in particular $0 \notin I$.

Next we know that $f(x) = N - x^{-2}$, and therefore we get $f'(x) = 2x^{-3}$, $f''(x) = -6x^{-4}$. It is obvious that $\nexists x \in \mathbb{R} : f'(x) = 0$, which means that $f'(x) \neq 0 \forall x \in I$ and so NR_1 is satisfied. Also as f'' is only discontinuous at $x = 0$ and $0 \notin I$, then $f''(x)$ is continuous $\forall x \in I$, meaning this satisfies NR_2 .

Now $M = \sup \left\{ \frac{1}{2} \left| \frac{2x^3}{6y^4} \right| : x, y \in I \right\}$, we can simplify the function we are trying to minimise to get $\frac{1}{6} \frac{x^3}{y^4}$. It is obvious that in order to maximise this function we should find the largest possible x and smallest possible y , as both are positive. Hence by taking $x = \sqrt{N^{-1}} + r$ and $y = 1$, then $M = \frac{1}{6}(2\sqrt{N^{-1}} - 1)^3 \leq \frac{1}{6}(2\sqrt{2} - 1)^3$.

Now we consider M_{ϵ_0} :

$$\begin{aligned} M_{\epsilon_0} &= \frac{1}{6}(2\sqrt{N^{-1}} - 1)^3(\sqrt{N} - 1) \\ &\leq \frac{1}{6}(2\sqrt{2} - 1)^3(\sqrt{2} - 1) \\ &\approx 0.42199376 \dots \\ &< 1 \end{aligned}$$

Therefore as $M_{\epsilon_0} < 1$ we have satisfied NR_3 , and as such we have quadratic convergence of our method to $\sqrt{N^{-1}}$. \square

3.5 Comparison of Methods

We have observed several methods that can be used to calculate Square Roots, and so now we will see how the methods compare to each other in practice. The exact root method that we first discussed is the hardest to compare to the other methods as it works in a very different manner. For now we will merely observe that it is an inefficient method that will be shown to take longer than the others.

Second we need to compare the different methods discussed for the Newton Square Root method. As the methods discussed work by the same mechanism of successive approximations, and have similar complexity for each iteration; then we will compare the efficiency of these methods by their computation time. To do this we will be testing 1000 values in the range $(0, 1000)$ and will calculate each of these values 100000 times, accurate to within a tolerance of 10^{-1} , for each method to give the most accurate results. The table below gives the calculated results:

	Total time:	Average time:	Minimum time:	Maximum time:
mpfr_newton_sqrt_v1	10.507s	0.010s	0.003s	0.016s
mpfr_newton_sqrt_v2	12.707s	0.012s	0.004s	0.021s
mpfr_newton_sqrt_v3	8.188s	0.008s	0.005s	0.016s

Here we see that our third method, as expected, is the fastest of the proposed methods and so we will use this method going forwards. One unexpected result is that the second method is actually slower than the first, which is likely due to the extra conversions, comparisons and method calls; this slows down the execution more than it is sped up by reduction in number of iterations required.

Now for the comparison of methods we will be comparing modified versions of Algorithms 3.2.4, 3.3.4 and 3.4.1, which will execute for a given number of steps, rather than testing for the approximate error. To do this we need to consider how many iterations each method needs to reach a particular number of decimal places of accuracy.

We have seen the required number of iterations for a tolerance $\tau = 10^{-k} : k \in \mathbb{N}$, for both the bisection and basic newton square root methods, and similar to the basic newton method, we can show that for the inverse newton method we are looking for $n \in \mathbb{N}$ that satisfies the following inequality:

$$n > \log_2 \left(\log_{\frac{1}{6}(\sqrt{2}-1)(2\sqrt{2}-1)^3} \left(\frac{\tau}{\sqrt{2}-1} \right) \right) - 1$$

This gives the following table:

$k : \tau = 10^k$	Bisection Method	Newton Method	Inverse Newton
5	16	4	4
10	33	5	5
100	332	8	9
1,000	3321	12	12
1,000,000	3219280	22	22

To show the above in action we have the table below which shows the convergence of all 3 methods to $\sqrt{0.75} \approx 0.86602540378$, for different numbers of iterations n with the bold digits being those correct:

n	bisectSquareRoot	NewtonSquareRoot	NewtonInvSquareRoot
0	0 .5000000000000000	1 .0000000000000000	0 .7500000000000000
1	0 .7500000000000000	0 .8750000000000000	0 .8437500000000000
2	0 .8750000000000000	0 .866071428571428	0 .865173339843750
3	0 .8125000000000000	0 .866025405007363	0 .866024146705512
4	0 .8437500000000000	0 .866025403784438	0 .866025403781701
5	0 .8593750000000000	0 .866025403784438	0 .866025403784438
6	0 .8671875000000000	0 .866025403784438	0 .866025403784438

If we compare the methods so that they guarantee an accuracy of 10 decimal places, then we will be able to see their relative efficiency. In particular we will again be testing the three methods using 1000 values in the range $(0, 1000)$, and calculating the square root of each of these values 10000 times for each method; further we will be including the digit by digit

method and the built-in C `sqrt` function. The results calculated are present in the following table:

	Total time:	Average time:	Minimum time:	Maximum time:
<code>root_digits_precise</code>	227.620s	0.227s	0.160s	0.429s
<code>bisect_sqrt</code>	2.520s	0.002s	0.002s	0.004s
<code>newton_sqrt</code>	1.028s	0.001s	0.000s	0.004s
<code>newton_inv_sqrt</code>	0.646s	0.000s	0.000s	0.001s
<code>builtin_sqrt</code>	0.072s	0.000s	0.000s	0.000s

Here we see the expected result that the digit by digit method is the least efficient method, taking two orders of magnitude more time than the second least efficient. We also see that while the two different newton methods are similar in time, and that even though they each performed the same number of iterations, the inverse square root method is the faster; this is due to the method having no division operations to perform. The quickest is of course the built-in `sqrt` function from C, this is due to an implementation that uses several low-level features of the C language to achieve the displayed level of performance.

In conclusion we can say that the best method that we have considered is Algorithm 3.4.1 which has rapid convergence to the sought square root, while also having fast execution. However if we are in a situation where we require large numbers of digits of accuracy, and yet do not have a suitable floating point types large enough to store these values, then the digit by digit method can be used to get an arbitrary number of digits of accuracy.

4 Trigonometric Functions

The trigonometric functions have been studied since antiquity, originally for their relation to triangles, which were incredibly important to early mathematical understanding. Presently the trigonometric functions have found applications in a vast array of problems from musical theory to satellite navigation.

Here we will discuss various methods for approximating the trigonometric functions \sin , \cos and \tan . Further we will explore the inverse trigonometric functions \sin^{-1} , \cos^{-1} and \tan^{-1} which also have many practical uses in modern life.

Trigonometric functions can be calculated using either degrees of an angle (e.g. $\sin(60^\circ) = \frac{\sqrt{3}}{2}$) or in radians (e.g. $\sin(\frac{\pi}{3}) = \frac{\sqrt{3}}{2}$). For this document we will only discuss the use of radians and consider that $\theta^{\text{rad}} = \theta^\circ \cdot \frac{\pi}{180}$ can be used to convert between the two units if needed.

4.1 Trigonometric Identities

Most readers will be well aware of the standard trigonometric identities: useful equalities that help in the analysis of trigonometric functions; this section will lay out such identities that will prove useful in this document. As with Section 2 this is not meant to be an exhaustive overview, merely a reminder and as such identities not listed here may be used in the document.


The first identities to consider are the basic ones taught in secondary school such as $\sin^2 x + \cos^2 x = 1$. In particular we are interested in the shifts, reflections and periods of \sin and \cos . Some of the relevant functions are included below:

$$\begin{array}{ll} \sin(-x) = -\sin(x) & \cos(-x) = \cos(x) \\ \sin(x + \frac{\pi}{2}) = \cos(x) & \cos(x + \frac{\pi}{2}) = -\sin(x) \\ \sin(x + \pi) = -\sin(x) & \cos(x + \pi) = -\cos(x) \\ \sin(x + 2\pi) = \sin(x) & \cos(x + 2\pi) = \cos(x) \end{array}$$

Another useful formula to remember is the sine rule, which is detailed in figure 4.1.1 as well as the combined angle formulas

$$\begin{aligned} \sin(x) &= \sin(x) \cos(y) \pm \sin(y) \cos(x) \\ \cos(x \pm y) &= \cos(x) \cos(y) \mp \sin(x) \sin(y) \\ \tan(x) &= \frac{\tan(x) \pm \tan(y)}{1 \mp \tan(x) \tan(y)} \end{aligned}$$

Figure 4.1.1: The Sine Rule



A final note in this section is the derivatives of the trigonometric functions, in particular

$$\frac{d}{dx} \sin(x) = \cos(x) \quad \frac{d}{dx} \cos(x) = -\sin(x) \quad \frac{d}{dx} \tan(x) = \sec^2(x)$$


will be useful later on in the development of methods.

4.2 Calculating π

Several of the methods in this section require that we already know the value of π , for example when we are applying several trig identities. Here we will briefly discuss several methods for calculating the value of π , so that we may use this value in later subsections.

The first method to consider is the method used by ancient mathematicians, such as the Greeks and Chinese[14][19, p. 106]. We know that if the radius of the circle is $\frac{1}{2}$, then the circumference of the circle is π , and the value is between the perimeters of the inner and outer polygon perimeters. The internal perimeter is $p_n = n \sin(\frac{\pi}{n})$ and the external perimeter is $P_n = n \tan(\frac{\pi}{n})$.

Figure 4.2.1: Ancient method of calculating π



As we know the values of $\tan(\frac{\pi}{6})$ and $\sin(\frac{\pi}{6})$, then we can calculate P_6 and p_6 . It has been shown that $P_{2n} = \frac{2p_n P_n}{p_n + P_n}$ and $p_{2n} = \sqrt{p_n P_{2n}}$ [9], which allows us to create an iterative method to approximate π , by taking the mid-point of the successive polygon perimeters.

Other common historical methods for approximating π are to use infinite series. One such method uses the series expansion of \tan^{-1} , which is discussed in detail below, where $\tan^{-1}(1) = \frac{\pi}{4}$. This gives the following approximation using N terms:

$$\pi = 4 \sum_{n=0}^N \frac{(-1)^n}{2n+1} = \sum_{n=0}^N \frac{8}{(4n+1)(4n+3)} \quad (4.2.1)$$

This sequence converges very slowly, with sub linear convergence, to the correct value. More modern methods have typically revolved around finding more rapidly converging infinite series, examples include Ramanujan's series[18]:

$$\frac{1}{\pi} = \frac{2\sqrt{2}}{9801} \sum_{n=0}^{\infty} \frac{(4n)!(1103 + 26390n)}{(k!)^n 396^{4n}} \quad (4.2.2)$$

or the Chudnovsky algorithm[3]:

$$\frac{1}{\pi} = 12 \sum_{n=0}^{\infty} \frac{(-1)^n (6n)!(13591409 + 545140134n)}{(3n)!(n!)^3 640320^{3n+\frac{3}{2}}} \quad (4.2.3)$$


This final series is extremely rapidly convergent to the value of $\frac{1}{\pi}$, for example just the first term gives π accurate to 13 decimal places while we can get π accurate to 1000 decimal places with summing just 71 terms. Compared to Equation 4.2.1 which takes the summation of 500 terms to achieve the same 1000 digits of accuracy.

To get large degrees of accuracy for π is extremely computer intensive and using the `mpfr` requires the number of bits of precision and number of terms to be set. This makes calculating π to a large number of decimal places, for example 1000000, computationally infeasible on a regular home computer. Therefore for our purposes we will use the pre calculated value of π to 1000000 decimal places as listed on Exploratorium.edu [6]

4.3 Geometric Method

The first method I will be discussing is a method based on geometric properties that are derived on a circle, and we will start by considering values of \cos in the range $[0, \frac{\pi}{2}]$. To do this we will consider the figure 4.3.1, which shows a unit circle.

Figure 4.3.1: Diagram showing angles to be dealt with



Here theta will be given in radians, and we can note that the labelled arc has length θ due the formula for the circumference of a circle. By using the following derivation we can find a formula for θ in terms of s :

$$\begin{aligned}
 s^2 &= \sin^2 \theta + (1 - \cos \theta)^2 \\
 &= (\sin^2 \theta + \cos^2 \theta) + 1 - 2 \cos \theta \\
 &= 2 - 2 \cos \theta \quad \text{By using } \sin^2 \theta + \cos^2 \theta = 1 \\
 \cos \theta &= 1 - \frac{s^2}{2}
 \end{aligned}$$

We will now consider a figure 4.3.2 which will allow us to calculate an approximate value of s .

We will first note that by an elementary geometry result we can know that the angle ABC is a right-angle; also we can consider that h is an approximation of $\frac{\theta}{2}$, which will become relevant later. Now because AC is a diameter of our circle then it's length is 2 and thus, by utilising Pythagoras' Theorem, we get that the length of AB is $\sqrt{AC^2 - BC^2} = \sqrt{4 - h^2}$.

From here we consider the area of triangle ABC , which can be calculated as $\frac{1}{2} \cdot h \cdot \sqrt{4 - h^2}$ and as $\frac{1}{2} \cdot 2 \cdot \frac{s}{2}$; by equating these two, squaring both sides and re-arranging we get that $s^2 = h^2(4 - h^2)$. We now have the basis for a method that will allow us to calculate $\cos \theta$.

To complete our method we will introduce a new line that is to h what h is to s as shown in figure 4.3.3.

We then see that if we repeat the steps above we get that $h^2 = \hat{h}^2(4 - \hat{h}^2)$, and it also follows that $\hat{h} \approx \frac{\theta}{4}$. Using this we can take an initial guess of $h_0 := \frac{\theta}{2^k}$, for some $k \in \mathbb{N}$, and then

Figure 4.3.2: Diagram detailing how to calculate s





Figure 4.3.3: Detailing the recursive steps



calculate $h_{n+1}^2 = h_n^2(4 - h_n^2)$ where $n \in [0, k] \cap \mathbb{Z}$; finally we calculate $\cos \theta = 1 - \frac{h_k^2}{2}$, giving the following algorithm:

Algorithm 4.3.1: Geometric calculation of \cos

```

1  geometric_cos ( $\theta \in [0, \frac{\pi}{2}], k \in \mathbb{N}$ )
2       $h_0 := \frac{\theta}{2^k}$ 
3       $n := 0$ 
4      while  $n < K$ :
5           $h_{n+1}^2 := h_n^2 \cdot (4 - h_n^2)$ 
6           $n \mapsto n + 1$ 
7      return  $1 - \frac{h_k^2}{2}$ 

```


Now we can use the above pseudo-code to calculate any trigonometric function value by using various trigonometric identities. First we suppose $\theta \in \mathbb{R}$, then we can repeatedly apply the identity $\cos \theta = \cos(\theta \pm 2\pi)$ to either add or subtract 2π until we have a value $\theta' \in [0, 2\pi)$. Once we have this value we can utilise the following assignment to calculate $\cos \theta$:

$$\cos \theta = \begin{cases} \cos \theta' & : \theta' \in [0, \frac{\pi}{2}] \\ -\cos(\pi - \theta') & : \theta' \in [\frac{\pi}{2}, \pi] \\ -\cos(\theta' - \pi) & : \theta' \in [\pi, \frac{3\pi}{2}] \\ \cos(2\pi - \theta') & : \theta' \in [\frac{3\pi}{2}, 2\pi] \end{cases}$$

Using Algorithm 4.3.1 we can also easily calculate both $\sin \theta$ and $\tan \theta$, by further use of trigonometric identities. In particular we note that $\sin \theta = \cos(\theta - \frac{\pi}{2})$ and $\tan \theta = \frac{\sin \theta}{\cos \theta}$. Hence we can now calculate the trigonometric function value of any angle.

We now wish to analyse the error of our approximation for \cos , as the other methods have errors that are derivative of the error for approximating \cos . Now Figure 4.3.4 shows an arc of a circle which creates chord x , with this we will be able to calculate the exact length of the chord and thus work on the error of our approximations.

Figure 4.3.4: Diagram to find actual arc approximation



To start we will note that $\phi = \frac{\pi - \theta}{2} = \frac{\pi}{2} - \frac{\theta}{2}$, and then by using the Sine Rule we get

$$\frac{x}{\sin \theta} = \frac{1}{\sin \phi} \implies x = \frac{\sin \theta}{\sin \phi}$$

Now we can recall the trigonometric identities for \sin , which gives $\sin \theta = 2 \sin \frac{\theta}{2} \cos \frac{\theta}{2}$, and also $\sin \phi = \cos \frac{\theta}{2}$. This allows us to see that

$$x = \frac{2 \sin \frac{\theta}{2} \cos \frac{\theta}{2}}{\cos \frac{\theta}{2}} = 2 \sin \frac{\theta}{2}$$

Therefore we see that h_n is approximating the chord length associated with angle $\theta 2^{n-k}$, and thus $\epsilon_n = |h_n - 2 \sin(\theta 2^{n-k-1})|$. Now as $h_0 = \theta 2^{-k} \approx 2 \sin(\theta 2^{-k-1})$ then it follows that $\exists \phi$ such that $h_0 = 2 \sin(\phi 2^{-k-1})$, from this we can see that $\phi = 2^{k+1} \sin^{-1}(\theta 2^{-k-1})$. We will use these facts to prove a couple of propositions.

Proposition 4.3.1. $h_n = 2 \sin(\phi 2^{n-k-1}) \forall n \in [0, k] \cap \mathbb{Z}$ where $\phi := 2^{k+1} \sin^{-1}(\theta 2^{-k-1})$.

Proof. Proceed by induction on $n \in [0, k] \cap \mathbb{Z}$.

$$\mathbf{H}(n): h_n = 2 \sin(\phi 2^{n-k-1})$$

$\mathbf{H}(0):$

$$\begin{aligned} 2 \sin(\phi 2^{-k-1}) &= 2 \sin(\sin^{-1}(\phi 2^{-k-1})) \\ &= \theta 2^{-k} \\ &= h_0 \end{aligned} \quad \text{by definition of } h_0$$

$\mathbf{H}(n) \implies \mathbf{H}(n+1):$

$$\begin{aligned} h_{n+1} &= h_n \sqrt{4 - h_n^2} \\ &= 2 \sin(\phi 2^{n-k-1}) \sqrt{4 - 4 \sin^2(\phi 2^{n-k-1})} && \text{by } \mathbf{H}(n) \\ &= 4 \sin(\phi 2^{n-k-1}) \cos(\phi 2^{n-k-1}) \\ &= 2 \sin(\phi 2^{n-k}) && \text{by the use of double angle formulas} \end{aligned}$$

□

Proposition 4.3.2. $h_n > 2 \sin(\theta 2^{n-k-1}) \forall n \in [0, k] \cap \mathbb{Z}$

Proof. We start by considering the expansion of the exact value of h_n .

Now as we know that $n \leq k$, then it follows that $\theta 2^{n-k-1} \leq \frac{1}{2}\theta$.

Also as $\theta \leq \frac{\pi}{2}$ we know that $\theta 2^{n-k-1} \leq \frac{\pi}{4}$.

We can also show that $\frac{1}{6}\theta^3 2^{n-3k-3} + \mathcal{O}(2^{-5k}) \leq \frac{\pi}{4}$, though the proof is omitted here for brevity; therefore we see that $\phi 2^{n-k-1} \leq \frac{\pi}{2}$, and obviously that $\phi 2^{n-k-1} > \theta 2^{n-k-1}$.

Hence, as \sin is an increasing function in the range $[0, \frac{\pi}{2}]$, we conclude that

$$h_n = 2 \sin(\phi 2^{n-k-1}) > 2 \sin(\theta 2^{n-k-1})$$

□

With these two propositions we can now consider the error of our approximation of \cos . First we will prove the following proposition regarding the error of the approximation of s :

Proposition 4.3.3. *If $\epsilon_n := |h_n - 2 \sin(\theta 2^{n-k-1})| \forall n \in [0, k] \cap \mathbb{Z}$, then $\epsilon_k < 2^k \epsilon_0$.*

Proof. $\epsilon_n = h_n - 2 \sin(\theta 2^{n-k-1})$ as $h_n > 2 \sin(\theta 2^{n-k-1})$ by Proposition 4.3.2.

Now we see that:

$$\begin{aligned} \epsilon_{n+1} &= h_{n+1} - 2 \sin(\theta 2^{n-k}) \\ &= h_n \sqrt{4 - h_n^2} - 4 \sin(\theta 2^{n-k-1}) \cos(\theta 2^{n-k-1}) \end{aligned}$$

If we consider the equation $\alpha\beta - \gamma\delta = (\alpha - \gamma) + \alpha(\beta - 1) - \gamma(\delta - 1)$ and apply it to our current formula we get:

$$\begin{aligned}
\epsilon_{n+1} &= (h_n - 2\sin(\theta 2^{n-k-1})) + h_n(\sqrt{4 - h_n^2} - 1) - 2\sin(\theta 2^{n-k-1})(2\cos(\theta 2^{n-k-1}) - 1) \\
&= \epsilon_n + h_n(\sqrt{4 - h_n^2} - 1) - 2\sin(\theta 2^{n-k-1})(2\cos(\theta 2^{n-k-1}) - 1) \\
&= 2\epsilon_n + h_n(\sqrt{4 - h_n^2} - 2) - 2\sin(\theta 2^{n-k-1})(2\cos(\theta 2^{n-k-1}) - 2) \\
&= 2\epsilon_n + h_n(\sqrt{4 - h_n^2} - 2) + 2\sin(\theta 2^{n-k-1})(2 - 2\cos(\theta 2^{n-k-1})) \\
&< 2\epsilon_n + h_n(\sqrt{4 - h_n^2} - 2\cos(\theta 2^{n-k-1})) \\
&< 2\epsilon_n + h_n\left(\sqrt{4 - 4\sin^2(\theta 2^{n-k-1})} - 2\cos(\theta 2^{n-k-1})\right) \\
&= 2\epsilon_n + h_n(2\cos(\theta 2^{n-k-1}) - 2\cos(\theta 2^{n-k-1})) \\
&= 2\epsilon_n
\end{aligned}$$

The inequalities in the above derivation arise from the fact that $h_n > 2\sin(\theta 2^{n-k-1})$ by Proposition 4.3.2.

Hence as we now know that $\epsilon_{n+1} < 2\epsilon_n$, we then see that $\epsilon_n < 2^n \epsilon_0$. Therefore we prove our statement that

$$\epsilon_k < 2^k \epsilon_0$$

□

Obviously $\epsilon_k = |h_k - s|$, and we can now use this to find the error of our final answer. First we will start by letting $\mathcal{C} := 1 - \frac{1}{2}h_k^2$ and note that analytically $\cos\theta = 1 - \frac{1}{2}s^2$. Therefore we will now consider $\epsilon_{\mathcal{C}} = |\mathcal{C} - \cos(\theta)|$:

$$\begin{aligned}
\epsilon_{\mathcal{C}} &= \left| 1 - \frac{h_k^2}{2} - 1 + \frac{s^2}{2} \right| \\
&= \frac{1}{2} |h_k^2 - s^2| \\
&= \frac{1}{2} |h_k h_k - 2\sin(\frac{\theta}{2})2\sin(\frac{\theta}{2})| \\
&= \frac{1}{2} (h_k h_k - 2\sin(\frac{\theta}{2})2\sin(\frac{\theta}{2})) \quad \text{as } 2\sin(\frac{\theta}{2}) < h_k \\
&= \frac{1}{2} (2\epsilon_k + h_k(h_k - 2) - 2\sin(\frac{\theta}{2})(2\sin(\frac{\theta}{2}) - 2)) \\
&< \frac{1}{2} (2\epsilon_k + h_k(h_k - 2\sin(\frac{\theta}{2}))) \\
&= \frac{1}{2} (2 + h_k)\epsilon_k \\
&= \frac{1}{2} (2 + 2\sin(\frac{\phi}{2}))\epsilon_k \\
&= (1 + \sin(\frac{\phi}{2}))\epsilon_k \\
&\leq 2\epsilon_k
\end{aligned}$$

As $\epsilon_{\mathcal{C}} \leq 2\epsilon_k$, then by Proposition 4.3.3 we see that $\epsilon_{\mathcal{C}} < 2^{k+1}\epsilon_0$. Now to consider ϵ_0 we first observe that $\epsilon_0 = \theta 2^{-k} - 2\sin\theta 2^{-k-1}$, and therefore we can conclude that:

$$\epsilon_{\mathcal{C}} < 2\theta - 2^{k+2}\sin(\theta 2^{-k-1})$$

It is not immediately obvious that $2\theta - 2^{k+2}\sin(\theta 2^{-k-1})$ is a useful upper bound for $\epsilon_{\mathcal{C}}$. However if we consider the series expansion of $\sin(x)$, shown in Section 4.4 to be $\sin(x) = x - \frac{1}{3}x^3 + \frac{1}{5}x^5 - \dots$, and substitute this into our equation we see that:

$$\begin{aligned}
\epsilon_C &< 2\theta - 2^{k+2}(\theta 2^{-k-1} - \frac{1}{3!}\theta^3 2^{-3k-3} + \frac{1}{5!}\theta^5 2^{-5k-5} - \dots) \\
&= 2\theta - 2\theta + \frac{1}{3}\theta^3 2^{-2k-1} - \frac{1}{5!}\theta^5 2^{-4k-3} + \dots \\
&= \frac{1}{3}\theta^3 2^{-2k-1} - \frac{1}{5!}\theta^5 2^{-4k-1} + \dots
\end{aligned}$$

Now obviously the last line tends towards zero as k tends to infinity, due to it being a formula of order $\mathcal{O}(2^{-2k-1})$. Therefore we know that $\forall \tau \in \mathbb{R}^+ \exists \mathcal{K} \in \mathbb{N} : \epsilon_{C,k} < \tau \forall k \in [\mathcal{K}, \infty) \cap \mathbb{Z}$. In particular, if we then wish to calculate $\cos \theta$ accurate to N decimal places then we are looking to find $k \in \mathbb{N}$ such that:

$$2\theta - 2^{k+2} \sin(\theta 2^{-k-1}) < 10^{-N} \implies 2^{k+2} \sin(\theta 2^{-k-1}) > 2\theta - 10^{-N}$$

For an example of the above in action we will be taking $\theta = 0.5$. The table below shows the minimum $k \in \mathbb{N}$ to guarantee N digits of accuracy in the result:

N	k
5	6
10	14
50	80
100	163
1000	1658

As can be seen the value of k required to achieve N digits of accuracy increases roughly linearly when $\theta = 0.5$. Testing for other values of θ reveals them to have similar required values for k , at least within the same order of each other.

Another consideration for Algorithm 4.3.1 is that we could "run it in reverse" to attain an algorithm for the inverse cosine function. To start take line 7 which is $C = 1 - \frac{1}{2}h_k^2$, which can be re-arranged to give $h_k^2 = 2 - 2C$, where we know C as our initial value.

Line 5 is a little more difficult, but by re-arranging we see that $h_n^4 - 4h_n^2 + h_{n+1}^2 = 0$, which can be solved via the quadratic formula to give $h_n^2 = 2 \pm \sqrt{4 - h_{n+1}^2}$. Now we can make the observation that if $x \in \mathbb{R}_0^+$, then $\cos^{-1}(-x) = \pi - \cos^{-1}(x)$ and so we can restrict our algorithm to only consider $x \in [0, 1]$. With this we know that $\theta \in [0, \frac{\pi}{2}]$, and thus $h_k \leq \sqrt{2}$. Therefore as $h_{n+1} > h_n \forall n \in [0, k-1] \cap \mathbb{Z}$ we see that $h_n^2 \leq 2 \forall n \in [0, k] \cap \mathbb{Z}$. This allows us to ascertain that to reverse Line 5 we perform $h_n^2 = 2 - \sqrt{4 - h_{n+1}^2}$.

Finally line 2 is reversed by returning the value $2^k h_0$; therefore we get the following algorithm for $\cos^{-1}(x)$ where $x \in [0, 1]$:

Algorithm 4.3.2: Geometric calculation of \cos^{-1}

```

1  geometric_aCos( $x \in [0, 1], k \in \mathbb{N}$ )
2       $h_k := 2 - 2x$ 
3       $n := k - 1$ 
4      while  $n \geq 0$ :
5           $h_n^2 := 2 - \sqrt{4 - h_{n+1}^2}$ 
6           $n \mapsto n - 1$ 
7      return  $2^k h_0$ 

```

Similar to the regular trigonometric functions we can use trigonometric identities to calculate the inverse trigonometric functions from \cos^{-1} . To start we recall that $\cos^{-1}(-x) = -\cos^{-1}(x)$ where $x \in [0, 1]$, then we can use the identities that $\sin^{-1}(x) = \frac{\pi}{2} - \cos^{-1}(x)$ and $\tan^{-1}(x) = \sin^{-1}\left(\frac{x}{\sqrt{x^2+1}}\right)$.

If we suppose that all operations in the method are accurately computed then Algorithm 4.3.2 is a computation with high accuracy. This is because there is no initial guess, such as in Algorithm 4.3.1, and so the only introduction of error is assuming that $2^k h_0 \approx \theta$. However as we discuss in detail in Section 3, calculating square roots is not a simple task and thus will introduce error to the method in general; therefore the accuracy of the method is roughly as accurate as our method of calculating square roots.

4.4 Taylor Series

If we consider our definition of a Maclaurin Series from Section 2.1.2, we can use this to approximate our Trigonometric Functions. Consider first $\cos \theta$, for which we know that $\frac{d}{d\theta} \cos \theta = -\sin \theta$; it then follows that $\frac{d^2}{d\theta^2} \cos \theta = -\cos \theta$, $\frac{d^3}{d\theta^3} \cos \theta = \sin \theta$ and $\frac{d^4}{d\theta^4} \cos \theta = \cos \theta$.

If we let $f(x) = \cos x$ and use the known values $\cos(0) = 1$ and $\sin(0) = 0$, then we see that:

$$f^{(n)}(0) = \begin{cases} 1 & : 4 \mid n \\ 0 & : 4 \mid n - 1 \\ -1 & : 4 \mid n - 2 \\ 0 & : 4 \mid n - 3 \end{cases}$$

By simplifying this by omitting the 0 coefficient terms we get the following series:

$$\sum_{n=0}^{\infty} \frac{(-1)^n}{(2n)!} x^{2n} = 1 - \frac{x^2}{2!} + \frac{x^4}{4!} - \dots \quad (4.4.1)$$

By using similar working we can get that the series associated with $\sin(x)$:

$$\sum_{n=0}^{\infty} \frac{(-1)^n}{(2n+1)!} x^{2n+1} = x - \frac{x^3}{3!} + \frac{x^5}{5!} - \dots \quad (4.4.2)$$

Before we go any further we need to consider when Equations 4.4.1 and 4.4.2 converge to their respective functions. Using the ratio test for series [a] and equation 4.4.1 we see that

$$\begin{aligned} L_C &= \lim_{n \rightarrow \infty} \left| \frac{a_{n+1}}{a_n} \right| \\ &= \lim_{n \rightarrow \infty} \left| \frac{\frac{(-1)^{n+1}}{(2n+2)!} x^{2n+2}}{\frac{(-1)^n}{(2n)!} x^{2n}} \right| \\ &= \frac{(2n)!}{(2n+2)!} |x|^2 \\ &= \frac{1}{(2n+2)(2n+1)} |x|^2 \end{aligned}$$

Now it is easy to see that, $L_C = 0$ for all values of x as the fractional component decreases as n increases and $|x|^2$ is a constant. Therefore we can conclude that Equation 4.4.1 converges

to $\cos(x)$ for all values of x . We can use a very similar deduction to show that Equation 4.4.2 converges to $\sin(x)$ for all values of x .

The above means that \cos and \sin can be approximated using Taylor Polynomials, in particular for a given $N \in \mathbb{N}$:

$$\cos x \approx \sum_{n=0}^N \frac{(-1)^n}{(2n)!} x^{2n} \quad \text{and} \quad \sin x \approx \sum_{n=0}^N \frac{(-1)^n}{(2n+1)!} x^{2n+1}$$

This allows us to create the following two methods for computing $\cos x$ and $\sin x$:

Algorithm 4.4.1: Taylor computation of \cos and \sin

```

1  taylor_cos( $x \in \mathbb{R}, N \in \mathbb{N}$ )
2       $\mathcal{C} := 0$ 
3       $n := 0$ 
4      while  $n < N$ :
5           $\mathcal{C} \mapsto \mathcal{C} + (-1)^n \cdot \frac{1}{(2n)!} x^{2n}$ 
6           $n \mapsto n + 1$ 
7      return  $\mathcal{C}$ 
8
9  taylor_sin( $x \in \mathbb{R}, N \in \mathbb{N}$ )
10      $\mathcal{S} := 0$ 
11      $n := 0$ 
12     while  $n < N$ :
13          $\mathcal{S} \mapsto \mathcal{S} + (-1)^n \cdot \frac{1}{(2n+1)!} x^{2n+1}$ 
14          $n \mapsto n + 1$ 
15     return  $\mathcal{S}$ 

```

As these two methods are obviously very similar and the fact that $\sin(x) = \cos(x - \frac{\pi}{2})$, we will continue by examining only the Taylor method for approximating \cos . We will assume that any calculations for \sin are transformed into a problem of finding a \cos value.

It should be noted that this \cos algorithm is particularly inefficient to calculate on a computer implementation; this is primarily due to the way in which the update of \mathcal{C} is calculated each loop.

In each loop we are calculating x^{2n} , which has a naive complexity of $\mathcal{O}(2n)$. However what we are actually calculating is $x^{2(n-1)} \cdot x^2$ and thus if we store the values of $x^{2(n-1)}$ and x^2 , the complexity of this step drops to $\mathcal{O}(1)$. Similarly we are also calculating $\frac{1}{(2n)!}$ in each loop which, by the same logic, is $\frac{1}{2(n-1)!} \cdot \frac{1}{(2n)(2n-1)}$, and we can use the same storage and update method as for x^{2n} .

As another step towards optimizing the algorithm we can start with an initial value of $\mathcal{C} = 1$, and then perform two updates of \mathcal{C} each loop until we reach or surpass N . This saves calculating $(-1)^n$ each loop, by explicitly performing two different calculations. Implementing all of the above gives us the following two updated methods:

Algorithm 4.4.2: Taylor computation of \cos optimised

```

1  taylor_cos( $x \in \mathbb{R}, N \in \mathbb{N}$ )
2       $\mathcal{C} := 1$ 

```

```

3       $x_2 := x^2$ 
4       $a := 1$ 
5       $b := 1$ 
6       $n := 1$ 
7      while  $n < N$ :
8           $a \mapsto a \cdot \frac{1}{(2n-1)(2n)}$ 
9           $b \mapsto b \cdot x_2$ 
10          $C \mapsto C - a \cdot b$ 
11          $a \mapsto a \cdot \frac{1}{(2n+1)(2n+2)}$ 
12          $b \mapsto b \cdot x_2$ 
13          $C \mapsto C + a \cdot b$ 
14          $n \mapsto n + 2$ 
15     return  $C$ 

```

As the next term of the polynomial is known definitively then we can see that it is very easy to calculate the error of our approximation. We see that

$$\begin{aligned}
 \epsilon_N &= |\cos(x) - \text{taylor_cos}(x, N)| \\
 &= \mathcal{O}(|x|^{N'+1}) \quad \text{where } N' \text{ is the smallest} \\
 &\quad \text{odd integer such that } N' \geq N \\
 &\leq \frac{1}{(2(N'+1))!} |x|^{N'+1} \\
 &\leq \frac{1}{(2(N+1))!} |x|^{N+1}
 \end{aligned}$$

If we place bounds on the value of \cos calculated as in Section 4.3, then we know that $|x| \leq \frac{\pi}{2}$, and thus we get the following bound for the error of our approximation:

$$\epsilon_N \leq \frac{\pi^{N'+1}}{2^{N'+1}(2(N'+1))!}$$

Thus if we find $N \in \mathbb{N}$ such that $\frac{\pi^{N+1}}{2^{N+1}(2(N+1))!} < \tau \in \mathbb{R}^+$ then we know that $\epsilon_N < \tau$. If we consider $\tau = 10^k$, then we can find $N \in \mathbb{N}$ such that our approximation is accurate to k decimal places. Below is a table which details some values of k and the corresponding minimum N to guarantee k decimal places of accuracy:

k	N
5	4
10	7
50	21
100	36
1000	233

Now for $\tan x$ we can either calculate both $\sin x$ and $\cos x$ using $\text{taylor_cos}(x, N)$ and divide the resulting value, or we can calculate $\tan x$ directly using a Taylor expansion.

In calculating the Maclaurin series for $\tan x$ we start by letting $\tan x = \sum_{n=0}^{\infty} a_n x^n$, and then noting that as $\tan x$ is an odd series then it's Maclaurin series only contains non-zero coefficients for odd powers of x [21]; therefore we get that $\tan x = \sum_{n=0}^{\infty} a_{2n+1} x^{2n+1} =$

$$a_1x + a_3x^3 + a_5x^5 + \dots$$

Next we consider that $\frac{d}{dx} \tan x = 1 + \tan^2 x$, and knowing the Maclaurin series form of $\tan x$ we get the following:

$$\begin{aligned} \sum_{n=0}^{\infty} (2n+1)a_{2n+1}x^{2n} &= 1 + \left(\sum_{n=0}^{\infty} a_{2n+1}x^{2n+1} \right)^2 \\ &= 1 + a_1^2x^2 + (2a_1a_3)x^4 + (2a_1a_5 + a_3^2)x^6 + \dots \end{aligned}$$

Considering the coefficients of powers on the right hand side of the above equation we see that $2a_1a_3 = a_1a_3 + a_3a_1 = a_1a_{4-1} + a_3a_{4-3}$ and $2a_1a_5 + a_3^2 = a_1a_5 + a_3a_3 + a_5a_1 = a_1a_{6-1} + a_3a_{6-3} + a_5a_{6-5}$. This indicates that our general form for the co-efficient of $2n$ on the right hand side is $\sum_{k=1}^n a_{2k-1}a_{2n-2k+1}$, and thus returning to our equation we get

$$a_1 + \sum_{n=1}^{\infty} (2n+1)a_{2n+1}x^{2n} = 1 + \sum_{n=1}^{\infty} \left(\sum_{k=1}^n a_{2k-1}a_{2n-2k+1} \right) x^{2n}$$

Using this we conclude that $a_1 = 1$ and $a_{2n+1} = \frac{1}{2n+1} \sum_{k=1}^n a_{2k-1}a_{2n-2k+1} \forall n \in \mathbb{N}$. We can note immediately that the calculation of any previous coefficients will provide no help in calculating later coefficients and so the entire sum must be calculated each loop, while also storing each co-efficient already calculated.

This means that the complexity to calculate coefficient a_{2n+1} is $\mathcal{O}(n)$ and will be the n^{th} such calculation, making the complexity of calculating n coefficients to be $\mathcal{O}(n^2)$. Comparing this to the `taylor_cos` method we see that to calculate up to n coefficients of both `cos` and `sin` has complexity $\mathcal{O}(n)$. Therefore it is more efficient to calculate `tan` by calculating both `cos` and `sin` using Algorithm 4.4.2, and performing division than directly using Taylor Polynomial approximation.

We would also like to be able to calculate the inverse trigonometric functions using this method, which means we need to find our Maclaurin series of the inverse trigonometric functions. The simplest of these is \tan^{-1} , where we start by recalling that $\frac{d}{dx} \tan^{-1} x = \frac{1}{1+x^2}$ and then by integrating both sides we get:

$$\begin{aligned} \tan^{-1} x &= \int \frac{1}{1+x^2} dx \\ &= \int (1 - (-x^2))^{-1} dx \\ &= \int \sum_{n=0}^{\infty} (-x^2)^n dx && \text{by Equation 2.1.2} \\ &= \int \sum_{n=0}^{\infty} (-1)^n x^{2n} dx \\ &= c + \sum_{n=0}^{\infty} \frac{(-1)^n}{2n+1} x^{2n+1} \end{aligned}$$

As $\tan^{-1}(0) = 0$ then we see that $c = 0$ and thus gives us the following formula for \tan^{-1} :

$$\tan^{-1} x = \sum_{n=0}^{\infty} \frac{(-1)^n}{2n+1} x^{2n+1}$$

Now due to the restrictions from Equation 2.1.2 the above is only valid for $x \in [-1, 1]$, but we know that the domain of \tan^{-1} is $x \in \mathbb{R}$. To fix this we will first recognise that $\tan^{-1}(-x) = -\tan^{-1}(x)$, so we can restrict our problem to $x \in \mathbb{R}_0^+$. Now if we take the double angle formula for \tan :

$$\tan(\alpha + \beta) = \frac{\tan(\alpha) + \tan(\beta)}{1 - \tan(\alpha)\tan(\beta)}$$

By substituting $\alpha = \tan^{-1}(x)$ and $\beta = \tan^{-1}(y)$ into the above then we get

$$\tan^{-1}(x) + \tan^{-1}(y) = \tan^{-1}\left(\frac{x+y}{1-xy}\right)$$

Using this, suppose we are looking for $\tan^{-1}(z)$ where $z \in (1, \infty)$ and let $y = 1$, then $\tan^{-1}(y) = \frac{\pi}{4}$. We can then re-arrange the equation $z = \frac{x+1}{1-x}$ to get $x = \frac{z-1}{z+1}$; finally as $z > 1$, then $0 < x < 1$. This allows us to calculate:

$$\tan^{-1}(z) = \frac{\pi}{4} + \tan^{-1}\left(\frac{z-1}{z+1}\right)$$

In the above the calculated value is in the range $[0, 1]$ and so it is valid to use a Taylor polynomial using our Maclaurin series above. This gives the following method

Algorithm 4.4.3: Taylor Method for \tan^{-1}

```

1  taylor_aTan ( $x \in [0, 1], N \in \mathbb{N}$ )
2       $\mathcal{T} := 0$ 
3       $x_2 := x^2$ 
4       $y := x$ 
5       $n := 0$ 
6      while  $n < N$ :
7           $\mathcal{T} \mapsto \mathcal{T} + \frac{1}{2n+1}y$ 
8           $y \mapsto y \cdot x_2$ 
9           $\mathcal{T} \mapsto \mathcal{T} - \frac{1}{2n+2}y$ 
10          $y \mapsto y \cdot x_2$ 
11          $n \mapsto n + 2$ 
12     return  $\mathcal{T}$ 

```

Similar to Algorithm 4.4.2 the error of Algorithm 4.4.3 is easy to calculate. We see that

$$\begin{aligned}
 \epsilon_N &= |\tan^{-1}(x) - \text{taylor_aTan}(x, N)| \\
 &\leq \frac{1}{2N+3} |x|^{2N+3} \\
 &\leq \frac{1}{2N+3} \quad \text{as } x \leq 1
 \end{aligned}$$

The next function we will consider is \sin^{-1} , which starts its derivation in much the same way as \tan^{-1} . First we start by recalling that $\frac{d}{dx} \sin^{-1}(x) = (1-x^2)^{-\frac{1}{2}}$, then by taking integrals of both sides we get the following derivation:

$$\begin{aligned}
\sin^{-1}(x) &= \int (1 - x^2)^{-\frac{1}{2}} dx \\
&= \int \sum_{n=0}^{\infty} \binom{-\frac{1}{2}}{n} (-x^2)^n \\
&= c + \sum_{n=0}^{\infty} (-1)^n \left(\prod_{k=1}^n \frac{-\frac{1}{2} - k + 1}{k} \right) \frac{x^{2n+1}}{2n+1} \\
&= c + \sum_{n=0}^{\infty} \frac{(-1)^n}{n!(2n+1)} \left(\prod_{k=1}^n \frac{\frac{1}{2} - k}{k} \right) x^{2n+1} \\
&= c + \sum_{n=0}^{\infty} \frac{(-1)^{2n}}{n!(2n+1)} \left(\prod_{k=1}^n \frac{2k-1}{2} \right) x^{2n+1} \\
&= c + \sum_{n=0}^{\infty} \frac{1}{n!(2n+1)2^n} \left(\prod_{k=1}^n 2k-1 \right) x^{2n+1} \\
&= c + \sum_{n=0}^{\infty} \frac{1}{n!(2n+1)2^n} (1 \times 3 \times 5 \times \cdots \times (2n-1)) x^{2n+1} \\
&= c + \sum_{n=0}^{\infty} \frac{1}{n!(2n+1)2^n} \times \frac{1 \times 2 \times 3 \times \cdots \times (2n)}{2 \times 4 \times \cdots \times (2n)} x^{2n+1} \\
&= c + \sum_{n=0}^{\infty} \frac{(2n)!}{(n!)^2 (2n+1) 4^n} x^{2n+1}
\end{aligned}$$

As $\sin^{-1}(0) = 0$ then we see that $c = 0$. Because the above is valid for $x \in (-1, 1)$, and we know the values of $\sin^{-1}(-1)$ and $\sin^{-1}(1)$, then we can have the following method for evaluating \sin^{-1} :

Algorithm 4.4.4: Taylor Method for \sin^{-1}

```

1  taylor_aSin( $x \in [-1, 1], N \in \mathbb{N}$ )
2      if  $x = 1$ :
3          return  $\frac{\pi}{2}$ 
4      if  $x = -1$ :
5          return  $-\frac{\pi}{2}$ 
6       $\mathcal{S} := x$ 
7       $x_2 := x^2$ 
8       $y := x$ 
9       $a := 1$ 
10      $b := 1$ 
11      $c := 1$ 
12      $n := 1$ 
13     while  $n < N$ :
14          $a \mapsto 2n \cdot (2n-1) \cdot a$ 
15          $b \mapsto n^2 \cdot b$ 
16          $c \mapsto 4 \cdot c$ 
17          $y \mapsto x_2 \cdot y$ 
18          $\mathcal{S} \mapsto \mathcal{S} + \frac{a}{b \cdot c \cdot (2n+1)} \cdot y$ 
19          $n \mapsto n + 1$ 

```

The error for this method is similar to the \tan^{-1} method, in that $\epsilon_N \leq \frac{(2(N+1))!}{((N+1)!(2N+1)4^{N+1})}$. Finally we note that $\cos^{-1}(x) = \frac{\pi}{2} - \sin^{-1}(x)$, and thus can be calculated from a value calculated with Algorithm 4.4.4.

4.5 CORDIC

CORDIC is an algorithm that stands for **C**oordinate **R**otation **D**igital **C**omputer[19, p. 138] and can be used to calculate many functions, including Trigonometric Values. The CORDIC algorithm works by utilising Matrix Rotations of unit vectors. This algorithm is less accurate than some other methods but has the advantage of being able to be implemented for fixed point real numbers in efficient ways using only addition and bit shifting.

CORDIC works by taking an initial value of $\mathbf{x}_0 = \begin{pmatrix} 1 \\ 0 \end{pmatrix}$ which can be rotated through an anti-clockwise angle of γ by the matrix

$$\begin{pmatrix} \cos \gamma & -\sin \gamma \\ \sin \gamma & \cos \gamma \end{pmatrix} = \frac{1}{\sqrt{1 + \tan^2 \gamma}} \begin{pmatrix} 1 & -\tan \gamma \\ \tan \gamma & 1 \end{pmatrix}$$

By taking taking smaller and smaller values of γ we can create an iterative process to find \mathbf{x}_n which converges, for a given $\theta \in (-\frac{\pi}{2}, \frac{\pi}{2})$, to

$$\begin{pmatrix} \cos \theta \\ \sin \theta \end{pmatrix}$$

To do this we repeatedly add and subtract our values for γ from θ to bring it as close to 0 as possible. For our purposes we wish to have a sequence $(\gamma_k : k \in [0, n] \cap \mathbb{Z})$ which will allow us to construct all angles in the range $(-\frac{\pi}{2}, \frac{\pi}{2})$ to within a known level of accuracy.

The way that this works can be thought of like a paper fan where each section is smaller than the last and to approximate the desired angle we repeatedly fold the angle back and forth. An visualisation of this is in figure 4.5.1, which shows three views of the CORDIC fan. The top left view is the unfolded fan, the top right is the fan folded to approximate the angle shown by the red line and the view at the bottom is a close up of the previous view.


While there are many possible choices for γ_k we wish to consider $(\gamma_k : k \in [0, n] \cap \mathbb{Z})$ such that $\tan \gamma_k = 2^{-k} \forall k \in [0, n] \cap \mathbb{Z}$. We can note that the powers of 2 have a useful property, in that if $m > n \in \mathbb{N}$ we see that $\sum_{k=n}^{m-1} 2^k = 2^m - 2^n$. We wish to show that our choice for γ_k have a similar property which will be useful in showing that they are a good choice for our CORDIC algorithm.

Proposition 4.5.1. *If $m \in \mathbb{Z}_0^+$ and $n \in \mathbb{Z}^+$ such that $m > n$ and $\gamma_k = \tan^{-1}(2^{-k}) \forall k \in \mathbb{Z}_0^+$, then $\gamma_m < \gamma_n + \sum_{k=m+1}^n \gamma_k$.*

Proof. We know that $2^{-m} = 2^{-n} + \sum_{k=m+1}^n 2^{-k}$, and thus by applying \tan^{-1} to both sides we get:

$$\tan^{-1} 2^{-m} = \gamma_m = \tan^{-1}(2^{-m-1} + 2^{-m-2} + \dots + 2^{-n} + 2^{-n})$$

Figure 4.5.1: The CORDIC fan



Let $a := 2^{-m-1} + 2^{-m-2} + \dots + 2^{-n} + 2^{-n}$ and $b := 2^{-m-2} + \dots + 2^{-n} + 2^{-n}$. Obviously $a < b$ and further we know that \tan^{-1} is continuous on $[a, b]$ and differentiable on (a, b) . Therefore we can apply the Mean Value Theorem[15] from calculus to find that

$$\exists c \in (a, b) : \frac{1}{c^2 + 1} = \frac{\tan^{-1}(b) - \tan^{-1}(a)}{b - a}$$

By re-arranging we see that

$$\begin{aligned} \tan^{-1}(b) &= \frac{2^{-m-1}}{c^2 + 1} + \tan^{-1}(a) \\ &< \frac{2^{-m-1}}{2^{-2m-2} + 1} + \tan^{-1}(a) \end{aligned}$$

It can be shown, by considering the series expansion of $\tan^{-1}(2^{-m-1})$, that $\frac{2^{-m-1}}{2^{-2m-2} + 1} < \tan^{-1}(2^{-m-1}) \forall m \in \mathbb{Z}_0^+$; therefore we get that:

$$\tan^{-1}(b) < \tan^{-1}(2^{-m-1}) + \tan^{-1}(a)$$

Following this and using the assumed value of γ_{m+1} , we see that:

$$\gamma_m < \gamma_{m+1} + \tan^{-1}(2^{-m-2} + \dots + 2^{-n} + 2^{-n})$$

By repeating the above process we eventually see that:

$$\gamma_m < \sum_{k=m+1}^{n-1} \gamma_k + \tan^{-1}(2^{-n} + 2^{-n})$$

In a similar manner we can repeat the above process with $a := \tan^{-1}(2^{-n})$ and $b := \tan^{-1}(2^{-n} + 2^{-n})$. This will show that:

$$\gamma_m < \gamma_n + \sum_{k=m+1}^n \gamma_k$$

□

Using the previous proposition we can then show that our γ_k have the property that every angle in $(-\frac{\pi}{2}, \frac{\pi}{2})$ can be approximated by either adding or subtracting successive γ_k to within a tolerance of γ_n .

Proposition 4.5.2. *If $\gamma_k = \tan^{-1}(2^{-k}) \forall k \in \mathbb{Z}$, then for any $n \in \mathbb{N}$*

$$\exists (c_k \in \{-1, 1\} : k \in [0, n] \cap \mathbb{Z}) : |\beta - \sum_{k=0}^n c_k \gamma_k| \leq \gamma_n \quad \forall \beta \in (-\frac{\pi}{2}, \frac{\pi}{2})$$

Proof. We let $\beta \in (-\frac{\pi}{2}, \frac{\pi}{2})$ and then will proceed by induction on $n \in \mathbb{N}$.

$$\mathbf{H}(n): \exists (c_k \in \{-1, 1\} : k \in [0, n] \cap \mathbb{Z}) : |\theta - \sum_{k=0}^n c_k \gamma_k| \leq \gamma_n$$

$\mathbf{H}(0)$: We have 4 cases to consider:

Case $\theta \in [0, \frac{\pi}{4})$: In this case $-\frac{\pi}{4} \leq \theta - \gamma_0 < 0$

Therefore $|\theta - \gamma_0| \leq \gamma_0$.

Case $\theta \in [\frac{\pi}{4}, \frac{\pi}{2})$: In this case $0 \leq \theta - \gamma_0 < \frac{\pi}{4}$

Therefore $|\theta - \gamma_0| \leq \gamma_0$.

Case $\theta \in (-\frac{\pi}{4}, 0)$: In this case $0 < \theta + \gamma_0 < \frac{\pi}{4}$

Therefore $|\theta - \gamma_0| < \gamma_0$.

Case $\theta \in (-\frac{\pi}{2}, -\frac{\pi}{4}]$: In this case $-\frac{\pi}{4} < \theta - \gamma_0 \leq 0$

Therefore $|\theta - \gamma_0| < \gamma_0$.

Therefore we see that $\mathbf{H}(0)$ holds true.

$\mathbf{H}(n) \implies \mathbf{H}(n+1)$:

By $\mathbf{H}(n) \exists (c_k \in \{-1, 1\} : k \in [0, n] \cap \mathbb{Z}) : |\theta - \sum_{k=0}^n c_k \gamma_k| \leq \gamma_n$; so let $\theta_n := \theta - \sum_{k=0}^n c_k \gamma_k$.

By Proposition 4.5.1 we know that $\gamma_n < 2\gamma_{n+1}$, and so we can proceed by case analysis:

Case $\theta_n \in [0, \gamma_{n+1})$:

$$-\gamma_{n+1} \leq \theta_n - \gamma_{n+1} < 0 \implies |\theta - \sum_{k=0}^{n+1} c_k \gamma_k| \leq \gamma_{n+1} \text{ where } c_{n+1} = -1.$$

Case $\theta_n \in [\gamma_{n+1}, \gamma_n)$:

$$0 \leq \theta_n - \gamma_{n+1} < \gamma_{n+1} \implies |\theta - \sum_{k=0}^{n+1} c_k \gamma_k| \leq \gamma_{n+1} \text{ where } c_{n+1} = -1.$$

Case $\theta_n \in [-\gamma_{n+1}, 0)$:

$$0 \leq \theta_n + \gamma_{n+1} < \gamma_{n+1} \implies |\theta - \sum_{k=0}^{n+1} c_k \gamma_k| \leq \gamma_{n+1} \text{ where } c_{n+1} = 1.$$

Case $\theta_n \in (-\gamma_n, -\gamma_{n+1})$:

$$-\gamma_{n+1} < \theta_n + \gamma_{n+1} < 0 \implies |\theta - \sum_{k=0}^{n+1} c_k \gamma_k| \leq \gamma_{n+1} \text{ where } c_{n+1} = 1.$$

Therefore as we have found a suitable c_n in all cases then we have shown that $H(n) \implies H(n+1)$. \square

With this proposition we see that our choice for γ_k is a good choice to use for the CORDIC algorithm as it covers the entire range of $(-\frac{\pi}{2}, \frac{\pi}{2})$.

Now, as stated before, the basis of our algorithm is to calculate $\begin{pmatrix} \cos \theta \\ \sin \theta \end{pmatrix}$ by using rotations of a unit vector. By putting our values for γ_k into our rotation matrix we get the following:

$$\begin{pmatrix} \cos \gamma_k & -\sin \gamma_k \\ \sin \gamma_k & \cos \gamma_k \end{pmatrix} = \frac{1}{\sqrt{1+2^{-2k}}} \begin{pmatrix} 1 & -2^{-k} \\ 2^{-k} & 1 \end{pmatrix}$$

Then if we take a current estimate of $\begin{pmatrix} \cos \theta \\ \sin \theta \end{pmatrix}$ at step k to be $\begin{pmatrix} x_k \\ y_k \end{pmatrix}$, we see that

$$\begin{pmatrix} \cos \gamma_k & -\sin \gamma_k \\ \sin \gamma_k & \cos \gamma_k \end{pmatrix} \begin{pmatrix} x_k \\ y_k \end{pmatrix} = \frac{1}{\sqrt{1+2^{-2k}}} \begin{pmatrix} x_k - 2^{-k}y_k \\ y_k + 2^{-k}x_k \end{pmatrix}$$

This gives a very simple formula for the update of x_k and y_k , which can be used as the basis of the CORDIC Algorithm.

As seen in our proof of Proposition 4.5.2, we can approximate our desire angle at step n by keeping a track of $\theta_n := \theta - \sum_{k=0}^{n-1} c_k \gamma_k$. At step n we then have $\theta_{n+1} = \theta_n - \gamma_n$ if $\theta_{n+1} \geq 0$, and $\theta_{n+1} = \theta_n + \gamma_n$ otherwise. This leads us to the general implementation of CORDIC for Trigonometric Functions:

Algorithm 4.5.1: General Cordic

```

1  CORDIC( $\theta \in (-\frac{\pi}{2}, \frac{\pi}{2}), n \in \mathbb{N}$ ):
2       $x := 1$ 
3       $y := 0$ 
4       $k := 0$ 
5      while  $k < n$ :
6          if  $\theta \geq 0$ :
7               $t := x$ 
8               $x \mapsto \frac{1}{\sqrt{1+2^{-2k}}}(x - 2^{-k}y)$ 
9               $y \mapsto \frac{1}{\sqrt{1+2^{-2k}}}(y + 2^{-k}t)$ 
10              $\theta \mapsto \theta - \tan^{-1}(2^{-k})$ 
11          else:
12               $t := x$ 
13               $x \mapsto \frac{1}{\sqrt{1+2^{-2k}}}(x + 2^{-k}y)$ 
14               $y \mapsto \frac{1}{\sqrt{1+2^{-2k}}}(y - 2^{-k}t)$ 
15               $\theta \mapsto \theta + \tan^{-1}(2^{-k})$ 
16           $k \mapsto k + 1$ 
17      return  $(x, y)^T$ 

```

There are few improvements we can make on the general algorithm, however if we start to consider implementations of the algorithm we can find several ways to make our algorithm more efficient.

We first consider the representation of our values in the program, and while the previous algorithms used a floating point `double` value, as described in Section 1.1, we will instead use a fixed point representation for CORDIC. If we have a fixed point representation of our values, then we are using an N bit integer to represent the value in question, with a fixed number of bits set aside for the integer part and the remainder for the fractional part. In this case the processes of addition, subtraction, and multiplication & division by powers of 2 is the same as that for integers.

In particular as our values never exceed the range of $(-2, 2)$, then we can use $N - 2$ bits of our N bit integer to be the fractional part; this gives us a maximum precision of 2^{2-N} . Further, as we are only performing multiplication and division by two, this operation can be performed by bit shifting the values, which is much quicker than actual integer multiplication.

Second we can pre calculate all of the values needed for the algorithm to trade storage space for a reduction in computational complexity. The values which we need to pre-calculate are $\gamma_k = \tan^{-1}(2^{-k})$ and $\frac{1}{\sqrt{1+2^{-2k}}}$ for $k \in [0, n) \cap \mathbb{Z}$. The first thing to note about this is that instead of calculating the multiplication $\frac{1}{\sqrt{1+2^{-2k}}}$ at each stage we can actually take this value out of the loops and pre-calculate $\prod_{k=0}^n \frac{1}{\sqrt{1+2^{-2k}}}$ for $k \in [0, n) \cap \mathbb{Z}$. Using these pre calculated products we can then replace $x := 1$ with $x := \prod_{k=0}^n \frac{1}{\sqrt{1+2^{-2k}}}$ in the initialisation stage.

Now to consider an actual implementation, suppose we are using the 16 bit integer `int16_t` to represent our values; which will have the leading two bits represent the integer part and the remaining 14 bits represent the fractional part. In this case the level of precision is $2^{-14} = 0.00006103515625$ and further we can show that as $\gamma_{14} = \tan^{-1}(2^{-14}) \approx 2^{-14}$; therefore the largest we will choose $n := 14$ to ensure the maximum possible accuracy, without performing excessive calculations

This means we can simplify our algorithm further by calculating only $\prod_{k=0}^{14} \frac{1}{\sqrt{1+2^{-2k}}}$ and $\tan^{-1}(2^{-k}) \forall k \in [0, 14] \cap \mathbb{Z}$. One further note is that these values then need to be converted to approximations in our 16 bit fixed point representation. The first value is:

$$\begin{aligned} \prod_{k=0}^{14} \frac{1}{\sqrt{1+2^{-2k}}} &= 0.60725293651701023412897124207973889082 \dots \\ &\approx 00.10011011011101_2 \\ &= 26dd_{16} \end{aligned}$$

Below is a table of all the angles in the relevant formats

γ_k	Exact Form	Binary	Hexadecimal
γ_0	0.7853981633...	00.11001001000011 ₂	3243 ₁₆
γ_1	0.4636476090...	00.01110110101100 ₂	1dac ₁₆
γ_2	0.2449786631...	00.00111110101101 ₂	0fad ₁₆
γ_3	0.1243549945...	00.00011111110101 ₂	07f5 ₁₆
γ_4	0.0624188099...	00.00001111111110 ₂	03fe ₁₆
γ_5	0.0312398334...	00.00000111111111 ₂	01ff ₁₆
γ_6	0.0156237286...	00.00000100000000 ₂	0100 ₁₆
γ_7	0.0078123410...	00.00000010000000 ₂	0080 ₁₆
γ_8	0.0039062301...	00.00000001000000 ₂	0040 ₁₆
γ_9	0.0019531225...	00.00000000100000 ₂	0020 ₁₆
γ_{10}	0.0009765621...	00.00000000010000 ₂	0010 ₁₆
γ_{11}	0.0004882812...	00.00000000001000 ₂	0008 ₁₆
γ_{12}	0.0002441406...	00.00000000000100 ₂	0004 ₁₆
γ_{13}	0.0001220703...	00.00000000000010 ₂	0002 ₁₆
γ_{14}	0.0000610351...	00.00000000000001 ₂	0001 ₁₆

This allows us to then write the following method in C to calculate both $\cos \beta$ and $\sin \beta$, provided $\beta \in [-\frac{\pi}{2}, \frac{\pi}{2}]$ is given in 16 bit fixed point representation:

```

1 int16_t *cordic_16(int16_t beta)
2 {
3     const int16_t GAMMA = {0x3243, 0x1dac, 0x0fad, 0x07f5, 0x03fe,
4                             0x01ff, 0x0100, 0x0080, 0x0040, 0x0020,
5                             0x0010, 0x0008, 0x0004, 0x0002, 0x0001};
6
7     int16_t x = 0x26dd, y = 0x0000, t, result;
8
9     for(int k = 0; k <= 14; ++k)
10    {
11        t = x;
12        if(beta >= 0)
13        {
14            beta -= GAMMA[k];
15            x = x - (y >> k);
16            y = y + (t >> k);
17        }
18        else
19        {
20            beta += GAMMA[k];
21            x = x + (y >> k);
22            y = y - (t >> k);
23        }
24    }
25
26    //This line is required by C to allow the value to be returned
27    result = malloc(2 * sizeof(int16_t));
28
29    result[0] = x;
30    result[1] = y;
31    return result;
32 }
```

As can easily be seen in the algorithm the number of calculations each iteration is constant, and the number of iterations is fixed at 15. This means that the algorithm is an $\mathcal{O}(1)$ algo-

rithm, and guarantees an answer accurate to 4 decimal places as $2^{-14} < 10^{-4}$. Further as the only calculations are integer addition, subtraction and bit shifting this method executes extremely quickly.

Similar methods exist for other fixed length formats such as using `int32_t` or `int64_t`. To examine in more detail how the method converges we will consider an implementation using `int64_t`, which will be approximating $\cos(0.5)$. The code used is included in the Appendix A.3 and can perform the calculations with $n \leq 63$. Below are some of the functions approximations for different values of n :

n	Output with bold accurate digits
1	0 .70710678118654757273731
2	0 .94868329805051376801827
3	0.8 4366148773210747346951
4	0 .90373783889353875853345
5	0.87 527458786899225984257
6	0.88 995346811933362385360
...	...
19	0.87758 301847694786257392
20	0.877582 10404530012649360
21	0.877582561 26152311971111
22	0.877582 78986933524468128
...	...
53	0.8775825618903727 5873943
54	0.877582561890372 64771712
55	0.8775825618903727 5873943
56	0.8775825618903727 5873943
...	...
63	0.8775825618903727 5873943

This table shows us several interesting features of the algorithm, the first being that while there are points at which a certain number of decimal places are guaranteed; before that point the number of decimal places of accuracy can vary, such as in the first few iterations. As we know that the error after n iterations is at most $\gamma_n = \tan^{-1}(2^{-n})$, then we can guarantee that we have at least d decimal places of accuracy if we use at least $\log_2(\cot(10^{-d}))$ iterations.

Second there are some values of n which have uncharacteristically close approximations of the actual value, such as the case when 21 iterations are used. This arises due to the algorithm finding a good approximation for β , but then successive numbers of iterations move away from this value, thus once more decreasing the number of decimal digits of accuracy.

Finally at the end of the table we see that from 55 iterations onwards, the results do not get any more accurate. It turns out this is due to the program converting the `int64_t` fixed point values into `double` values, which typically have a precision of around 2^{-55} . If we instead modify the program to use a more precise floating point representation we see that the 53-56 section of the table becomes:

n	Output with bold accurate digits
53	0.8775825618903727 3965747
54	0.877582561890372 68653156
55	0.87758256189037271 298609
56	0.8775825618903727 2621336

This is much more in line with what we would expect to see from the known error of the algorithm.

Now another use of CORDIC is to effectively run it in reverse, which will allow us to calculate the Inverse Trigonometric functions. To do this we will start by considering the method for calculating \tan^{-1} , and then use trigonometric identities to calculate both \cos^{-1} and \sin^{-1} .

To accomplish this we will be fixing some initial values for $\sin \theta$ and $\cos \theta$, and then running the CORDIC algorithm to move the approximation of $\sin \theta$ towards zero. In doing this we will effectively run our algorithm in reverse, and if we keep track of the angles that we rotate through we can find \tan^{-1} .

We know that $\tan \theta = \frac{\sin \theta}{\cos \theta}$, which means that if we have a current approximation $\begin{pmatrix} x_k \\ y_k \end{pmatrix}$ then $\frac{y_k}{x_k} \approx \tan \theta$. Using this, if we have an input of $\tan \theta = z$ then we can take our initial values to be $x_0 := \frac{1}{2}$ and $y_0 := \frac{z}{2}$. This has the desired property that $\frac{y_0}{x_0} = z$, and if we have y_n tending to 0 then the angle we approximate in the process will be θ .

If we again consider a 16 bit fixed point implementation for our algorithm we can implement it as follows:

```

1 | int16_t *cordic_atan_16(int16_t z)
2 | {
3 |     const int16_t GAMMA = {0x3243, 0x1dac, 0x0fad, 0x07f5, 0x03fe,
4 |                             0x01ff, 0x0100, 0x0080, 0x0040, 0x0020,
5 |                             0x0010, 0x0008, 0x0004, 0x0002, 0x0001};
6 |
7 |     int16_t x = 0x2000, y = z >> 1, t, theta;
8 |
9 |     for(int k = 0; k <= 14; ++k)
10 |    {
11 |        t = x;
12 |        if(y < 0)
13 |        {
14 |            theta -= GAMMA[k];
15 |            x = x - (y >> k);
16 |            y = y + (t >> k);
17 |        }
18 |        else
19 |        {
20 |            theta += GAMMA[k];
21 |            x = x + (y >> k);
22 |            y = y - (t >> k);
23 |        }
24 |    }
25 |
26 |    return theta;
27 | }
```

Similar to our considerations when dealing with the Taylor method of calculating \tan^{-1} , we need to ensure that the input value is not too large, and so can perform the same transformations to the value to ensure we are always calculating a value in the range $[0, 1)$. Using this we can then use the identities $\sin^{-1}(z) = \tan^{-1}(\frac{z}{\sqrt{1-z^2}})$ and $\cos^{-1}(z) = \tan^{-1}(\frac{\sqrt{1-z^2}}{z})$.

Obviously there are basic exceptional values that need to be checked for, in particular $\cos^{-1}(0) = \frac{\pi}{2}$, and $\sin^{-1}(\pm 1) = \pm \frac{\pi}{2}$. If these values are checked prior to the actual calculation, then we are never dividing by 0 and $z \in [-1, 1] \cap \mathbb{Z}$, thus we have a complete algorithm that calculates the inverse Trigonometric Functions.

This method, like the CORDIC method for the regular Trigonometric Functions, has an approximation that is accurate to within γ_n . Thus for our 16 bit implementation, the output will be accurate to within an error of $2^{-14} = 0.00006103515625$, in particular guaranteeing at least 4 decimal places of accuracy. A final note is that the Inverse Trigonometric Functions, again much like the regular CORDIC algorithm, is an $\mathcal{O}(1)$ algorithm with simple calculations, making the algorithm extremely efficient.

4.6 Comparison of Methods

We have observed three different methods for calculating the Trigonometric Functions, as well as their inverses, so should compare their efficiency and accuracy properties.

First we will compare how quickly each algorithm approaches the correct value for different inputs of n , and using $\theta = 0.5$. The comparison will use `double` values for computation, so that all three methods may be equally compared. The table below compares the convergence of $\cos \theta$, with the bold digits being the correct digits found:

n	geometric_Cos(0.5, n)	taylor_Cos(0.5, n)	CORDIC(0.5, n)
1	0.87 6953125000000	1.0000000000000000	0.707106781186547
2	0.877 426177263259	0.877 604166666666	0.948683298050513
3	0.8775 43526076081	0.877 604166666666	0.843661487732107
4	0.8775 72806699400	0.87758256 2158978	0.903737838893538
5	0.87758 0123327654	0.87758256 2158978	0.875274587868992
6	0.877581 952264380	0.877582561890373	0.889953468119333
7	0.877582 409484792	0.877582561890373	0.882719918613777
8	0.8775825 23789035	0.877582561890372	0.879022003513595
9	0.87758255 2365041	0.877582561890372	0.877152884812089
10	0.87758255 9509040	0.877582561890372	0.878089122532394

This table demonstrates that `taylor_Cos` has the fastest convergence, and also demonstrates the staggered increase in accuracy as each step of the algorithm calculates two updates to $\cos \theta$, and thus the output only gets more accurate every other value of n . The `geometric_Cos` method has the second best convergence, while the CORDIC algorithm lags behind, having inconsistent convergence as measured in correct digits.

Next we will note that all algorithms can guarantee 10 digits of accuracy in a fixed number of steps. In particular we can guarantee 10 digits of accuracy for `geometric_Cos` when $n \geq 16$, `taylor_Cos` when $n \geq 8$ and CORDIC when $n \geq 34$. Using the lower bounds of each of these

values for n we can directly compare the speed of the methods.

To compare the methods we will be testing 1000 random values in the range $[0, \frac{\pi}{2})$ for which we will calculate the cosine of with each method 100000 times. This will then also be compared to the standard C implementation of the `cos` function, available in `math.h`. The results of my personal testing follow, where the given times are for individual values, not individual method execution times:

	Total time:	Average time:	Minimum time:	Maximum time:
<code>geometric_cos</code>	16.029s	0.016s	0.015s	0.022s
<code>taylor_cos</code>	7.937s	0.007s	0.007s	0.013s
<code>cordic_cos</code>	21.471s	0.021s	0.020s	0.030s
<code>builtin_cos</code>	0.243s	0.000s	0.000s	0.000s

These values show that the fastest algorithm that we have discussed is algorithm 4.4.2, while the slowest is the CORDIC algorithm. However all of our algorithms are much less efficient than the built-in `cos` function of C. It turns out this discrepancy is due to inefficient implementation as the `cos` function also uses a Taylor approximation[10], but it is implemented in a much lower level method that optimises the execution of the code.

Next we will compare our methods for the Inverse Trigonometric Functions, starting with how they converge to the correct value, as detailed in the following table:

n	<code>geometric_aCos(0.5, n)</code>	<code>taylor_aCos(0.5, n)</code>	<code>CORDIC(0.5, n)</code>
1	2.351 425307918200	2.270 796326794896	2.356 194490192344
2	2.347 503635391542	2.327 962993461563	1.892 546881191538
3	2.346 521397812842	2.340 568243461563	2.137 525544318402
4	2.346 275724597314	2.344 244774711563	2.261 880538865164
5	2.346 214299177873	2.345 470795757570	2.324 299348861121
6	2.34619 8942378459	2.345 913166442261	2.355 539182291389
7	2.34619 5103149716	2.346 081295659538	2.339 915453670913
8	2.34619 4143336564	2.3461 47594614218	2.347 727794731014
9	2.346193 903386887	2.3461 74467628018	2.343 821564599047
10	2.3461938 43452078	2.3461 85594784405	2.345 774687115525

Here we see for the inverse trigonometric functions the convergence speed has been altered with the Geometric method now having the fastest convergence; the Taylor Method converges much slower and the CORDIC method is more stable. One interesting behaviour that emerges for larger values of n in the `geometric_aCos` is demonstrated in the following table:

n	<code>geometric_aCos(0.5, n)</code>
13	2.34619382 2083380897
14	2.34619381 2716280469
15	2.34619373 7779483257
...	...
22	2.346 097524754926944
23	2.341 202123910687049
24	2.351 023238547698124

This behaviour arises due to the use of `double` to calculate values of very small magnitude, this causes the value to become effectively 0 and thus lead to the inaccuracies seen. If we use a higher precision representation for the calculations we get the following table instead:

n	<code>geometric_aCos(0.5, n)</code>
13	2.346193823 718087586
14	2.3461938234 83759158
15	2.3461938234 25177051
...	...
22	2.3461938234056 50874
23	2.346193823405649 980
24	2.346193823405649 757

With this we see that Algorithm 4.3.2 continues in the same pattern as before and is actually correct. So we may time our functions again to compare their efficiency. To do this we will again use 1000 random values, this time in the range $(-1, 1)$, each of which we will calculate \cos^{-1} using each method 100000 times. We note that the algorithms can guarantee 10 decimal places of accuracy for different values of n , in particular `geometric_aCos` when $n \geq 18$, `taylor_aCos` when $n \geq 30$ and CORDIC when $n \geq 50$.

	Total time:	Average time:	Minimum time:	Maximum time:
<code>geometric_cos</code>	27.273s	0.027s	0.026s	0.033s
<code>taylor_cos</code>	14.358s	0.014s	0.014s	0.018s
<code>cordic_cos</code>	29.142s	0.029s	0.028s	0.032s
<code>builtin_cos</code>	2.143s	0.002s	0.001s	0.005s

Again this table shows that the Taylor method is the quickest of those analysed and the CORDIC method is the slowest, however they also both are much slower than the built in methods. One thing to note is that the inverse trigonometric functions are simply less efficient to calculate, as can be seen in the execution time of the built-in method, which appears to be two orders of magnitude greater than the corresponding trigonometric method.

We conclude that for most implementations the Taylor method is the most appropriate method to use to ensure a high accuracy quickly. However the CORDIC algorithm is of use when more advanced features such as floating point type values, or hardware multipliers are not present; further it is possible to create hardware implementations of the CORDIC algorithm which can even further speed up the calculations.

5 Logarithms and Exponentials

Exponentiation is the operation of calculating x^y where x and y are members of some field, for the purposes of this document we will be considering $x, y \in \mathbb{R}$. This operation is widely used by many different branches of mathematics and industry, for example many real world phenomena can be modelled by exponentials[12]; we therefore need to calculate x^y quickly and efficiently.

The first thing we consider is that x^y when $x \in \mathbb{R}^-$ and $y \in \mathbb{R} \setminus \mathbb{Z}$ is not well-defined on \mathbb{R} , and requires consideration of the function on the complex plane. Due to this we will not be considering negative numbers to non-integer bases; in particular, unless stated otherwise, we

will be assuming that $x \in \mathbb{R}_0^+$.

Now we also know that $x^{-y} = \frac{1}{x^y}$ when $y \in \mathbb{R}$, and as such we will also be restricting this section to the assumption that $y \in \mathbb{R}_0^+$. Further we consider the following facts:

$$x^0 = 1 \quad \forall x \in \mathbb{R}_0^+$$

$$0^y = 0 \quad \forall y \in \mathbb{R}^+$$

If we take out these known trivial cases then we can restrict this section to considering only $(x, y) \in (\mathbb{R}^+)^2$.

Now if we have $y \in \mathbb{R}^+$ then it follows that $\exists(a, b) \in \mathbb{Z}_0^+ \times [0, 1)$ such that $y = a + b$. This allows us to use the identity that $x^{m+n} = x^m x^n$ to consider the following two cases separately:

$$x^a : a \in \mathbb{Z}_0^+ \tag{5.0.1}$$

$$x^b : b \in [0, 1) \tag{5.0.2}$$

5.1 Calculating x^a

As we know that $a \in \mathbb{Z}_0^+$, then we know that $x^a = \underbrace{x \times \cdots \times x}_a$; i.e. the problem is equivalent to finding x multiplied with itself a times. As we are only dealing with $a \in \mathbb{Z}_0^+$, then we will be considering $x \in \mathbb{R}$ as we can calculate exponentials of negative numbers.

The naive way to go about calculating x^a is to simply perform the multiplication of x by itself a times. The algorithm for that can be seen below:

Algorithm 5.1.1: Naive integer exponentiation

```

1  naive_int_exp( $x \in \mathbb{R}, a \in \mathbb{Z}_0^+$ ):
2       $n := 0$ 
3       $z := 1$ 
4      while  $n < a$ :
5           $z \mapsto x \cdot z$ 
6      return  $z$ 

```

This algorithm is very simple and has complexity of $\mathcal{O}(a)$, which is a reasonably low complexity, but still has the chance to grow large as a grows. Instead we can consider a more informed approach, in particular we know that either $2 \mid a$ or $2 \nmid a$, which then gives us the following:

$$x^a = \begin{cases} (x^2)^{\frac{a}{2}} & : 2 \mid a \\ x \cdot (x^2)^{\frac{a-1}{2}} & : 2 \nmid a \end{cases}$$

We can use this fact to build a recursive method of calculating x^a , where we repeatedly call the method from within itself. To ensure the method ends correctly we need to identify a base case for the recursion, i.e. where the process stops and returns the correct value. We can see that eventually the above will reach the point where $a = 0$, in which case we know that $x^0 = 1$; this will be the base case of our recursion.

We want to ensure that the algorithm will terminate, which we can do by seeing that it terminates when $a = 0$ and then considering $a \in \mathbb{Z}^+$. Now if $2 \mid a$ then $\frac{a}{2} \in \mathbb{Z}^+$ and also $\frac{a}{2} < a$, similarly if $2 \nmid a$ then $\frac{a-1}{2} \in \mathbb{Z}_0^+$ because $a \geq 1$ and also $\frac{a-1}{2} < a$. Thus we see that the sequence produced by $a \in \mathbb{Z}^+$ is a strictly decreasing sequence that is bounded below by 0 and thus we must eventually reach 0, meaning the algorithm terminates.

Instead of a recursive algorithm that calls itself, the algorithm below is an iterative version which performs the same function:

Algorithm 5.1.2: Exponentiation by squaring

```

1  exp_by_squaring( $x \in \mathbb{R}, a \in \mathbb{Z}_0^+$ ):
2       $n := a$ 
3       $z := 1$ 
4       $\hat{x} := x$ 
5      while  $n > 0$ :
6          if  $2 \nmid n$ :
7               $z \mapsto \hat{x} \times z$ 
8               $n \mapsto n - 1$ 
9               $\hat{x} \mapsto \hat{x}^2$ 
10              $n \mapsto \frac{n}{2}$ 
11     return  $z$ 

```

This algorithm is much more efficient than Algorithm 5.1.1 due to the number of times the inner loop is executed. The inner loop drives a towards 0 by dividing by 2 each step, this means that as $a = \mathcal{O}(2^{\log_2(a)})$, then this goal is achieved in only $\log_2(a)$ loops. Therefore the complexity of this algorithm is $\mathcal{O}(\log_2(a))$, which is an improvement upon the previous algorithm's complexity of $\mathcal{O}(a)$.

To see this difference in efficiency in action the following table shows the times taken for each method when comparing 1000 different pairs of values $(x, a) \in [0, 10] \times ([0, 100] \cap \mathbb{Z})$. With these values we calculated x^a using both methods 100000 times to get the following results:

	Total time:	Average time:	Minimum time:	Maximum time:
naive_int_exp	16.800s	0.016s	0.000s	0.037s
squaring_int_exp	2.593s	0.002s	0.000s	0.004s

With this we will move on to further subsections as there are few improvements that can be made on an $\mathcal{O}(\log_2(a))$ algorithm, particularly in this instance.

5.2 Calculating x^b

If we have $b \in (0, 1)$, then we obviously can't use the our previous subsection for calculating x^y . The most common way of calculating such exponentiation is by considering that $x = e^{\ln(x)}$ and thus $x^b = (e^{\ln(x)})^b = e^{b \ln(x)}$; however this now raises the problem of how to calculate both e^α and $\ln(\beta)$. The following will deal with how to calculate these values and thus use them in conjunction to calculate x^b .

5.3 Naive Method

The mathematical constant e has been known since the early 1600s and was originally calculated by Jacob Bernoulli, and was studied by Leonhard Euler, where it appeared in Euler's Mechanica in 1736. While several possible equivalent definitions of e exist the most common such definition is that $e := \lim_{n \rightarrow \infty} (1 + \frac{1}{n})^n$.

If we now consider the definition of e and also consider e^x , then we can show that $e^c = \lim_{n \rightarrow \infty} (1 + \frac{x}{n})^n$. This gives us our first basic method of how to calculate e^x :

Algorithm 5.3.1: Basic Method for calculating e^x

```

1  basic_exp( $x \in \mathbb{R}, n \in \mathbb{N}$ )
2  return  $(1 + \frac{x}{n})^n$ 

```

If we consider $(1 + \frac{x}{n})^n$ as a function of a continuous n then we can find the following derivation:

$$\begin{aligned}
 \frac{d}{dn} \left[\left(1 + \frac{x}{n}\right)^n \right] &= \left(1 + \frac{x}{n}\right)^n \frac{d}{dn} \left[n \ln\left(1 + \frac{x}{n}\right) \right] \\
 &= \left(1 + \frac{x}{n}\right)^n \left(\frac{d}{dn} [n] \ln\left(1 + \frac{x}{n}\right) + n \frac{d}{dn} \left[\ln\left(1 + \frac{x}{n}\right) \right] \right) \\
 &= \left(1 + \frac{x}{n}\right)^n \left(\ln\left(1 + \frac{x}{n}\right) + \frac{n}{1 + \frac{x}{n}} \frac{d}{dn} \left[1 + \frac{x}{n} \right] \right) \\
 &= \left(1 + \frac{x}{n}\right)^n \left(\ln\left(1 + \frac{x}{n}\right) - \frac{x}{n + x} \right) \\
 &= \frac{\left(1 + \frac{x}{n}\right)^n}{x + n} ((x + n) \ln\left(1 + \frac{x}{n}\right) - x)
 \end{aligned}$$

By the last line of this we can see that because $(x, n) \in (\mathbb{R}^+)^2$ then $\ln(1 + \frac{x}{n}) > 0$ and thus we conclude that $(x + n) \ln(1 + \frac{x}{n}) - x > 0$. Therefore we see that $\frac{d}{dn} \left[\left(1 + \frac{x}{n}\right)^n \right] > 0$ for all $(x, n) \in \mathbb{R}^{+2}$, and in particular this means that $(1 + \frac{x}{n})^n < (1 + \frac{x}{n+1})^{n+1} \forall n \in \mathbb{N}$.

One consequence of this is that $(1 + \frac{x}{n})^n < e^x \forall n \in \mathbb{N}$, therefore we can define the error of algorithm 5.3.1 as $\epsilon_n := |e^x - (1 + \frac{x}{n})^n| = e^x - (1 + \frac{x}{n})^n$. Now as $\lim_{n \rightarrow \infty} (1 + \frac{x}{n})^n = e^x$ then we see that $\lim_{n \rightarrow \infty} \epsilon_n = 0$, and thus our algorithm is correct and valid for approximating e^x .

Next we see that this method, while simple, approximates e^x very poorly. In particular the table below shows the approximation of $e^{0.75}$ for different values of n , where the bold digits are the correctly approximated digits.

n	Approximation of $e^{0.75}$
1	1.80000000000000000044
10	2.1 58924997272786787
100	2.2 18468215957572747
1000	2.22 4829248807374831
10000	2.225 469716120127850
100000	2.2255 33806810873500
1000000	2.225540 216319864358
10000000	2.2255408 57275162929
100000000	2.22554092 1370736781
1000000000	2.225540927 780294606

With this table we see that the method very poorly approximates e^x , requiring a very large n to get just a few digits of accuracy. While this does not require more calculations from the method, requiring this large a value of n can lead to inaccuracies in the implementation of the algorithm using `double` data types in C.

In general there are better methods of approximating e^x and also $\ln(x)$, which while requiring more calculations are much more accurate than the most basic method presented here.

5.4 Taylor Series Method

If we take the elementary result from calculus that $\frac{d}{dx}e^x = e^x$, then we can calculate the Maclaurin series of e^x . By the definition of a Maclaurin series we know that the series expansion of e^x about 0 is

$$\sum_{k=0}^{\infty} \frac{d^{(k)}}{dx^k}[e^x](0) \frac{x^k}{k!}$$

As $\frac{d^{(k)}}{dx^k}[e^x] = e^x \forall k \in \mathbb{Z}_0^+$ and $e^0 = 1$ then we see that the series becomes

$$\sum_{k=0}^{\infty} \frac{x^k}{k!}$$

Using this we see that $e^x \approx \sum_{k=0}^n \frac{x^k}{k!}$, which gives the following method for approximating e^x :

Algorithm 5.4.1: Taylor Method for calculating e^x

```

1  taylor_exp( $x \in \mathbb{R}, n \in \mathbb{Z}_0^+$ )
2       $t = 1$ 
3       $z = 1$ 
4       $k = 1$ 
5      while  $k < n$ :
6           $t \mapsto \frac{t \cdot x}{n}$ 
7           $z \mapsto z + t$ 
8           $k \mapsto k + 1$ 
9      return  $z$ 
```

This allows us to calculate e^x more efficiently, and we can see that the error of the approximation is easy $\epsilon_n := |e^x - \sum_{k=0}^n \frac{x^k}{k!}| \leq \frac{|x|^{n+1}}{(n+1)!}$ for all $n \in \mathbb{Z}_0^+$. While we can't guarantee the size of x in general we will consider $x \in (0, 1)$ for the purposes of analysing this function.

As $x \in (0, 1)$ then it follows that $x < 1$ and thus we can see that $\epsilon_n < \frac{1}{n!} \forall n \in \mathbb{Z}_0^+$. Using this we can see that to use our method such that the error is at most $\tau_d := 10^{-d}$, then we need to find $n \in \mathbb{Z}_0^+ : \frac{1}{n!} < \tau_d$. The table below shows some values for (n, d) pairs such that n is the smallest positive integer such that $\frac{1}{n!} < \tau_d$:

$d \in \mathbb{N}$	$\arg \min \{n \in \mathbb{N} : n! > 10^d\}$
1	4
10	14
100	70
1000	450

Therefore we can guarantee 100 digits of accuracy with an input of $n \geq 70$ and 1000 digits of accuracy with $n \geq 450$, this is much less than our previous method where an input of $n = 1000$ only gave 2 decimal places of accuracy.

The inverse of the function $z = e^x$ is the logarithm function $\ln(z) = x$, which we can again consider for Taylor Series expansion. First we will show the result from calculus that $\frac{d}{dx}[\ln(x)] = \frac{1}{x}$:

Proposition 5.4.1.

$$\frac{d}{dx}[\ln(x)] = \frac{1}{x}$$

Proof. We will prove this from the first principles using the definition that $\frac{d}{dx}[f(x)] = \lim_{h \rightarrow 0} \frac{f(x+h) - f(x)}{h}$

$$\begin{aligned} \frac{d}{dx}[\ln(x)] &= \lim_{h \rightarrow 0} \frac{\ln(x+h) - \ln(x)}{h} \\ &= \lim_{h \rightarrow 0} \frac{\ln(1 + \frac{h}{x})}{h} \\ &= \lim_{h \rightarrow 0} \ln \left(\left(1 + \frac{h}{x}\right)^{\frac{1}{h}} \right) \end{aligned}$$

If we let $u := \frac{h}{x}$, then we get that $ux = h$ and $\frac{1}{h} = \frac{1}{ux}$. Also $\lim_{h \rightarrow 0} u = 0$, and so we get the following:

$$\begin{aligned} \frac{d}{dx}[\ln(x)] &= \lim_{u \rightarrow 0} \ln \left((1+u)^{\frac{1}{ux}} \right) \\ &= \frac{1}{x} \lim_{u \rightarrow 0} \ln \left((1+u)^{1/u} \right) \end{aligned}$$

If we now let $n := \frac{1}{u}$ and consider that $\lim_{u \rightarrow 0} n = \infty$, then our derivative becomes:

$$\begin{aligned} \frac{d}{dx} \ln(x) &= \frac{1}{x} \lim_{n \rightarrow \infty} \ln \left(\left(1 + \frac{1}{n}\right)^n \right) \\ &= \frac{1}{x} \ln \left(\lim_{n \rightarrow \infty} \left(1 + \frac{1}{n}\right)^n \right) \\ &= \frac{1}{x} \ln(e) && \text{by the definition of } e \\ &= \frac{1}{x} \end{aligned}$$

□

Now we know that $\frac{d^k}{dx^k} \left[\frac{1}{x} \right] = (-1)^k k! x^{-k-1}$, and thus we can build up a Taylor Series expansion. In this case, rather than centring the series about $x = 0$ for a Maclaurin series we can instead centre the series around $x = 1$ which gives the following series expansion for $\ln(x)$:

$$\begin{aligned}
\sum_{k=0}^{\infty} \frac{\frac{d^k}{dx^k}[\ln(x)](1)}{k!} (x-1)^k &= \ln(1) + \sum_{k=1}^{\infty} \frac{\frac{d^{k-1}}{dx^{k-1}}[x^{-1}](1)}{k!} (x-1)^k \\
&= \sum_{k=1}^{\infty} \frac{[(-1)^{k-1}(k-1)!x^{-k}](1)}{k!} (x-1)^k \\
&= \sum_{k=1}^{\infty} \frac{(-1)^{k-1}}{k} (x-1)^k \\
&= - \sum_{k=1}^{\infty} \frac{(1-x)^k}{k}
\end{aligned}$$

We know that $\ln(x) = - \sum_{k=1}^{\infty} \frac{(1-x)^k}{k}$ when the series $\sum_{k=1}^{\infty} \frac{(1-x)^k}{k}$ converges. We thus need to know when the sum converges.

Proposition 5.4.2. *The series $\sum_{k=1}^{\infty} \frac{(1-x)^k}{k}$ converges when $x \in (0, 2)$.*

Proof. Let $a_k := \frac{(1-x)^k}{k}$. We will proceed by using the ratio test to show when the series converges absolutely. The test states that the series converges when $\lim_{k \rightarrow \infty} \left| \frac{a_{k+1}}{a_k} \right| < 1$.

Now we can consider the following derivation:

$$\begin{aligned}
\lim_{k \rightarrow \infty} \left| \frac{a_{k+1}}{a_k} \right| &= \lim_{k \rightarrow \infty} \left| \frac{\frac{1}{k+1}(1-x)^{k+1}}{\frac{1}{k}(1-x)^k} \right| \\
&= \lim_{k \rightarrow \infty} \left| \frac{k}{k+1}(1-x) \right| \\
&= |1-x| \lim_{k \rightarrow \infty} \left| \frac{k}{k+1} \right| \\
&= |1-x|
\end{aligned}$$

Therefore our series converges when:

$$\begin{aligned}
|1-x| < 1 &\iff -1 < 1-x < 1 \\
&\iff -1 < x-1 < 1 \\
&\iff 0 < x < 2
\end{aligned}$$

Hence $\sum_{k=1}^{\infty} \frac{(1-x)^k}{k}$ converges when $x \in (0, 2)$. □

Now as we can't know if $x \in (0, 2)$ then we can consider that $\forall x \in \mathbb{R}^+ \exists (a, b) \in [\frac{1}{2}, 1) \times \mathbb{Z} : x = a \cdot 2^b$; thus we see that $\ln(x) = \ln(a \cdot 2^b) = b \ln(2) + \ln(a)$. As previously noted in Section 1.1 this operation, while theoretically complex, is simple to calculate for most computers by how they represent floating point values.

With this we can then use the following method to approximate $\ln(x)$ by the Taylor polynomial $-\sum_{k=1}^n \frac{(1-x)^k}{k}$:

Algorithm 5.4.2: Taylor Method for calculating $\ln(x)$

```

1  taylor_nat_log( $x \in \mathbb{R}^+, n \in \mathbb{N}$ ):
2      Find  $(a, b) \in [\frac{1}{2}, 1) \times \mathbb{Z}$  such that  $x = a \cdot 2^b$ 
3       $y := 1 - a$ 
4       $t := y$ 
5       $z := y$ 
6       $k := 1$ 
7      while  $k < n$ :
8           $t \mapsto t \cdot y$ 
9           $z \mapsto z + \frac{t}{k}$ 
10          $k \mapsto k + 1$ 
11     return  $b \ln(2) - z$ 

```

The first thing to consider for the above method is how to calculate $\ln(2)$. It is not possible to directly calculate $\ln(2)$ using the above algorithm as $2 = \frac{1}{2} \cdot 2^2$, however $\frac{1}{2} = \frac{1}{2} \cdot 2^0$ and so we do not need to know $\ln(2)$ to calculate $\ln(\frac{1}{2})$. We can see that $\ln(2) = -\ln(\frac{1}{2})$, and so we can calculate our constant value $\ln(2)$ to be used in the algorithm by using the algorithm itself.

Now similar to previous Taylor approximations the final error of our approximation using the above method is $\epsilon_n := |\ln(x) - \text{taylor_log}(x, n)|$. As the next term of the approximation would be $\frac{(1-x)^n}{n}$, then we know that $\epsilon_n \leq \left| \frac{(1-a)^n}{n} \right|$; further we know that $a \in [\frac{1}{2}, 1)$ and thus $\epsilon_n < \frac{1}{2^n n}$.

Using this approximation we can see that if we wish to guarantee d decimal places of accuracy then it suffices to find $n \in \mathbb{N}$ such that $\frac{1}{2^n n} < 10^{-d} \implies 2^n n > 10^d$. As $n \in \mathbb{N}$ then $2^n < 2^n n$ and so we merely need to find $n \in \mathbb{N}$ such that $2^n > 10^d$ to guarantee d decimal places of accuracy. Some example values are included in the table below:

$d \in \mathbb{N}$	$\arg \min \{n \in \mathbb{N} : 2^n > 10^d\}$
1	4
10	34
100	333
1000	3322

As we now have Taylor methods for approximating both e^x and $\ln(x)$, then we can use the two to derive a Taylor method of calculating x^y and $\log_x(y)$. To start we will consider $x^y = e^{y \ln(x)}$ and $x = a \cdot 2^b$, giving the solution as $x^y = e^{y(b \ln(2) + \ln(a))}$. Similarly we note that $\log_x(y) = \frac{\ln(y)}{\ln(x)}$, and if we consider that $x = a \cdot 2^b$ and $y = c \cdot 2^d$, then we see that $\log_x(y) = \frac{d \ln(2) + \ln(c)}{b \ln(2) + \ln(a)}$. Below are the Taylor methods for approximating these functions:

Algorithm 5.4.3: Taylor Method for calculating x^y and $\log_x(y)$

```

1  taylor_log( $x \in \mathbb{R}^+, y \in \mathbb{R}^+, n \in \mathbb{N}$ ):
2       $a := \text{taylor\_nat\_log}(y, n)$ 
3       $b := \text{taylor\_nat\_log}(x, n)$ 
4      return  $\frac{a}{b}$ 
5
6  taylor_pow( $x \in \mathbb{R}^+, y \in \mathbb{R}, n \in \mathbb{N}$ ):
7       $a := \text{taylor\_nat\_log}(x, n)$ 
8       $a \mapsto y \cdot a$ 

```

To test the convergence of the Taylor methods above we are going to test calculations of $7.3^{4.8}$, $7.3^{-4.8}$, $0.21^{4.8}$, $7.3^{0.21}$, $\log_{7.3}(4.8)$, $\log_{0.21}(4.8)$ and $\log_{7.3}(0.21)$. These values are calculated for several different values of n with the bold digits representing the correct values in the tables below:

n	$7.3^{4.8}$	$7.3^{-4.8}$	$0.21^{4.8}$	$7.3^{0.21}$
1	1.0000000000	1.0000000000	1.0000000000	1.0000000000
2	10.561319400	-8.561319400	-6.422212933	1.4183077237
3	56.076838311	36.990949511	21.518877680	1.5046585363
4	200.85920964	-107.8118783	-48.47602784	1.5167171202
5	546.24576990	237.58122696	82.710783892	1.5179778747
6	1205.3726532	-421.5471761	-113.8668463	1.5180831956
7	2253.5829747	626.66342673	131.57101558	1.5180905223
8	3682.4131809	-802.1668232	-131.0877429	1.5180909589
9	5386.6141612	902.03416303	114.86315726	1.5180909816
10	7193.4074522	-904.7591286	-89.85299062	1.5180909827
...
20	13901.238666	-11.00988984	-0.092958315	1.5180909827
...
40	13929.955484	0.0000717862	0.0005580236	1.5180909827
...
80	13929.955484	0.0000717877	0.0005580236	1.5180909827

As we can see in the table the taylor_pow does not converge perfectly, and may even diverge from the correct value for small values of n ; however we see that the methods do converge for large values of n . This behaviour is due to the values being outside the restrictions used in the analysis of the functions.

n	$\log_{7.3}(4.8)$	$\log_{0.21}(4.8)$	$\log_{7.3}(0.21)$
1	0.8431178860	-1.086107266	-0.776274970
2	0.8431178860	-1.086107266	-0.776274970
3	0.8045021618	-1.025878600	-0.784207957
4	0.7938608884	-1.011309817	-0.784982875
5	0.7906472231	-1.007102721	-0.785071082
6	0.7896173849	-1.005776909	-0.785082036
7	0.7892739993	-1.005337682	-0.785083473
8	0.7891562591	-1.005187460	-0.785083668
9	0.7891150494	-1.005134935	-0.785083695
10	0.7891003970	-1.005116266	-0.785083699
...
50	0.7890920869	-1.005105681	-0.785083699

This shows that taylor_log converges better than taylor_exp, however part of this is due to the values tested having magnitudes close to 1. Answers with a larger or smaller magnitudes tend to converge slower, which can be seen in the table for taylor_exp. The value that had best convergence in the taylor_exp table had an answer of about 1.5 and all other values tested had answers that were several orders of magnitude different from 1.

5.5 Hyperbolic Series Method

There are more efficient series which can be used to find \ln , which converge quicker than the Taylor approximation. One such method is to consider the Hyperbolic Trigonometric function \tanh . We start by considering the definition that $\tanh(x) := \frac{e^x - e^{-x}}{e^x + e^{-x}}$, and then find a formula for $\tanh^{-1}(x)$:

$$\begin{aligned}
 z = \frac{e^x - e^{-x}}{e^x + e^{-x}} &\implies z = \frac{e^{2x} - 1}{e^{2x} + 1} \\
 &\implies ze^{2x} + z = e^{2x} - 1 \\
 &\implies e^{2x}(1 - z) = 1 + z \\
 &\implies e^{2x} = \frac{1 + z}{1 - z} \\
 &\implies e^x = \left(\frac{1 + z}{1 - z} \right)^{\frac{1}{2}} \\
 &\implies x = \frac{1}{2} \ln \left(\frac{1 + z}{1 - z} \right)
 \end{aligned}$$

Using this we can see that $2 \tanh^{-1} \left(\frac{x-1}{x+1} \right) = \ln(x)$, and we can use the Taylor Expansion of \tanh^{-1} to approximate \ln .

Now to attain the Taylor series for $\tanh^{-1}(x)$ we can use the same method as when we calculated the Taylor series for \ln . The exact calculations are omitted, but the end result is that we get that:

$$\tanh^{-1}(x) = \sum_{n=0}^{\infty} \frac{x^{2n+1}}{2n+1} \quad \forall x \in \mathbb{R}^+$$

And thus by using this series we get the result that:

$$\ln(x) = 2 \sum_{n=0}^{\infty} \frac{1}{2n+1} \left(\frac{x-1}{x+1} \right)^{2n+1} \quad \forall x \in \mathbb{R}^+$$

The implementation of this is similar to previous implementations of series approximations of a function and is detailed below:

Algorithm 5.5.1: Hyperbolic series method for \ln

```

1  hyperbolic_nat_log( $x \in \mathbb{R}^+, n \in \mathbb{Z}_0^+$ ):
2       $a := \frac{x-1}{x+1}$ 
3       $b := a^2$ 
4       $c := a$ 
5       $k := 0$ 
6      while  $k \leq n$ :
7           $a \mapsto a \cdot b$ 
8           $c \mapsto c + \frac{a}{2k+1}$ 
9           $k \mapsto k + 1$ 
10     return  $2 \cdot c$ 

```

Using this we see that if we have $\epsilon_n := |\ln(x) - \text{hyperbolic_nat_log}(x, n)|$, then we know that $\epsilon_n \leq \frac{1}{2n+3} \left| \frac{x-1}{x+1} \right|^{2n+3}$. If we consider restricting our calculations to $x \in [\frac{1}{2}, 1)$ by using the same calculations as shown for algorithm 5.4.2, then we can see that $|x - 1| \leq \frac{1}{2}$ and $|x + 1| \geq \frac{3}{2}$; therefore $\epsilon_n \leq \frac{1}{3^{2n+3}(2n+3)}$.

By considering the final simplification that $\epsilon_n < \frac{1}{3^{2n+3}}$, then if we wish to have $\epsilon_n < \tau \in \mathbb{R}^+$ it suffices to find $n \in \mathbb{N}$ such that $\frac{1}{3^{2n+3}} < \tau$. In particular we consider when $\tau = 10^{-d}$ which will guarantee d decimal places of accuracy, below is a table showing the smallest $n \in \mathbb{N}$ that guarantees d decimal places of accuracy:

$d \in \mathbb{N}$	$\arg \min \{n \in \mathbb{N} : 3^{2n+3} > 10^d\}$
1	1
10	8
100	104
1000	1047

As can be seen in the table, fewer iterations are needed to approximate $\ln(x)$ to the same degree of accuracy using hyperbolic series as when using the Taylor series. Further, the calculations performed each iteration are very similar in complexity, both being $\mathcal{O}(1)$, and thus we can expect that algorithm 5.5.1 will execute faster than 5.4.2.

5.6 Continued fractions

Another method for evaluating e^x is the use of continued fractions, which are a way of approximating real functions by a rational number[13] with a recursive structure. Such fractions have been studied for many years and can be used to rationally approximate functions. Some examples of continued fractions for real numbers are[13, p. 266]:

$$e = 2 + \frac{1}{1 + \frac{1}{2 + \frac{1}{1 + \frac{1}{1 + \frac{1}{4 + \ddots}}}}}$$

$$\pi = 3 + \frac{1}{7 + \frac{1}{15 + \frac{1}{1 + \frac{1}{292 + \frac{1}{1 + \ddots}}}}}$$

In general a continued fraction for a number $x \in \mathbb{R}$ has the form:

$$b_0 + \frac{a_1}{b_1 + \frac{a_2}{b_2 + \frac{a_3}{b_3 + \ddots}}} \quad (5.6.1)$$

As the writing of continued fractions in the above manner takes up a lot of room and has a degree of ambiguity we will use the following notation:

$$\mathbf{K}_{n=1}^{\infty} \frac{a_n}{b_n} := \frac{a_1}{b_1 + \frac{a_2}{b_2 + \frac{a_3}{b_3 + \frac{a_4}{b_4 + \ddots}}}} \quad (5.6.2)$$

Therefore we can re-write Equation 5.6.1 as $b_0 + \mathbf{K}_{n=1}^{\infty} \frac{a_n}{b_n}$.

One of the most useful formulas regarding continued fractions was formulated by Leonhard Euler[5, Ch. 18], and deals with the sum $a_0 + a_0 a_1 + a_0 a_1 a_2 + \cdots + (a_0 \cdots a_n) = \sum_{i=0}^n (\prod_{j=0}^i a_j)$. The formula derived by Euler is known as Euler's Continued Fraction Formula and is as follows:

$$\sum_{i=0}^n \left(\prod_{j=0}^i a_j \right) = \mathbf{K}_{i=0}^n \frac{\alpha_i}{\beta_i} \text{ where } \alpha_i := \begin{cases} a_0 & : i = 0 \\ -a_i & : i \in [1, n] \cap \mathbb{Z} \end{cases} \quad (5.6.3)$$

$$\beta_i := \begin{cases} 1 & : i = 0 \\ 1 + a_i & : i \in [1, n] \cap \mathbb{Z} \end{cases}$$

Many Taylor series have a structure that is compatible with equation 5.6.3 and so can be approximated by a continued fraction in this way. In particular we are looking at $e^x = \sum_{k=0}^n \frac{x^k}{k!}$ where we note that we can re-write the series as $e^x = 1 + \sum_{i=1}^n (\prod_{j=1}^i \frac{x}{j})$ and therefore by using Euler's Continued Fraction Formula we see that:

$$e^x = 1 + \frac{x}{1 - \frac{\frac{1}{2}x}{1 + \frac{1}{2}x - \frac{\frac{1}{3}x}{\ddots - \frac{\frac{1}{n-1}x}{1 + \frac{1}{n-1}x - \frac{\frac{1}{n}x}{1 + \frac{1}{n}x}}}}} \quad (5.6.4)$$

$$= 1 + \mathbf{K}_{i=1}^n \frac{\alpha_i}{\beta_i} \quad \text{where } \alpha_i := \begin{cases} x & : i = 1 \\ -\frac{1}{i}x & : i \in [2, n] \cap \mathbb{Z} \end{cases}$$

$$\beta_i := \begin{cases} 1 & : i = 1 \\ 1 + \frac{1}{i}x & : i \in [2, n] \cap \mathbb{Z} \end{cases}$$

We want to simplify the above equation to remove the fractional coefficients. If we consider multiplying α_1 by some constant c_1 , then to have an equivalent fraction we would have to multiply it's denominator by c_1 ; in practice this means multiplying β_1 and α_2 by c_1 . We can suppose that we could continue in a similar manner for constants c_2, c_3, \dots multiplying $\alpha_2, \alpha_3, \dots$

Proposition 5.6.1. *If we have a continued fraction $b_0 + \mathbf{K}_{i=1}^n \frac{a_i}{b_i}$ and constants $(c_i : i \in [1, n] \cap \mathbb{Z})$, then:*

$$b_0 + \mathbf{K}_{i=1}^n \frac{a_i}{b_i} = b_0 + \mathbf{K}_{i=1}^n \frac{c_{i-1}c_i a_i}{c_i b_i}$$

where $c_0 := 1$, for any $n \in \mathbb{N}$.

Proof. We will proceed by induction on $n \in \mathbb{N}$.

$$\mathbf{H}(n): b_0 + \mathbf{K}_{i=1}^n \frac{a_i}{b_i} = b_0 + \mathbf{K}_{i=1}^n \frac{c_{i-1}c_i a_i}{c_i b_i}$$

$\mathbf{H}(1):$

$$\begin{aligned} b_0 + \frac{c_0 c_1 a_1}{c_1 b_1} &= b_0 + \frac{c_1 a_1}{c_1 b_1} && \text{as } c_0 = 1 \\ &= b_0 + \frac{a_1}{b_1} && \text{as required} \end{aligned}$$

$\mathbf{H}(n) \implies \mathbf{H}(n+1):$

$$\begin{aligned} b_0 + \mathbf{K}_{i=1}^{n+1} \frac{c_{i-1}c_i a_i}{c_i b_i} &= b_0 + \left(\mathbf{K}_{i=1}^n \frac{c_{i-1}c_i a_i}{c_i b_i} \right)_+ \frac{c_n c_{n+1} a_{n+1}}{c_{n+1} b_{n+1}} \\ &= b_0 + \left(\mathbf{K}_{i=1}^n \frac{c_{i-1}c_i a_i}{c_i b_i} \right)_+ c_n \frac{a_{n+1}}{b_{n+1}} \end{aligned}$$

Now let us define b'_i as:

$$\begin{cases} b_n + \frac{a_{n+1}}{b_{n+1}} & : i = n \\ b_i & : i \neq n \end{cases}$$

Therefore we can continue and see that:

$$\begin{aligned} b_0 + \mathbf{K}_{i=1}^{n+1} \frac{c_{i-1}c_i a_i}{c_i b_i} &= b_0 + \mathbf{K}_{i=1}^n \frac{c_{i-1}c_i a_i}{c_i b'_i} \\ &= b_0 + \mathbf{K}_{i=1}^n \frac{a_i}{b'_i} && \text{by } \mathbf{H}(n) \\ &= b_0 + \mathbf{K}_{i=1}^{n+1} \frac{a_i}{b_i} \end{aligned}$$

□

Using this proposition we can see that if we have the sequence (c_1, c_2, \dots, c_n) defined as $c_i = i$ and apply it to our sequence for e^x we get that:

$$\begin{aligned} e^x &= 1 + \frac{x}{1 - \frac{x}{2 + x - \frac{2x}{\ddots - \frac{(n-1)x}{n-1 + x - \frac{x}{n+x}}}}} \end{aligned} \tag{5.6.5}$$

$$= 1 + \mathbf{K}_{i=1}^n \frac{\alpha_i}{\beta_i} \quad \text{where } \alpha_i := \begin{cases} x & : i = 1 \\ -(i-1)x & : i \in [2, n] \cap \mathbb{Z} \end{cases}$$

$$\beta_i := \begin{cases} 1 & : i = 1 \\ x + i & : i \in [2, n] \cap \mathbb{Z} \end{cases}$$

This is a much simpler continued fraction, but evaluating it would still be computationally expensive due to the repeated division operations; to get around this we can consider what are known as the convergents of a continued fraction. It is obvious that if we use a continued fraction to approximate some value z by the continued fraction $b_0 + \mathbf{K}_{i=1}^n \frac{a_i}{b_i}$, then there are

some $A_n, B_n \in \mathbb{N}$ such that $z = \frac{A_n}{B_n}$.

To start we will define $A_{-1} := 1$ and $B_{-1} := 0$, and consider the case when $n = 0$; for this case $z = b_0$, which means that $A_0 = b_0$ and $B_0 = 1$. For the case when $n = 1$ we have $z = b_0 + \frac{a_1}{b_1}$, which when rearranged is $z = \frac{b_0 b_1 + a_1}{b_1}$. This means that $A_1 = b_0 b_1 + a_1 = b_1 A_0 + a_1 A_{-1}$ and $B_1 = b_1 = b_1 B_0 + a_1 B_{-1}$, and for the case when $n = 2$ we get the similar result that $A_2 = b_0 b_1 b_2 + a_2 b_0 + a_1 b_2 = b_2 A_1 + a_2 A_0$ and $B_2 = b_1 b_2 + a_2 = b_2 B_1 + a_2 B_0$.

It is actually true that this relationship continues for all $n \in \mathbb{N}$, and thus we get what are known as the Fundamental Recurrence Formulas for continued fractions:

$$\begin{aligned} A_{-1} &= 1 & B_{-1} &= 0 \\ A_0 &= b_0 & B_0 &= 1 \\ A_{n+1} &= b_{n+1} A_n + a_{n+1} A_{n-1} & B_{n+1} &= b_{n+1} B_n + a_{n+1} B_{n-1} \quad \forall n \in \mathbb{Z}_0^+ \end{aligned}$$

Using this and our simplified continued fraction for e^x we can use the following method to approximate e^x by using a continued fraction up to a_n, b_n where $n \geq 2$:

Algorithm 5.6.1: Continued fraction for e^x

```

1  cont_frac_exp( $x \in \mathbb{R}, n \in \mathbb{N}$ ):
2       $A_1 := x + 1$ 
3       $B_1 := 1$ 
4       $A_2 := x^2 + 2x + 2$ 
5       $B_2 := 2$ 
6       $a := -x$ 
7       $b := 2 + x$ 
8       $k := 2$ 
9      while  $k < n$ :
10          $a \mapsto a - x$ 
11          $b \mapsto b + 1$ 
12          $A_{k+1} := b A_k + a A_{k-1}$ 
13          $B_{k+1} := b B_k + a B_{k-1}$ 
14          $k \mapsto k + 1$ 
15  return  $\frac{A_k}{B_k}$ 

```

One observation of the above algorithm, when implemented on a computer, is that if we pre-generate b_i and a_i for $i \in [2, n] \cap \mathbb{Z}$ then the calculations of A_i and B_i are independent. This means that, if supported by the computer, both A_i and B_i could be computed in parallel. This may allow an implementation of the algorithm to be more efficient than one that computes the function in sequence.

While continued fractions are useful for approximating functions, it is difficult to evaluate the error of their output analytically. One result is that if $a_n = 1 \forall n \in \mathbb{N}$ when approximating some value z , then $|z - \frac{A_n}{B_n}| \leq \frac{1}{|B_{n+1} B_n|}$ [13]. If we transform the continued fraction of e^x into this form by using Proposition 5.6.1, then we get that:

$$e^x = 1 + \frac{1}{\left(-\frac{1}{x}\right) + \frac{1}{\left(-\frac{2+x}{2x}\right) + \frac{1}{\left(-\frac{3+x}{3x}\right) + \ddots}}}$$

By using a computer to implement the calculations for a test value of $x = 1$, we see that $\frac{1}{B_5 B_6} = 0.009131261889664\dots$ and $\frac{1}{B_{10} B_{11}} = 0.000041307209877\dots$; thus we can guarantee two decimal place of accuracy with `cont_frac_exp(1, 5)` and 4 with `cont_frac_exp(1, 10)`. However if we instead have $x = 14$ then $\frac{1}{B_{10} B_{11}} = 0.314711263190806\dots$ and convergence is similarly poor for negative values.

Further computations show that convergence of $x \in (0, 1)$ is better than the convergence when $x = 1$, and thus we can use the identities and conversions to ensure good convergence. In particular if $x \in \mathbb{Z}^+$ then we can calculate the reciprocal of `cont_frac_exp(-x, n)` and if $x \in (1, \infty)$ we use the identity that $x = a \cdot 2^b$; with this we see that $e^x = (e^a)^{2^b}$ and $2^b \in \mathbb{Z}^+$.

As $a \in (0, 1)$ and $2^b \in \mathbb{Z}^+$ then we can calculate $z = 2^a$ using algorithm 5.6.1. Then we can calculate z^{2^b} using algorithm 5.1.2, to find our approximation of e^x . Performing the calculation in this way allows us to use the our continued fraction method to guarantee fast convergence, and the $\mathcal{O}(1)$ integer exponential algorithm to guarantee the correct approximation without increasing the algorithmic complexity of the calculations by more than a constant factor.

With this restriction in place we know that algorithm 5.6.1 converges at least as quickly as it does for $x = 1$, and thus we can use its convergence to guarantee the convergence of our method. Below is a table that shows the minimum n needed to achieve the associated d decimal places of accuracy:

d	Minimum n to guarantee d decimal places of accuracy
1	2
10	22
100	235
1000	2386

An alternative continued fraction for e^x that arises from the generalized hyper geometric series[11] is:

$$e^x = \frac{1}{1 - \frac{x}{1 + \frac{x}{2 - \frac{x}{3 + \frac{x}{4 - \frac{x}{5 + \frac{x}{6 - \ddots}}}}}}} \quad (5.6.6)$$

$$= \mathbf{K}_{i=1}^n \frac{\alpha_i}{\beta_i}$$

$$\text{where } \alpha_i := \begin{cases} 1 & : i = 1 \\ -x & : i = 2 \\ (-1)^{i-1} \lfloor \frac{i-1}{2} \rfloor x & : i \in [3, \infty) \cap \mathbb{Z} \end{cases}$$

$$\beta_i := \begin{cases} 1 & : i = 1 \\ i - 1 & : i \in [2, \infty) \cap \mathbb{Z} \end{cases}$$

Due to the $(-1)^{i-1} \lfloor \frac{i-1}{2} \rfloor$ factor in the definition of α_i it is more efficient to perform two updates each step rather than one. Below is the implementation of this method:

Algorithm 5.6.2: Continued fraction for e^x version 2

```

1  cont_frac_exp_v2 ( $x \in \mathbb{R}, n \in \mathbb{N}$ ):
2       $A_1 := 1$ 
3       $B_1 := 1$ 
4       $A_2 := 1$ 
5       $B_2 := 1 - x$ 
6       $a := 1$ 
7       $b := 1$ 
8       $k := 2$ 
9      while  $k < n$ :
10          $a \mapsto xa$ 
11          $b \mapsto b + 1$ 
12          $A_{k+1} := bA_k + aA_{k-1}$ 
13          $B_{k+1} := bB_k + aB_{k-1}$ 
14          $k \mapsto k + 1$ 
15          $b \mapsto b + 1$ 
16          $A_{k+1} := bA_k - aA_{k-1}$ 
17          $B_{k+1} := bB_k - aB_{k-1}$ 
18          $k \mapsto k + 1$ 
19     return  $\frac{A_k}{B_k}$ 

```

The fraction needed to analyse this method is again found by using proposition 5.6.1, and is:

$$\begin{array}{c}
1 \\
\hline
1 + \frac{1}{\hline -\frac{1}{x} + \frac{1}{\hline -2 + \frac{1}{\hline \frac{3}{x} + \frac{1}{\hline 2 + \frac{1}{\hline -\frac{5}{x} + \frac{1}{\hline -2 + \ddots}}}}}}
\end{array}$$

By implementing this we get similar results to above, particularly there is rapid convergence for $x \in (0, 1)$. Further the convergence of values in $x \in (0, 1)$ is more rapid than $x = 1$ and so we can use the convergence of $x = 1$ as an upper bound of our method. Below is the table showing the smallest $n \in \mathbb{N}$ needed to ensure d decimal places of accuracy for some $d \in \mathbb{N}$:

d	Minimum n to guarantee d decimal places of accuracy
1	4
10	12
100	61
1000	405

As can be seen the convergence of 5.6.6 appears to be significantly faster than that of 5.6.5 and one might be satisfied by this, however an even better solution exists.

As the fraction 5.6.6 can be shown to converge for all values of x to e^x then we can consider the even and odd convergents. The even convergents of the sequence are $\frac{A_0}{B_0}, \frac{A_2}{B_2}, \frac{A_4}{B_4}, \dots$, while the odd convergents are $\frac{A_1}{B_1}, \frac{A_3}{B_3}, \frac{A_5}{B_5}, \dots$. As $\lim_{n \rightarrow \infty} \frac{A_n}{B_n} = e^x$, then $\lim_{n \rightarrow \infty} \frac{A_{2n}}{B_{2n}} = \lim_{n \rightarrow \infty} \frac{A_{2n+1}}{B_{2n+1}} = e^x$; the following proposition[13, p. 86] gives an explicit form for the odd and even convergents.

Proposition 5.6.2. *If $z = K_{i=1}^{\infty} \frac{a_i}{1}$, then the limit of the odd convergent of z is:*

$$x_{odd} = a_1 - \frac{a_1 a_2}{1 + a_2 + a_3 - \frac{a_3 a_4}{1 + a_4 + a_5 - \frac{a_5 a_6}{1 + a_6 + a_7 - \ddots}}}$$

while the limit of the even convergent is:

$$x_{even} = \frac{a_1}{1 + a_2 - \frac{a_2 a_3}{1 + a_3 + a_4 - \frac{a_4 a_5}{1 + a_5 + a_6 - \ddots}}}$$

Proof. Omitted □

If we apply proposition 5.6.1 to 5.6.6, to achieve the form $K_{i=1}^{\infty} \frac{a_i}{1}$ then we end up with the following fraction:

$$e^x = \frac{1}{1 + \frac{-x}{1 + \frac{\frac{1}{2}x}{1 + \frac{-\frac{1}{6}x}{1 + \frac{\frac{1}{6}x}{1 + \frac{-\frac{1}{10}x}{1 + \frac{\frac{1}{10}x}{1 + \frac{-\frac{1}{14}x}{1 + \ddots}}}}}}} \quad (5.6.7)$$

Now if we apply proposition 5.6.2 to the above fraction we see that:

$$e^x = 1 + \frac{x}{1 - x + \frac{1}{2}x + \frac{\frac{1}{12}x^2}{1 - \frac{1}{6}x + \frac{1}{6}x + \frac{\frac{1}{60}x^2}{1 - \frac{1}{10}x + \frac{1}{10}x + \frac{\frac{1}{140}x^2}{\ddots}}}} \quad (5.6.8)$$

Finally by simplifying and applying proposition 5.6.1 one more time we reach the following continued fraction for e^x :

$$e^x = 1 + \frac{2x}{2 - x + \frac{x^2}{6 + \frac{x^2}{10 + \frac{x^2}{14 + \ddots}}}} \quad (5.6.9)$$

$$= 1 + \mathbf{K}_{i=1}^{\infty} \frac{\alpha_i}{\beta_i} \quad \text{where } \alpha_i := \begin{cases} 2x & : i = 1 \\ x^2 & : i \in [2, \infty) \cap \mathbb{Z} \end{cases}$$

$$\beta_i := \begin{cases} 2 - x & : i = 1 \\ 4i - 2 & : i \in [2, \infty) \cap \mathbb{Z} \end{cases}$$

If we implement this method by using the Fundamental Recurrence Formula then we get the following:

Algorithm 5.6.3: Continued fraction for e^x version 3

```

1  cont_frac_exp_v3 ( $x \in \mathbb{R}, n \in \mathbb{N}$ ):
2     $A_0 := 1$ 
3     $B_0 := 1$ 
4     $A_1 := 2 + x$ 
5     $B_1 := 2 - x$ 
6     $a := x^2$ 
7     $b := 2$ 
8     $k := 1$ 
9    while  $1 < n$ :
10      $b \mapsto b + 4$ 

```

```

11       $A_{k+1} := bA_k + aA_{k-1}$ 
12       $B_{k+1} := bB_k + ab_{k-1}$ 
13       $k \mapsto k + 1$ 
14  return  $\frac{A_k}{B_k}$ 

```

As with the previous two continued fraction methods of approximating e^x we can apply proposition 5.6.1 to 5.6.9 to find the following equivalent continued fraction:

$$1 + \frac{1}{\frac{1}{x} - \frac{1}{2} + \frac{1}{\frac{12}{x} + \frac{1}{\frac{5}{x} + \frac{1}{\frac{28}{x} + \dots}}}}$$

Again a computer was used to evaluate B_k of the above fraction, which gave the expected results of quick convergence for $x \in (0, 1)$ and more rapid convergence for $x \in (0, 1)$ than $x = 1$. Using this the table below was generated to show the minimum $n \in \mathbb{N}$ that guarantees d digits of accuracy:

d	Minimum n to guarantee d decimal places of accuracy
1	2
10	6
100	30
1000	202

This has the fastest theoretical convergence of the three methods, and thus is expected to perform the best.

5.7 Comparison of Methods

We have introduced several methods for calculating both logarithms and exponentials in this chapter, and considered their theoretical convergence; we now look at a direct comparison of the methods as implemented in C.

The first consideration is which values to use while comparing methods. While all the methods converge for all values, or can be made to by using transformations of the inputs and outputs, most methods converge best for small values. Therefore values being tested will typically be in the range of $[0.5, 1)$.

The first methods to be compared here are the versions of the continued fraction method discussed previously. Below we have the outputs of different versions of the program for different values of n , with the bold digits being the correctly approximated digits.

n	cont_frac_exp_v1	cont_frac_exp_v2	cont_frac_exp_v3
1	1 .94499999999999	3.33333333333333	2 .0769230769230
2	2 .00216666666666	3.33333333333333	2 .0132689987937
3	2 .01217083333333	2 .0769230769230	2 .0137543842848
4	2 .01357141666666	2 .0054200542005	2 .0137527042253
5	2 .0137348180555	2 .0132689987937	2 .0137527074744
6	2 .0137511581944	2 .0137906192914	2 .0137527074704
7	2 .0137525879565	2 .0137543842848	2 .0137527074704
8	2 .0137526991603	2 .0137526161232	2 .0137527074704
9	2 .0137527069445	2 .0137527042253	2 .0137527074704
10	2 .0137527074399	2 .0137527076056	2 .0137527074704

As can be seen here the first two methods have similar convergence, however despite having a very poor theoretical convergence the first method converges better than the second version. Further, it is obvious that the third method has the fastest convergence, and thus should be the one to use in further comparisons.

Now we can compare the speed of the Taylor and continued fraction methods of calculating exponential values. For this we will use 1000 values in the range $[\frac{1}{2}, 1)$ and calculate each 100000 times to compare the speed of the method. We will be using values of n which guarantee 10 decimal places of accuracy, in particular $n = 14$ for `taylor_exp` and $n = 6$ for `cont_frac_exp_v3`.

The results of the tests run on my computer are included in the table below alongside those for the built in `exp` function in `math.h`:

	Total time:	Average time:	Minimum time:	Maximum time:
<code>taylor_exp</code>	12.430s	0.012s	0.012s	0.019s
<code>cont_frac_exp</code>	4.741s	0.004s	0.004s	0.007s
<code>builtin_exp</code>	2.608s	0.002s	0.002s	0.004s

This shows that the continued fractions method of evaluating exponential functions is almost three times as efficient as the standard Taylor series method. However both fall short of the built in method, despite the hyperbolic series method being a close second. This is likely due to a lower level implementation of the exponential function with various highly efficient programming practices implemented to optimize the code execution speed.

However one consideration is that if we instead test values in the range $[-5, 50]$, then while both `taylor_exp` and `cont_frac_exp` have similar results the total time for `cont_frac_exp` becomes 9.347s. This discrepancy is due to the additional calculations needed by `cont_frac_exp` so that it evaluates only values in the range $[\frac{1}{2}, 1)$ for a quicker convergence.

The two methods discussed to evaluate \ln have their convergence for different values of n shown below, where they are approximating the value 0.7, with the bold digits representing the correctly approximated digits:

n	taylor_nat_log	hyperbolic_nat_log
1	-0.300000000000	-0.356604925707
2	-0.300000000000	-0.356673383305
3	-0.345000000000	-0.356674906089
4	-0.354000000000	-0.356674942973
5	-0.356025000000	-0.356674943913
6	-0.356511000000	-0.356674943938
7	-0.356632500000	-0.356674943938
8	-0.356663742857	-0.356674943938
9	-0.356671944107	-0.356674943938
10	-0.356674131107	-0.356674943938

We can see here that the hyperbolic method converges a lot faster than the Taylor method; one particular note is that the hyperbolic series accurately approximates the first 12 decimal places of $\ln(0.7)$ accurately in just 6 iterations while the Taylor method only achieves 6 decimal places after 10 iterations.

To further test the two methods the table below shows the timings of calculating 1000 values in the range $[0.02, 50]$, each of which will be calculated to 10 decimal places 100000 times by each method. To guarantee 10 decimal places of accuracy with `taylor_log` we can use $n = 34$ and $n = 8$ for `hyperbolic_log`, below is the table that displays the results alongside the results for the built in `log` function:

	Total time:	Average time:	Minimum time:	Maximum time:
taylor_log	22.247s	0.022s	0.021s	0.026s
hyperbolic_log	7.742s	0.007s	0.007s	0.009s
builtin_log	3.438s	0.003s	0.003s	0.005s

Here we can see that the hyperbolic method of approximating $\ln(x)$ is the better of the two methods discussed, around three times faster in execution. While the built in function is, as to be expected, the fastest executing function, `hyperbolic_log` is not far behind, implying that `builtin_log` may use an optimized version of `hyperbolic_log`.

Finally we get to comparing the general exponential, x^y , and logarithm, $\log_x(y)$, functions. First we will test the convergence of the two variations of the $\log_x(y)$ function for different values of n , using $(x, y) = (1.5, 15)$:

n	taylor_log	hyperbolic_log
1	6.1155499597569	6.6784758082659
2	6.1155499597569	6.6788677803210
3	6.5747854684469	6.6788734950163
4	6.6587865280763	6.6788735857263
5	6.6748050386470	6.6788735872409
6	6.6780194099644	6.6788735872671
7	6.6786895339230	6.6788735872675
8	6.6788331533503	6.6788735872675
9	6.6788645711387	6.6788735872675
10	6.6788715529216	6.6788735872675

This table clearly demonstrates that `hyperbolic_log` converges faster to the correct value than `taylor_log` as expected. The table below shows the convergence of `taylor_pow` and `improved_pow` for the input of $(x, y) = (1.115, 15)$:

n	<code>taylor_pow</code>	<code>improved_pow</code>
1	1.00000000000000	5.7430173458025
2	4.7597077083991	5.1163939774264
3	5.9158698156503	5.1185134154921
4	5.6528248111124	5.1182823832710
5	5.3825631874287	5.1182688605223
6	5.2383576844918	5.1182679322812
7	5.1703487304639	5.1182678673534
8	5.1401274582517	5.1182678627291
9	5.1272612067887	5.1182678623951
10	5.1219353833296	5.1182678623708
...
20	5.1182684126550	5.1182678623688
...
30	5.1182678624756	5.1182678623688
...
40	5.1182678623689	5.1182678623688

Both of these methods for the general exponential function have slow convergences, though the improved method does converge faster. This implies that there may be a more efficient method for approximating x^y .

Next we will consider the actual speed of both the logarithm and exponential functions presented. One note is that C does not have a general log function for an arbitrary base in `math.h`, and so to implement this we will use `log(y)/log(x)` for the built in logarithm. All of the functions will have values of $(x, y) \in (0, 2] \times [0, 3)$ and will use a value of n sufficient to calculate their answer accurate to 10 decimal places. Below are the calculations for 1000 random values calculated 10000 times for each method:

	Total time:	Average time:	Minimum time:	Maximum time:
<code>taylor_log</code>	4.750s	0.004s	0.004s	0.008s
<code>hyperbolic_log</code>	1.589s	0.001s	0.001s	0.002s
<code>builtin_log</code>	0.690s	0.000s	0.000s	0.000s
<code>taylor_pow</code>	6.956s	0.006s	0.006s	0.007s
<code>improved_pow</code>	2.456s	0.002s	0.002s	0.003s
<code>builtin_pow</code>	0.787s	0.000s	0.000s	0.001s

Again we see that the methods that we showed to be theoretically superior, do in fact have superior execution speeds; however our methods still fail to match the execution speed of those built into C.

Overall we can conclude that if we were to want to implement calculating logarithms of a number then the hyperbolic series method is the best choice discussed, while the best choice

for evaluating exponentials is the continued fraction method.

The special case of the exponentiation by squaring is worth considering in the case where a computer only supports integers. This is because the algorithm will still work for integer only values, while most of the others will not, and has a computational complexity of $\mathcal{O}(1)$.

6 Conclusion

In this document we set out to consider different methods of calculating common functions that one may find on a calculator, as such we succeeded and now have a deeper understanding of these functions. We have also gained an insight into how many calculators or computers may operate in calculating these functions.

In studying the root functions we have seen that while there are various methods available the most efficient method is the inverse newton square root method. This method converges quadratically to the required root and has a faster computation time than the standard Newton method, due to the lack of division operations. We saw that both of these methods outstripped the linear convergence of the bisection method, which while simple and efficient in the computational complexity sense, takes many more steps to achieve comparable accuracy and so is less efficient in computational time.

The digit by digit method of calculating square roots is interesting but ultimately of little practical value for modern computers due to its poor efficiency, though it has the interesting accuracy property of generating precisely one new digit each iteration. Its integer square root counterpart on the other hand is particularly interesting due to its $\mathcal{O}(1)$ computational complexity and reliance on simple integer operations, and has possible practical applications if square roots are only needed to be accurate only to their integer part.

The root functions were successfully implemented in C, and when implemented with MPFR they were able to give answers accurate to arbitrary precision. In particular we were able to accurately compute $\sqrt{2}$ accurate to 1000000 decimal places in a reasonably short span of time.

The trigonometric functions are an engaging topic to study and in doing so we found several very different methods for approximating their values. The geometric method studied is conceptually simple, but turned out to be complex to analyse, the end result giving a method that had a low computational complexity per iteration but required many iterations to achieve accuracy comparable to other methods.

The Taylor method for trigonometric functions was found to be the most efficient method, once the range of inputs was restricted. Further this method was easy to analyse the accuracy of due to the nature of the Taylor series, making it simple to guarantee a given degree of accuracy.

The CORDIC algorithm was the least efficient of the methods analysed, but as mentioned earlier, it still has its place. In particular CORDIC is still useful for simple systems that do not have the capability to handle floating point values, or for which the floating point operations take a long time to compute. Further the CORDIC algorithm has the capability to be directly implemented in hardware which would guarantee its use as being the most efficient method.

Finally, in the analysis of the Logarithmic and Exponential functions we saw more methods, ranging from the trivial and naive, to the detailed and reasoned. As expected the more considered methods that took advantage of aspects of the functions being approximated had better results than those that did not.

The Taylor methods for approximating both logarithms and exponentials were good starting points, as the methods were conceptually simple with low computational complexity for each iteration. Similar to what was witnessed in the analysis of the trigonometric functions, it was very simple to calculate the number of iterations required for a given accuracy, which is a desirable property to have.

Unlike the trigonometric section there was no one method that could be used to the efficiency of both the exponential function and the logarithm function. However, the two methods considered, both gave significant increases in efficiency over the Taylor method. The analysis of the two resulting methods showed that they both converged at a faster rate than the standard Taylor method, and neither of the methods was significantly more computationally complex at each iteration.

A final note is that while our analysis has shown when one algorithm is better than another, and even achieved good computational times, they still fall short of the built in versions from the standard C libraries. This is due to either the libraries using even methods other than those discussed here, the libraries utilising low level programming techniques to speed up computation, or a combination of the two.

References

- [1] C. Aliprantis and O. Burkinshaw. *Principles of Real Analysis*. 2nd ed. Academic Press Limited, 1990. ISBN: 0120502550.
- [2] N. Artemiadis. *History of Mathematics. From a Mathematician's Vantage Point*. Trans. by N. Sofronidis. AMS, 2004. ISBN: 0821834037.
- [3] D. Chudnovsky and G. Chudnovsky. "The Computation of Classical Constants". In: *PNAS* (Aug. 1989). URL: <http://www.pnas.org/content/86/21/8178.full.pdf>.
- [4] Institute of Electrical and Electronics Engineers. *IEEE Standard for Floating-Point Arithmetic*. 2008. ISBN: 9780738157535. URL: <https://standards.ieee.org/findstds/standard/754-2008.html>.
- [5] L. Euler. *Introductio in analysin infinitorum*. 1748.
- [6] Exploratorium.edu. *A million digits of Pi*. URL: http://www.exploratorium.edu/pi/pi_archive/Pi10-6.html.
- [7] A. Greenbaum and T. Chartier. *Numerical Methods. Design, analysis, and computer implementation of algorithms*. Princeton University Press, 2012. ISBN: 97806911511229.
- [8] T. Henderson. "Cryptography and Complexity". 2012. URL: <http://hackthology.com/pdfs/crypto-complexity.pdf>.
- [9] E. Howard. *An Introduction to the History of Mathematics*. 6th ed. Saunders College Publishing, 1992. ISBN: 9780880294188.

- [10] IBM. *Implementation of cos in math.h*. URL: https://sourceware.org/git/?p=glibc.git;a=blob;f=sysdeps/ieee754/dbl-64/s_sin.c;hb=HEAD#l281.
- [11] W. Jones and W. Thron. *Continued Fractions. Analytic Theory and Applications*. Addison-Wesley Publishing Company, 1980.
- [12] L. Jordan. *Exponential Functions in the Real World*. URL: <https://www.sophia.org/tutorials/exponential-functions-in-the-real-world--3>.
- [13] L. Lorentzen and H. Waadeland. *Continued Fractions. Volume 1: Convergence Theory*. Ed. by C. Chui. 2nd ed. Atlantis Press, 2008. ISBN: 9789078677079.
- [14] B. McKeeman. "The Computation of Pi by Archimedes". MathWorks File Exchange. Nov. 2010. URL: <http://www.mathworks.com/matlabcentral/fileexchange/29504-the-computation-of-pi-by-archimedes/content/html/ComputationOfPiByArchimedes.html#37>.
- [15] F. Morgan. *Real Analysis*. AMS, 2005. ISBN: 0821836706.
- [16] Computer History Museum. *The Babbage Engine*. URL: <http://www.computerhistory.org/babbage/>.
- [17] R. Parris. "Elementary Functions and Calculators". URL: <http://math.exeter.edu/rparris/peanut/cordic.pdf>.
- [18] S. Ramanujan. "Modular equations and approximations to pi". In: *Quart J Math: Oxford* (1914).
- [19] G. Rising. *Inside Your Calculator. From Simple Programs to Significant Insights*. Wiley, 2007. ISBN: 978047011408.
- [20] J. Stewart. *Multivariable Calculus*. 6th ed. Brooks Cole, June 2007. ISBN: 9780495011637.
- [21] C. Stover. *Odd Function*. Wolfram Alpha. URL: <http://mathworld.wolfram.com/OddFunction.html>.
- [22] B. Taylor. *Methodus Incrementorum Directa et Inversa*. Direct and Reverse Methods of Incrementation. 1715.
- [23] *The GNU MPFR Library*. URL: <http://www.mpfr.org/>.
- [24] *The GNU Multiple Precision Arithmetic Library*. URL: <https://gmplib.org/>.
- [25] *Yale Babylonian Collection Images and Analysis*. URL: <http://www.math.ubc.ca/~cass/Euclid/ylbc/ylbc.html>.

A Code

In this appendix I list the entirety of the code which implement the algorithms discussed in the body of this document.

A.1 General Code

General Utilities File:

File : utilities.c

```
1 || #include <gmp.h>
2 || #include <mpfr.h>
```

```

3 | #include "utilities.h"
4 |
5 | const double ROOT_2      = 1.4142135623730950488016887242096980785696718753;
6 | const double ROOT_2_INV = 0.7071067811865475244008443621048490392848359376;
7 |
8 | inline unsigned int d(unsigned int D)
9 | {
10 |     return D > 10 ? 10 : D;
11 | }
12 |
13 | void inline mpfr_digits_to_tolerance(unsigned int D, mpfr_t T)
14 | {
15 |     mpfr_init_set_ui(T, 10, MPFR_RNDN);
16 |     mpfr_pow_ui(T, T, D, MPFR_RNDN);
17 |     mpfr_ui_div(T, 1, T, MPFR_RNDN);
18 | }

```

Trigonometric Utilities File:

File : trig_utilities.c

```

1 | #include <assert.h>
2 | #include "trig_utilities.h"
3 |
4 | TRIG_FIXED_TYPE double_to_fixed(double d)
5 | {
6 |     assert(d < 2 && d >= -2);
7 |     return (TRIG_FIXED_TYPE) (d * (TRIG_FIXED_TYPE)CONVERSION_VALUE);
8 | }
9 |
10 | double fixed_to_double(TRIG_FIXED_TYPE t)
11 | {
12 |     return (double)t / (TRIG_FIXED_TYPE)CONVERSION_VALUE;
13 | }

```

Header Files for Utilities:

File : utilities.h

```

1 | #ifndef UTILITIES_HEADER
2 |     #define UTILITIES_HEADER
3 |
4 |     #include <mpfr.h>
5 |
6 |     #define ROOT_2_INFILE "root_2_digits.txt"
7 |     #define ROOT_2_INV_INFILE "root_2_inv_digits.txt"
8 |
9 |     extern const double ROOT_2, ROOT_2_INV;
10 |     extern mpfr_t MPFR_ROOT_2, MPFR_ROOT_2_INV,
11 |                 MPFR_ONE, MPFR_HALF, MPFR_THREE_HALF, MPFR_TWO;
12 |
13 |     unsigned int d(unsigned int);
14 |     void mpfr_digits_to_tolerance(unsigned int, mpfr_t);
15 | #endif

```

File : trig_utilities.h

```

1 | #ifndef TRIG_UTILITIES
2 |     #define TRIG_UTILITIES
3 |
4 |     #include <mpfr.h>

```

```

5 | #include "trig_fixed.h"
6 |
7 | #define PI      3.1415926535897932384626433832795028841971693993751058
8 | #define HALF_PI 1.5707963267948966192313216916397514420985846996875529
9 | #define TWO_PI  6.2831853071795864769252867665590057683943387987502116
10 | #define PI_INFILE "pi_digits.txt"
11 |
12 | extern mpfr_t MPFR_PI, MPFR_HALF_PI, MPFR_TWO_PI;
13 |
14 | TRIG_FIXED_TYPE double_to_fixed(double);
15 | double fixed_to_double(TRIG_FIXED_TYPE);
16 | #endif

```

File : log_exp_utilities.h

```

1 | #ifndef LOG_EXP_UTILITIES_HEADER
2 | #define LOG_EXP_UTILITIES_HEADER
3 |
4 | #include <mpfr.h>
5 |
6 | #define NAT_LOG_2 0.693147180559945309417232121458176568075500134360
7 | #define E_CONST   2.718281828459045235360287471352662497757247093699
8 |
9 | #define NAT_LOG_2_INFILE "nat_log_2_digits.txt"
10 | #define E_CONST_INFILE  "e_digits.txt"
11 |
12 | extern mpfr_t MPFR_NAT_LOG_2, MPFR_E_CONST;
13 |
14 | #endif

```

Makefile for the project:

File : makefile

```

1 | #Compiler and basic flags
2 | CC=gcc
3 | CFLAGS=-std=c11 -g
4 |
5 | #Directories used
6 | OBJDIR=obj
7 | OUTDIR=out
8 | LIBDIR=lib
9 | TSTDIR=test
10 |
11 | #Library linking options
12 | MPFRLIB=-lgmp -lmpfr
13 | MATHLIB=$(MPFRLIB) -lm
14 | UTILLIB=-L$(LIBDIR) -lutil
15 |
16 | #Used to compile a given file with the main program
17 | EXE=-D COMPILE_MAIN
18 |
19 | #The lists of files that are used or created
20 | INFILES =bisect_root.c exact_root.c newton_inv_sqrt.c newton_sqrt.c \
21 |         geometric_trig.c geometric_inv_trig.c taylor_trig.c \
22 |         taylor_inv_trig.c cordic_trig.c int_exp.c taylor_exp_log.c \
23 |         hyperbolic_log.c cont_frac_exp.c
24 | TESTFILES=test_newton_sqrt.c test_trig_methods.c test_inv_trig_methods.c \
25 |          test_sqrt_methods.c test_int_exp_methods.c test_exp_methods.c \
26 |          test_log_methods.c test_log_pow_methods.c

```



```

27 OUTFILES=$(addprefix $(OUTDIR)/, $(INFILES:.c=.out))
28 OBJECTS=$(addprefix $(OBJDIR)/, $(INFILES:.c=.o))
29 TESTS=$(addprefix $(TSTDIR)/, $(TESTFILES:.c=.out))
30 LU=$(LIBDIR)/libutil.a
31
32 #Default option to build all the files
33 all: $(OBJECTS) $(LU) $(OUTFILES) $(TESTS)
34
35 #Cleans the workspace, best used for a fresh start
36 clean:
37     rm $(OBJDIR)/*.o $(OUTDIR)/*.out $(LIBDIR)/*.a
38
39 #The following few are used to build the utilities library
40 $(OBJDIR)/util.o: utilities.c utilities.h
41     $(CC) $(CFLAGS) -c $< -o $@
42
43 $(OBJDIR)/util_trig.o: trig_utilities.c trig_utilities.h
44     $(CC) $(CFLAGS) -c $< -o $@
45
46 $(OBJDIR)/util_test.o: $(TSTDIR)/test_utilities.c $(TSTDIR)/test_utilities.h
47     $(CC) $(CFLAGS) -c $< -o $@
48
49 $(LU): $(OBJDIR)/util.o $(OBJDIR)/util_trig.o $(OBJDIR)/util_test.o
50     ar -cr $@ $(OBJDIR)/util*.o
51
52 #How to compile a c file to an object file
53 $(OBJECTS): $(notdir $(@:.o=.c)) $(notdir $(@:.o=.h)) utilities.h \
54     trig_utilities.h
55     $(CC) $(CFLAGS) -c $(notdir $(@:.o=.c)) -o $@
56
57 #The following are how to compile each of the OUTFILES
58 # The accompanying entry is a shorthand for the first
59
60 ### SQUARE ROOT FILES ###
61 $(OUTDIR)/bisect_root.out: bisect_root.c bisect_root.h utilities.h $(LU)
62     $(CC) $(CFLAGS) $(EXE) bisect_root.c \
63     $(MATHLIB) $(UTILLIB) -o $@
64 bisect_root: $(OUTDIR)/bisect_root.out
65
66 $(OUTDIR)/exact_root.out: exact_root.c exact_root.h utilities.h $(LU)
67     $(CC) $(CFLAGS) $(EXE) exact_root.c \
68     $(MPFRLIB) $(UTILLIB) -o $@
69 exact_root: $(OUTDIR)/exact_root.out
70
71 $(OUTDIR)/newton_inv_sqrt.out: newton_inv_sqrt.c newton_inv_sqrt.h \
72     utilities.h $(LU)
73     $(CC) $(CFLAGS) $(EXE) newton_inv_sqrt.c \
74     $(MATHLIB) $(UTILLIB) -o $@
75 newton_inv_sqrt: $(OUTDIR)/newton_inv_sqrt.out
76
77 $(OUTDIR)/newton_sqrt.out: newton_sqrt.c newton_sqrt.h exact_root.h \
78     utilities.h $(OBJDIR)/exact_root.o $(LU)
79     $(CC) $(CFLAGS) $(EXE) newton_sqrt.c $(OBJDIR)/exact_root.o \
80     $(MATHLIB) $(UTILLIB) -o $@
81 newton_sqrt: $(OUTDIR)/newton_sqrt.out
82
83 ### TRIGONOMETRIC FILES ###
84 $(OUTDIR)/geometric_trig.out: geometric_trig.c trig_utilities.h \

```

```

85         geometric_trig.h $(LU)
86 $(CC) $(CFLAGS) $(EXE) geometric_trig.c $(UTILLIB)\
87 $(MATHLIB) -o $@
88 geometric_trig: $(OUTDIR)/geometric_trig.out
89
90 $(OUTDIR)/geometric_inv_trig.out: geometric_inv_trig.c trig_utilities.h \
91         geometric_inv_trig.h $(LU)
92 $(CC) $(CFLAGS) $(EXE) geometric_inv_trig.c $(UTILLIB)\
93 $(MATHLIB) -o $@
94 geometric_inv_trig: $(OUTDIR)/geometric_inv_trig.out
95
96 $(OUTDIR)/taylor_trig.out: taylor_trig.c trig_utilities.h \
97         taylor_trig.h $(LU)
98 $(CC) $(CFLAGS) $(EXE) taylor_trig.c $(UTILLIB)\
99 $(MATHLIB) -o $@
100 geometric_trig: $(OUTDIR)/taylor_trig.out
101
102 $(OUTDIR)/taylor_inv_trig.out: taylor_inv_trig.c trig_utilities.h \
103         taylor_inv_trig.h $(LU)
104 $(CC) $(CFLAGS) $(EXE) taylor_inv_trig.c $(UTILLIB)\
105 $(MATHLIB) -o $@
106 geometric_trig: $(OUTDIR)/taylor_inv_trig.out
107
108 $(OUTDIR)/cordic_trig.out: cordic_trig.c trig_utilities.h trig_fixed.h \
109         cordic_trig.h $(LU)
110 $(CC) $(CFLAGS) $(EXE) cordic_trig.c $(UTILLIB) $(MATHLIB) -o $@
111 cordic_trig: $(OUTDIR)/cordic_trig.out
112
113 ## LOGARITHM AND EXPONENTIAL FILES ##
114 $(OUTDIR)/int_exp.out: int_exp.c int_exp.h $(LU)
115 $(CC) $(CFLAGS) $(EXE) int_exp.c $(UTILLIB) $(MPFRLIB) -o $@
116 int_exp: $(OUTDIR)/int_exp.out
117
118 $(OUTDIR)/taylor_exp_log.out: taylor_exp_log.c taylor_exp_log.h $(LU) \
119         int_exp.h $(OBJDIR)/int_exp.o \
120         log_exp_utilities.h
121 $(CC) $(CFLAGS) $(EXE) taylor_exp_log.c $(OBJDIR)/int_exp.o \
122 $(UTILLIB) $(MATHLIB) -o $@
123 taylor_exp_log: $(OUTDIR)/taylor_exp_log.out
124
125 $(OUTDIR)/hyperbolic_log.out: hyperbolic_log.c hyperbolic_log.h $(LU) \
126         log_exp_utilities.h
127 $(CC) $(CFLAGS) $(EXE) hyperbolic_log.c $(UTILLIB) $(MATHLIB) -o $@
128 hyperbolic_log: $(OUTDIR)/hyperbolic_log.out
129
130 $(OUTDIR)/cont_frac_exp.out: cont_frac_exp.c cont_frac_exp.h int_exp.h \
131         hyperbolic_log.h $(LU)
132 $(CC) $(CFLAGS) $(EXE) cont_frac_exp.c $(OBJDIR)/int_exp.o \
133 $(OBJDIR)/hyperbolic_log.o $(UTILLIB) $(MATHLIB) -o $@
134 cont_frac_exp: $(OUTDIR)/cont_frac_exp.out
135
136 ## TESTING FILES ##
137 $(TSTDIR)/test_newton_sqrt.out: $(TSTDIR)/test_newton_sqrt.c \
138         $(TSTDIR)/test_utilities.h \
139         $(OBJDIR)/newton_sqrt.o $(OBJDIR)/exact_root.o \
140         $(LU)
141 $(CC) $(CFLAGS) $< $(OBJDIR)/newton_sqrt.o $(OBJDIR)/exact_root.o \
142 $(MPFRLIB) $(UTILLIB) -l. -o $@

```

```

143 test_newton_sqrt: $(TSTDIR)/test_newton_sqrt.out
144
145 $(TSTDIR)/test_trig_methods.out: $(TSTDIR)/test_trig_methods.c \
146     $(TSTDIR)/test_utilities.h \
147     $(OBJDIR)/geometric_trig.o $(OBJDIR)/taylor_trig.o \
148     $(OBJDIR)/cordic_trig.o $(LU)
149 $(CC) $(CFLAGS) $< $(OBJDIR)/geometric_trig.o $(OBJDIR)/taylor_trig.o \
150     $(OBJDIR)/cordic_trig.o $(MATHLIB) $(UTILLIB) -l. -o $(@)
151 test_trig_methods: $(TSTDIR)/test_trig_methods.out
152
153 $(TSTDIR)/test_inv_trig_methods.out: $(TSTDIR)/test_inv_trig_methods.c \
154     $(TSTDIR)/test_utilities.h \
155     $(OBJDIR)/geometric_inv_trig.o \
156     $(OBJDIR)/taylor_inv_trig.o \
157     $(OBJDIR)/cordic_trig.o $(LU)
158 $(CC) $(CFLAGS) $< $(OBJDIR)/geometric_inv_trig.o \
159     $(OBJDIR)/taylor_inv_trig.o \
160     $(OBJDIR)/cordic_trig.o $(MATHLIB) $(UTILLIB) -l. -o $(@)
161 test_inv_trig_methods: $(TSTDIR)/test_inv_trig_methods.out
162
163 $(TSTDIR)/test_sqrt_methods.out: $(TSTDIR)/test_sqrt_methods.c \
164     $(TSTDIR)/test_utilities.h \
165     $(OBJDIR)/exact_root.o $(OBJDIR)/bisect_root.o \
166     $(OBJDIR)/newton_sqrt.o $(OBJDIR)/newton_inv_sqrt.o $(LU)
167 $(CC) $(CFLAGS) $< $(OBJDIR)/exact_root.o $(OBJDIR)/bisect_root.o \
168     $(OBJDIR)/newton_sqrt.o $(OBJDIR)/newton_inv_sqrt.o \
169     $(MATHLIB) $(UTILLIB) -l. -o $(@)
170 test_sqrt_methods: $(TSTDIR)/test_sqrt_methods.out
171
172 $(TSTDIR)/test_int_exp_methods.out: $(TSTDIR)/test_int_exp_methods.c \
173     $(TSTDIR)/test_utilities.h $(OBJDIR)/int_exp.o $(LU)
174 $(CC) $(CFLAGS) $< $(OBJDIR)/int_exp.o $(MPFRLIB) $(UTILLIB) -l. -o $@
175 test_int_exp_methods: $(TSTDIR)/test_int_exp_methods.out;
176
177 $(TSTDIR)/test_exp_methods.out: $(TSTDIR)/test_exp_methods.c \
178     $(TSTDIR)/test_utilities.h \
179     $(OBJDIR)/taylor_exp_log.o \
180     $(OBJDIR)/cont_frac_exp.o \
181     $(OBJDIR)/int_exp.o $(LU)
182 $(CC) $(CFLAGS) $< $(OBJDIR)/taylor_exp_log.o \
183     $(OBJDIR)/cont_frac_exp.o $(OBJDIR)/int_exp.o \
184     $(OBJDIR)/hyperbolic_log.o $(MATHLIB) $(UTILLIB) -l. -o $@
185 test_exp_methods: $(TSTDIR)/test_exp_methods.out
186
187 $(TSTDIR)/test_log_methods.out: $(TSTDIR)/test_log_methods.c \
188     $(TSTDIR)/test_utilities.h \
189     $(OBJDIR)/taylor_exp_log.o \
190     $(OBJDIR)/hyperbolic_log.o \
191     $(OBJDIR)/int_exp.o $(LU)
192 $(CC) $(CFLAGS) $< $(OBJDIR)/taylor_exp_log.o $(OBJDIR)/int_exp.o \
193     $(OBJDIR)/hyperbolic_log.o $(MATHLIB) $(UTILLIB) -l. -o $@
194 test_log_methods: $(TSTDIR)/test_log_methods.out
195
196 $(TSTDIR)/test_log_pow_methods.out: $(TSTDIR)/test_log_pow_methods.c \
197     $(TSTDIR)/test_utilities.h $(OBJDIR)/taylor_exp_log.o \
198     $(OBJDIR)/hyperbolic_log.o $(OBJDIR)/int_exp.o \
199     $(OBJDIR)/cont_frac_exp.o $(LU)
200 $(CC) $(CFLAGS) $< $(OBJDIR)/taylor_exp_log.o $(OBJDIR)/int_exp.o \

```

```

201 $(OBJDIR)/hyperbolic_log.o $(OBJDIR)/cont_frac_exp.o \
202 $(MATHLIB) $(UTILLIB) -l. -o $@
203 test_log_pow_methods: $(TSTDIR)/test_log_pow_methods.out

```

A.2 Square Root Code

Code for Exact Square Root Methods:

File : exact_root.c

```

1 #include <gmp.h>
2 #include <mpfr.h>
3 #include <stdlib.h>
4 #include <stdio.h>
5 #include <inttypes.h>
6
7 #include "utilities.h"
8 #include "exact_root.h"
9
10 char *root_digits_precise(char *N, unsigned int D)
11 {
12     //Counter variables
13     unsigned int i, a;
14     //The offset value used to set the correct character's value
15     unsigned int o = 0;
16     //Real and Integer types from GMP and MPFR used for precision
17     mpfr_t Yr, Nr, T, tmpr_0, tmpr_1;
18     mpz_t P, tmpz, Yz;
19
20     //Allocates memory for the number of digits requested plus 5 to be safe
21     char *R = malloc((D+5) * sizeof(*R));
22
23     //Sets Nr from the provided string representing N
24     mpfr_init_set_str(Nr, N, 10, MPFR_RNDN);
25     mpfr_init(Yr);
26     mpfr_init(tmpr_0);
27     mpfr_init(tmpr_1);
28
29     //P will be used to keep track of the current partial solution
30     mpz_init_set_ui(P, 0);
31     mpz_init(Yz);
32     mpz_init(tmpz);
33
34     //T represents the power of the 10 that the current digit represents
35     //T is initially floor(n/2) where  $N = K \cdot 10^n$  and K is in [0, 10)
36     //T is of the form  $10^t$ 
37     mpfr_init(T);
38     //The mpfr_log10 function is used here as placeholder
39     // this may be replaced by my own log function later
40     mpfr_log10(T, Nr, MPFR_RNDN);
41     mpfr_div_ui(T, T, 2, MPFR_RNDN);
42     mpfr_floor(T, T);
43     //Similar to log, but for exponentiation
44     mpfr_exp10(T, T, MPFR_RNDN);
45
46     //This takes into account numbers of the form 0.x
47     if(mpfr_cmp_ui(T, 1) < 0)
48     {

```

```

49     R[0] = '0';
50     R[1] = '.';
51     //Offset set to 2 to indicate there are 2 pre-assigned characters
52     o = 2;
53 }
54
55 //Main loop
56 for(i=0; i <= D; i++)
57 {
58     //Calculates 10^(2t) and 20*P
59     mpfr_mul(tmpr_0, T, T, MPFR_RNDN);
60     mpz_mul_ui(tmpz, P, 20);
61     //tmpr_1 is used to prevent re-calculation later on
62     mpfr_set_ui(tmpr_1, 0, MPFR_RNDN);
63
64     /*
65     This loop stops when any digit produces a Y too large or all
66     digits have been considered.
67     In both cases a will be one greater than required and thus
68     must be decremented afterwards
69     */
70     for(a=1; a <= 9; a++)
71     {
72         //Calculates N - (20*P + a)*a*10^(2t)
73         mpz_add_ui(Yz, tmpz, a);
74         mpz_mul_ui(Yz, Yz, a);
75         mpfr_mul_z(Yr, tmpr_0, Yz, MPFR_RNDN);
76
77         if(mpfr_cmp(Yr, Nr) > 0)
78             //The exit condition for the loop has been met
79             break;
80         else
81             //tmpr_1 updated to remove the need for re-calculation
82             mpfr_set(tmpr_1, Yr, MPFR_RNDN);
83     }
84
85     //Decrements a and adds the correct digit to the result string
86     a--;
87     R[i+o] = DIGITS[a];
88
89     //Reduces Nr by the largest Yr found in the previous loop
90     mpfr_sub(Nr, Nr, tmpr_1, MPFR_RNDN);
91
92     //Break if an exact solution is found
93     //Note that due to the representation of floating point numbers it
94     // is possible to have found an exact solution with a positive
95     // remainder that is very close to zero. Unfortunately there is
96     // no way to test for this without knowing, the exact precision
97     // of the input beforehand.
98     if(mpfr_cmp_ui(Nr, 0) == 0)
99     {
100         //This loop adds 0s to a string where an exact solution has
101         // been found but needs right padding with zeros.
102         while(mpfr_cmp_ui(T, 1) > 0)
103         {
104             R[++i + o] = '0';
105             mpfr_div_ui(T, T, 10, MPFR_RNDN);
106         }

```

```

107     break;
108 }
109
110 //Calculates  $P = 10 * P + a$ 
111 mpz_mul_ui(P, P, 10);
112 mpz_add_ui(P, P, a);
113 //Calculates  $T = T/10 \Rightarrow 10^t \rightarrow 10^{(t-1)}$ 
114 mpfr_div_ui(T, T, 10, MPFR_RNDN);
115
116 //If we have dropped below  $10^0$  for the first time then add
117 // a '.' to the result string and increase the offset to 1
118 //This case only occurs if no '.' is in the string already
119 if(o == 0 && mpfr_cmp_ui(T, 1) < 0)
120 {
121     o = 1;
122     R[i+o] = '.';
123 }
124 }
125
126 //Adds a null character to terminate the string
127 R[i+o+1] = '\0';
128 return R;
129 }
130
131 //The use of uintmax_t gives the largest number of unsigned integers
132 // for which this function will work with.
133 uintmax_t uint_sqrt(uintmax_t num)
134 {
135     //Represents the value of  $2r(2^m)$ , where r is the
136     // current known part of the integer root
137     uintmax_t res = 0;
138     //Represents the largest power of  $(2^m)^2 = 4^m$ , the initial value
139     // is calculated as 011...11 XOR 0011...11 as the size
140     // of uintmax_t is not known beforehand
141     uintmax_t bit = (UINTMAX_MAX >> 1) ^ (UINTMAX_MAX >> 2);
142
143     //Finds the largest power of 4 that is at most 'num' in value
144     while(bit > num)
145         bit >>= 2;
146
147     //while(bit) is equivalent to while(bit > 0) for unsigned integers
148     while(bit)
149     {
150         //Checks the two cases for updating 'res' and 'num'
151         if(num >= res + bit)
152         {
153             // 'num' is used to keep track of the difference between
154             // r and the original value, N, that was to be rooted.
155             num -= res + bit;
156             //This calculates 'res'  $\rightarrow 2(r + 2^m) * 2^{(m-1)}$  using addition
157             // and bitshifting
158             res = (res >> 1) + bit;
159         }
160         //In the other case 'res'  $\rightarrow 2r(2^{m-1})$ 
161         else
162             res >>= 1;
163
164         //Move on to the next lower power of 2

```

```

165     bit >>= 2;
166 }
167
168 //Returns the integer part of the square root
169 return res;
170 }
171
172 #ifdef COMPILE_MAIN
173 int main(int argc, char **argv)
174 {
175     uintmax_t N;
176     unsigned int p, d;
177     char *R;
178
179     if (argc == 1)
180     {
181         printf("Usage: %s [a/b] [arguments]", argv[0]);
182         exit(1);
183     }
184
185     switch(argv[1][0])
186     {
187         case 'a':
188             if(argc == 5 &&
189                 sscanf(argv[3], "%u", &d) == 1 &&
190                 sscanf(argv[4], "%u", &p) == 1)
191             {
192                 mpfr_set_default_prec(p);
193                 printf("sqrt(%s) ~= \n\t%s", argv[2],
194                     root_digits_precise(argv[2], d));
195             }
196             else
197                 printf("Usage: %s a <N=Number to sqrt> "
198                     "<d=Number of significant digits> "
199                     "<p=bits of precision to use>\n", argv[0]);
200             break;
201
202         case 'b':
203             if(argc == 3 &&
204                 sscanf(argv[2], "%u" SCNuMAX, &N) == 1)
205                 printf("int_sqrt(%" PRIuMAX ") ~= %" PRIuMAX "\n",
206                     N, uint_sqrt(N));
207             else
208                 printf("Usage: %s b <N=Positive integer to sqrt>\n",
209                     argv[0]);
210             break;
211
212         default:
213             printf("Usage: %s [a/b] [arguments]", argv[0]);
214     }
215 }
216 #endif

```

Code for the Bisection Methods:

File : bisect_root.c

```

1 | #include <stdio.h>
2 | #include <stdlib.h>

```

```

3 #include <gmp.h>
4 #include <mpfr.h>
5 #include <assert.h>
6 #include <math.h>
7
8 #include "bisect_root.h"
9 #include "utilities.h"
10
11 #define INIT_CONSTANTS mpfr_init_set_ui(MPFR_ONE, 1, MPFR_RNDN); \
12     mpfr_init_set_d(MPFR_HALF, 0.5, MPFR_RNDN);
13
14 mpfr_t MPFR_ONE, MPFR_HALF;
15
16 double bisect_sqrt(double N, double T)
17 {
18     assert(N >= 0);
19     assert(T >= 0);
20
21     int e;
22     double a, b, x, f;
23
24     N = frexp(N, &e);
25     if(e%2)
26     {
27         N /= 2;
28         e += 1;
29     }
30
31     //Sets the initial values of a and b
32     a = 0;
33     b = 1;
34
35     x = 0.5*(a + b);
36     f = x*x - N;
37
38     //fabs(f) > T is our approximation of
39     // f != 0, by using the given tolerance
40     while(fabs(f) > T && b - a > T)
41     {
42         //Update of the bounds a and b
43         if (f < 0)
44             a = x;
45         else
46             b = x;
47
48         //Update of x and f
49         x = 0.5*(a + b);
50         f = x*x - N;
51     }
52
53     return ldexp(x, e / 2);
54 }
55
56 double bisect_sqrt_it(double N, unsigned int l)
57 {
58     assert(N >= 0);
59
60     int e;

```



```

61 double a, b, x, f;
62
63 N = frexp(N, &e);
64 if (e%2)
65 {
66     N /= 2;
67     e += 1;
68 }
69
70 //Sets the initial values of a and b
71 a = 0;
72 b = 1;
73
74 x = 0.5*(a + b);
75 f = x*x - N;
76
77 //fabs(f) > T is our approximation of
78 // f != 0, by using the given tolerance
79 for(int i = 0; i < I; ++i)
80 {
81     //Update of the bounds a and b
82     if (f < 0)
83         a = x;
84     else
85         b = x;
86
87     //Update of x and f
88     x = 0.5*(a + b);
89     f = x*x - N;
90 }
91
92 return ldexp(x, e / 2);
93 }
94
95 double iPow(double x, unsigned int n)
96 {
97     double r = 1;
98     while(n--)
99         r *= x;
100     return r;
101 }
102
103 double bisect_nRoot(double N, double T, unsigned int n)
104 {
105     assert(N >= 0);
106     assert(T >= 0);
107     //Ensures that none of the trivial cases are requested
108     assert(n >= 2);
109
110     //Runs the more optimal bisect_sqrt if n == 2
111     if(n == 2)
112         return bisect_sqrt(N, T);
113
114     double a, b, x, f;
115
116     //Sets the initial values of a and b
117     a = 0;
118     //This statement is equivalent to

```

```

119 // if(N<1) b=1; else b=N;
120 b = N < 1 ? 1 : N;
121
122 x = 0.5*(a + b);
123 f = iPow(x, n) - N;
124
125 //fabs(f) > T is our approximation of
126 // f != 0, by using the given tolerance
127 while(fabs(f) > T && b - a > T)
128 {
129     //Update of the bounds a and b
130     if (f < 0)
131         a = x;
132     else
133         b = x;
134
135     //Update of x and f
136     x = 0.5*(a + b);
137     f = iPow(x, n) - N;
138 }
139
140 return x;
141 }
142
143 void mpfr_bisect_sqrt(mpfr_t R, mpfr_t N, mpfr_t T)
144 {
145     if(mpfr_cmp_ui(N, 0) < 0)
146     {
147         fprintf(stderr, "The value to square root must be non-negative\n");
148         exit(-1);
149     }
150     if(mpfr_cmp_ui(T, 0) < 0)
151     {
152         fprintf(stderr, "The tolerance must be non-negative\n");
153         exit(-1);
154     }
155
156     mpfr_exp_t e;
157     mpfr_t a, b, x, f, d, fab, n;
158
159     mpfr_init(n);
160     mpfr_frexp(&e, n, N, MPFR_RNDN);
161     if(e%2)
162     {
163         mpfr_div_ui(n, n, 2, MPFR_RNDN);
164         e += 1;
165     }
166
167     //Set a == 0
168     mpfr_init_set_ui(a, 0, MPFR_RNDN);
169
170     //Set b == 1
171     mpfr_init_set_ui(b, 1, MPFR_RNDN);
172
173     //Set x = (a + b)/2
174     mpfr_init(x);
175     mpfr_add(x, a, b, MPFR_RNDN);
176     mpfr_mul(x, x, MPFR_HALF, MPFR_RNDN);

```

```

177
178 //Set  $f = x^2 - N$  and  $fab = |f|$ 
179 mpfr_init(f);
180 mpfr_init(fab);
181 mpfr_mul(f, x, x, MPFR_RNDN);
182 mpfr_sub(f, f, N, MPFR_RNDN);
183 mpfr_abs(fab, f, MPFR_RNDN);
184
185 //Set  $d = b - a$ 
186 mpfr_init(d);
187 mpfr_sub(d, b, a, MPFR_RNDN);
188
189 while(mpfr_cmp(fab, T) > 0 && mpfr_cmp(d, T) > 0)
190 {
191     //Update the bounds, a and b
192     if(mpfr_cmp_ui(f, 0) < 0)
193         mpfr_set(a, x, MPFR_RNDN);
194     else
195         mpfr_set(b, x, MPFR_RNDN);
196
197     //Update x
198     mpfr_add(x, a, b, MPFR_RNDN);
199     mpfr_mul(x, x, MPFR_HALF, MPFR_RNDN);
200
201     //Update f and fab
202     mpfr_mul(f, x, x, MPFR_RNDN);
203     mpfr_sub(f, f, n, MPFR_RNDN);
204     mpfr_abs(fab, f, MPFR_RNDN);
205 }
206
207 printf("beep");
208 mpfr_mul_2si(R, x, e/2, MPFR_RNDN);
209 }
210
211 void mpfr_bisect_nRoot(mpfr_t R, mpfr_t N, mpfr_t T, unsigned int n)
212 {
213     if(mpfr_cmp_ui(N, 0) < 0)
214     {
215         fprintf(stderr, "The value to square root must be non-negative\n");
216         exit(-1);
217     }
218     if(mpfr_cmp_ui(T, 0) < 0)
219     {
220         fprintf(stderr, "The tolerance must be non-negative\n");
221         exit(-1);
222     }
223     assert(n >= 2);
224
225     mpfr_t a, b, x, f, d, fab;
226
227     //Set  $a = 0$ 
228     mpfr_init_set_ui(a, 0, MPFR_RNDN);
229
230     //Set  $b = \max\{1, N\}$ 
231     mpfr_init(b);
232     mpfr_max(b, MPFR_ONE, N, MPFR_RNDN);
233
234     //Set  $x = (a + b)/2$ 

```

```

235 mpfr_init(x);
236 mpfr_add(x, a, b, MPFR_RNDN);
237 mpfr_mul(x, x, MPFR_HALF, MPFR_RNDN);
238
239 //Set f = x^2 - N
240 mpfr_init(f);
241 mpfr_init(fab);
242 mpfr_pow_ui(f, x, n, MPFR_RNDN);
243 mpfr_sub(f, f, N, MPFR_RNDN);
244 mpfr_abs(fab, f, MPFR_RNDN);
245
246 //Set d = b - a
247 mpfr_init(d);
248 mpfr_sub(d, b, a, MPFR_RNDN);
249
250 while(mpfr_cmp(fab, T) > 0 && mpfr_cmp(d, T) > 0)
251 {
252     //Update the bounds, a and b
253     if(mpfr_cmp_ui(f, 0) < 0)
254         mpfr_set(a, x, MPFR_RNDN);
255     else
256         mpfr_set(b, x, MPFR_RNDN);
257
258     //Update x
259     mpfr_add(x, a, b, MPFR_RNDN);
260     mpfr_mul(x, x, MPFR_HALF, MPFR_RNDN);
261
262     //Update f
263     mpfr_pow_ui(f, x, n, MPFR_RNDN);
264     mpfr_sub(f, f, N, MPFR_RNDN);
265     mpfr_abs(fab, f, MPFR_RNDN);
266 }
267
268 mpfr_set(R, x, MPFR_RNDN);
269 }
270
271 #ifdef COMPILE_MAIN
272 int main(int argc, char** argv)
273 {
274     double N, T;
275     unsigned int n, D, p;
276     mpfr_t Nr, Tr, R;
277     int c;
278     char sf[50];
279
280     if (argc == 1)
281     {
282         printf("Usage: %s [a/b/c/d/e] [arguments]\n", argv[0]);
283         exit(1);
284     }
285
286     switch(argv[1][0])
287     {
288         case 'a':
289             if (argc == 5 &&
290                 sscanf(argv[2], "%lf", &N) == 1 &&
291                 sscanf(argv[3], "%lf", &T) == 1 &&
292                 sscanf(argv[4], "%u", &D) == 1)

```

```

293     printf("sqrt(%.1f) =~ %.1f\n", d(D), N, D, bisect_sqrt(N, T));
294 else
295     printf("Usage: %s a <N=Value to sqrt> "
296           "<T=Tolerance> <D=Number of digits to display>\n",
297           argv[0]);
298     break;
299
300 case 'b':
301     if (argc == 6 &&
302         sscanf(argv[2], "%.1f", &N) == 1 &&
303         sscanf(argv[3], "%.1f", &T) == 1 &&
304         sscanf(argv[4], "%u", &n) == 1 &&
305         sscanf(argv[5], "%u", &D) == 1)
306         printf("%u_Root(%.1f) =~ %.1f\n",
307               n, d(D), N, D, bisect_nRoot(N, T, n));
308     else
309         printf("Usage: %s b <N=Value to root> <T=Tolerance>"
310               "<n=nth Root> <D=Number of digits to display>\n",
311               argv[0]);
312     break;
313
314 case 'c':
315     if (argc == 5 &&
316         sscanf(argv[3], "%u", &D) == 1 &&
317         sscanf(argv[4], "%u", &p) == 1)
318     {
319         mpfr_set_default_prec(p);
320         INIT_CONSTANTS
321
322         if (mpfr_init_set_str(Nr, argv[2], 10, MPFR_RNDN) == 0)
323         {
324             mpfr_init(R);
325             //Sets the tolerance to Tr = 10^-D
326             mpfr_digits_to_tolerance(D, Tr);
327
328             //Generates the required format string
329             sprintf(sf, "sqrt(%%.%%uRNf) =~\n\t%%.%%uRNf\n", d(D), D);
330
331             mpfr_bisect_sqrt(R, Nr, Tr);
332             mpfr_printf(sf, Nr, R);
333         }
334     else
335         printf("Usage: %s c <N=Value to sqrt> "
336               "<D=Number of digits to calculate to> "
337               "<p=bits of precision>\n", argv[0]);
338     }
339     else
340         printf("Usage: %s c <N=Value to sqrt> "
341               "<D=Number of digits to calculate to> "
342               "<p=bits of precision>\n", argv[0]);
343     break;
344
345 case 'd':
346     if (argc == 6 &&
347         sscanf(argv[3], "%u", &D) == 1 &&
348         sscanf(argv[4], "%u", &n) == 1 &&
349         sscanf(argv[5], "%u", &p) == 1)
350     {

```

```

351     mpfr_set_default_prec(p);
352     INIT_CONSTANTS
353
354     if (mpfr_init_set_str(Nr, argv[2], 10, MPFR_RNDN) == 0)
355     {
356         mpfr_init(R);
357
358         //Sets the tolerance to Tr = 10^-D
359         mpfr_digits_to_tolerance(D, Tr);
360
361         //Generates the required format string
362         sprintf(sf, "%u_Root(%%uRNf) =~\n\t%%uRNf\n", d(D), D);
363
364         mpfr_bisect_nRoot(R, Nr, Tr, n);
365         mpfr_printf(sf, n, Nr, R);
366     }
367     else
368         printf("Usage: %s d <N=Value to root> "
369              "<D=Number of digits to calculate to> "
370              "<n=nth root> <p=bits of precision>\n", argv[0]);
371 }
372 else
373     printf("Usage: %s d <N=Value to root> "
374          "<D=Number of digits to calculate to> "
375          "<n=nth root> <p=bits of precision>\n", argv[0]);
376 break;
377
378 case 'e':
379     if (argc == 5 &&
380         sscanf(argv[2], "%lf", &N) == 1 &&
381         sscanf(argv[3], "%u", &p) == 1 &&
382         sscanf(argv[4], "%u", &D) == 1)
383         printf("sqrt(%.*lf) =~ %.*lf\n", d(D), N, D,
384              bisect_sqrt_it(N, p));
385     else
386         printf("Usage: %s a <N=Value to sqrt> "
387              "<I=iterations> <D=Number of digits to display>\n",
388              argv[0]);
389     break;
390 default:
391     printf("Usage: %s [a/b/c/d/e] [arguments]", argv[0]);
392 }
393 }
394 #endif

```

Code for Newton Square Root Methods:

File : newton_root.c

```

1  #include <stdio.h>
2  #include <stdlib.h>
3  #include <gmp.h>
4  #include <mpfr.h>
5  #include <assert.h>
6  #include <math.h>
7
8  #include "utilities.h"
9  #include "exact_root.h"
10 #include "newton_sqrt.h"

```

```

11
12 #define INIT_CONSTANTS mpfr_init_set_d(MPFR_HALF, 0.5, MPFR_RNDN); \
13     in = fopen(ROOT_2_INFILE, "r"); \
14     mpfr_init(MPFR_ROOT_2); \
15     mpfr_inp_str(MPFR_ROOT_2, in, 10, MPFR_RNDN); \
16     fclose(in); \
17     in = fopen(ROOT_2_INV_INFILE, "r"); \
18     mpfr_init(MPFR_ROOT_2_INV); \
19     mpfr_inp_str(MPFR_ROOT_2_INV, in, 10, MPFR_RNDN); \
20     fclose(in);
21
22 mpfr_t MPFR_ROOT_2, MPFR_ROOT_2_INV, MPFR_HALF;
23
24 double newton_sqrt_v1(double N, double T)
25 {
26     assert(N >= 0);
27     assert(T >= 0);
28
29     double x, px, d;
30
31     x = N > 1 ? N : 1;
32     d = 1000000;
33
34     while(d > T)
35     {
36         px = x;
37         x = 0.5 * (x + N/x);
38         d = fabs(x - px);
39     }
40     return x;
41 }
42
43 double newton_sqrt_v2(double N, double T)
44 {
45     assert(N >= 0);
46     assert(T >= 0);
47
48     double x, px, d;
49
50     if(N >= 4)
51         x = uint_sqrt((unsigned long) N);
52     else
53         x = N > 1 ? N : 1;
54
55     d = 1000000;
56
57     while(d > T)
58     {
59         px = x;
60         x = 0.5 * (x + N/x);
61         d = fabs(x - px);
62     }
63
64     return x;
65 }
66
67 double newton_sqrt_v3(double N, double T)
68 {

```

```

69  assert(N >= 0);
70  assert(T >= 0);
71
72  int e;
73  double x, px, d;
74
75  N = frexp(N, &e);
76
77  x = 1;
78  d = 1000000;
79
80  while(d > T)
81  {
82      px = x;
83      x = 0.5 * (x + N/x);
84      d = fabs(x - px);
85  }
86
87  if(e%2)
88      x *= e > 0 ? ROOT_2 : ROOT_2_INV;
89  return ldexp(x, e / 2);
90 }
91
92 double newton_sqrt_v3_it(double N, unsigned int l)
93 {
94     assert(N >= 0);
95
96     int e;
97     double x;
98
99     N = frexp(N, &e);
100
101     x = 1;
102
103     for(int i = 0; i < l; ++i)
104         x = 0.5 * (x + N/x);
105
106     if(e%2)
107         x *= e > 0 ? ROOT_2 : ROOT_2_INV;
108     return ldexp(x, e / 2);
109 }
110
111
112 void mpfr_newton_sqrt_v3(mpfr_t R, mpfr_t N, mpfr_t T)
113 {
114     mpfr_t x, px, d, t, n;
115     mpfr_exp_t e;
116
117     mpfr_init(n);
118     mpfr_frexp(&e, n, N, MPFR_RNDN);
119
120     mpfr_init_set_ui(x, 1, MPFR_RNDN);
121     mpfr_init(px);
122     mpfr_init_set_ui(d, 1000000, MPFR_RNDN);
123     mpfr_init(t);
124
125     while(mpfr_cmp(d, T) > 0)
126     {

```



```

127     mpfr_set(px, x, MPFR_RNDN);
128     mpfr_div(t, n, x, MPFR_RNDN);
129     mpfr_add(x, x, t, MPFR_RNDN);
130     mpfr_mul(x, MPFR_HALF, x, MPFR_RNDN);
131     mpfr_sub(d, x, px, MPFR_RNDN);
132     mpfr_abs(d, d, MPFR_RNDN);
133 }
134
135 if(e%2)
136     if(e > 0)
137         mpfr_mul(x, MPFR_ROOT_2, x, MPFR_RNDN);
138     else
139         mpfr_mul(x, MPFR_ROOT_2_INV, x, MPFR_RNDN);
140 mpfr_mul_2si(R, x, e/2, MPFR_RNDN);
141 }
142
143 #ifdef COMPILE_MAIN
144 int main(int argc, char **argv)
145 {
146     double N, T;
147     unsigned int n, D, p;
148     mpfr_t Nr, Tr, R;
149     int c;
150     char sf[50];
151     FILE *in;
152
153     if(argc==1)
154     {
155         printf("Usage: %s [a/b/c/d/e] <Arguments>\n", argv[0]);
156         exit(1);
157     }
158
159     switch(argv[1][0])
160     {
161         case 'a':
162             if (argc == 5 &&
163                 sscanf(argv[2], "%lf", &N) == 1 &&
164                 sscanf(argv[3], "%lf", &T) == 1 &&
165                 sscanf(argv[4], "%u", &D) == 1)
166                 printf("sqrt(%.1f) ~ %.1f\n", d(D), N, D, newton_sqrt_v1(N, T));
167             else
168                 printf("Usage: %s a <N=Value to sqrt> "
169                     "<T=Tolerance> <D=Number of digits to display>\n",
170                     argv[0]);
171             break;
172
173         case 'b':
174             if (argc == 5 &&
175                 sscanf(argv[2], "%lf", &N) == 1 &&
176                 sscanf(argv[3], "%lf", &T) == 1 &&
177                 sscanf(argv[4], "%u", &D) == 1)
178                 printf("sqrt(%.1f) ~ %.1f\n", d(D), N, D, newton_sqrt_v2(N, T));
179             else
180                 printf("Usage: %s b <N=Value to sqrt> "
181                     "<T=Tolerance> <D=Number of digits to display>\n",
182                     argv[0]);
183             break;
184

```

```

185 case 'c':
186     if (argc == 5 &&
187         sscanf(argv[2], "%lf", &N) == 1 &&
188         sscanf(argv[3], "%lf", &T) == 1 &&
189         sscanf(argv[4], "%u", &D) == 1)
190         printf("sqrt(%.*lf) ~ %.*lf\n", d(D), N, D, newton_sqrt_v3(N, T));
191     else
192         printf("Usage: %s c <N=Value to sqrt> "
193             "<T=Tolerance> <D=Number of digits to display>\n",
194             argv[0]);
195     break;
196
197 case 'd':
198     if (argc == 5 &&
199         sscanf(argv[3], "%u", &D) == 1 &&
200         sscanf(argv[4], "%u", &p) == 1)
201     {
202         mpfr_set_default_prec(p);
203         INIT_CONSTANTS
204
205         if (mpfr_init_set_str(Nr, argv[2], 10, MPFR_RNDN) == 0)
206         {
207             mpfr_init(R);
208
209             mpfr_digits_to_tolerance(D, Tr);
210
211             sprintf(sf, "sqrt(%%.%%uRNf) ~\n\t%%.%%uRNf\n", d(D), D);
212
213             mpfr_newton_sqrt_v3(R, Nr, Tr);
214             mpfr_printf(sf, Nr, R);
215         }
216     else
217         printf("Usage: %s d <N=Value to sqrt> "
218             "<D=Number of digitsto calculate to> "
219             "<p=bits of precision>\n", argv[0]);
220     }
221     else
222         printf("Usage: %s d <N=Value to sqrt> "
223             "<D=Number of digitsto calculate to> "
224             "<p=bits of precision>\n", argv[0]);
225     break;
226
227 case 'e':
228     if (argc == 5 &&
229         sscanf(argv[2], "%lf", &N) == 1 &&
230         sscanf(argv[3], "%u", &p) == 1 &&
231         sscanf(argv[4], "%u", &D) == 1)
232         printf("sqrt(%.*lf) ~ %.*lf\n", d(D), N, D,
233             newton_sqrt_v3_it(N, p));
234     else
235         printf("Usage: %s e <N=Value to sqrt> "
236             "<I=Iterations> <D=Number of digits to display>\n",
237             argv[0]);
238     break;
239 default:
240     printf("Usage: %s [a/b/c/d/e] <Arguments>\n", argv[0]);
241 }
242 }

```

Code for Newton Inverse Square Root Methods:

File : newton_inv_sqrt.c

```

1  #include <stdio.h>
2  #include <stdlib.h>
3  #include <gmp.h>
4  #include <mpfr.h>
5  #include <assert.h>
6  #include <math.h>
7
8  #include "utilities.h"
9  #include "newton_inv_sqrt.h"
10
11 #define INIT_CONSTANTS mpfr_init_set_d(MPFR_THREE_HALF, 1.5, MPFR_RNDN); \
12     in = fopen(ROOT_2_INFILE, "r"); \
13     mpfr_init(MPFR_ROOT_2); \
14     mpfr_inp_str(MPFR_ROOT_2, in, 10, MPFR_RNDN); \
15     fclose(in); \
16     in = fopen(ROOT_2_INV_INFILE, "r"); \
17     mpfr_init(MPFR_ROOT_2_INV); \
18     mpfr_inp_str(MPFR_ROOT_2_INV, in, 10, MPFR_RNDN); \
19     fclose(in);
20
21 mpfr_t MPFR_ROOT_2, MPFR_ROOT_2_INV, MPFR_THREE_HALF;
22
23 double newton_inv_sqrt(double N, double T)
24 {
25     assert(N >= 0);
26     assert(T >= 0);
27
28     int e;
29     double x, px, d, NN, N_2;
30
31     NN = N;
32     N = frexp(N, &e);
33     N_2 = 0.5*N;
34
35     x = 1;
36     d = 1000000;
37
38     while(d > T)
39     {
40         px = x;
41         x = x * (1.5 - N_2*x*x);
42         d = fabs(x - px);
43     }
44
45     if(e%2)
46         x *= e > 0 ? ROOT_2_INV : ROOT_2;
47     x *= NN;
48     return ldexp(x, -e / 2);
49 }
50
51 double newton_inv_sqrt_it(double N, unsigned int l)
52 {
53     assert(N >= 0);

```

```

54
55     int e;
56     double x, NN, N_2;
57
58     NN = N;
59     N = frexp(N, &e);
60     N_2 = 0.5*N;
61
62     x = 1;
63
64     for(int i = 0; i < I; ++i)
65         x = x * (1.5 - N_2*x*x);
66
67     if(e%2)
68         x *= e > 0 ? ROOT_2_INV : ROOT_2;
69     x *= NN;
70     return ldexp(x, -e / 2);
71 }
72
73 void mpfr_newton_inv_sqrt(mpfr_t R, mpfr_t N, mpfr_t T)
74 {
75     mpfr_t x, px, d, t, n, n_2;
76     mpfr_exp_t e;
77
78     mpfr_init(n);
79     mpfr_frexp(&e, n, N, MPFR_RNDN);
80     mpfr_init_set(n_2, n, MPFR_RNDN);
81     mpfr_div_ui(n_2, n, 2, MPFR_RNDN);
82
83     mpfr_init_set_ui(x, 1, MPFR_RNDN);
84     mpfr_init(px);
85     mpfr_init_set_ui(d, 1000000, MPFR_RNDN);
86     mpfr_init(t);
87
88     while(mpfr_cmp(d, T) > 0)
89     {
90         mpfr_set(px, x, MPFR_RNDN);
91         mpfr_mul(t, x, x, MPFR_RNDN);
92         mpfr_mul(t, t, n_2, MPFR_RNDN);
93         mpfr_sub(t, MPFR_THREE_HALF, t, MPFR_RNDN);
94         mpfr_mul(x, x, t, MPFR_RNDN);
95         mpfr_sub(d, x, px, MPFR_RNDN);
96         mpfr_abs(d, d, MPFR_RNDN);
97     }
98
99     if(e%2)
100         if(e > 0)
101             mpfr_mul(x, MPFR_ROOT_2_INV, x, MPFR_RNDN);
102         else
103             mpfr_mul(x, MPFR_ROOT_2, x, MPFR_RNDN);
104     mpfr_mul(x, x, N, MPFR_RNDN);
105     mpfr_mul_2si(R, x, -e/2, MPFR_RNDN);
106 }
107
108 #ifdef COMPILE_MAIN
109 int main(int argc, char **argv)
110 {
111     double N, T;

```

```

112 unsigned int n, D, p;
113 mpfr_t Nr, Tr, R;
114 int c;
115 char sf[50];
116 FILE *in;
117
118 if(argc == 1)
119 {
120     printf(" Usage: %s [a/b/c] <Arguments>\n", argv[0]);
121     exit(1);
122 }
123
124 switch(argv[1][0])
125 {
126     case 'a':
127         if (argc == 5 &&
128             sscanf(argv[2], "%lf", &N) == 1 &&
129             sscanf(argv[3], "%lf", &T) == 1 &&
130             sscanf(argv[4], "%u", &D) == 1)
131             printf(" sqrt(%.*lf) =~ %.*lf\n", d(D), N, D, newton_inv_sqrt(N, T));
132         else
133             printf(" Usage: %s a <N=Value to sqrt> "
134                 "<T=Tolerance> <D=Number of digits to display>\n",
135                 argv[0]);
136         break;
137
138     case 'b':
139         if (argc == 5 &&
140             sscanf(argv[3], "%u", &D) == 1 &&
141             sscanf(argv[4], "%u", &p) == 1)
142         {
143             mpfr_set_default_prec(p);
144             INIT_CONSTANTS
145
146             if(mpfr_init_set_str(Nr, argv[2], 10, MPFR_RNDN) == 0)
147             {
148                 mpfr_init(R);
149
150                 mpfr_digits_to_tolerance(D, Tr);
151
152                 sprintf(sf, " sqrt(%%.%%uRNf) =~\n\t%.%%uRNf\n", d(D), D);
153
154                 mpfr_newton_inv_sqrt(R, Nr, Tr);
155                 mpfr_printf(sf, Nr, R);
156             }
157         else
158             printf(" Usage: %s b <N=Value to sqrt> "
159                 "<D=Number of digits to calculate to> "
160                 "<p=bits of precision>\n", argv[0]);
161         }
162     else
163         printf(" Usage: %s b <N=Value to sqrt> "
164             "<D=Number of digits to calculate to> "
165             "<p=bits of precision>\n", argv[0]);
166         break;
167
168     case 'c':
169         if (argc == 5 &&

```

```

170         sscanf(argv[2], "%lf", &N) == 1 &&
171         sscanf(argv[3], "%u", &p) == 1 &&
172         sscanf(argv[4], "%u", &D) == 1)
173     printf(" sqrt(%.*lf) ~ %.*lf\n", d(D), N, D,
174           newton_inv_sqrt_it(N, p));
175     else
176         printf(" Usage: %s a <N=Value to sqrt> "
177               "<I=iterations> <D=Number of digits to display>\n",
178               argv[0]);
179     break;
180
181     default:
182         printf(" Usage: %s [a/b/c] <Arguments>\n", argv[0]);
183     }
184 }
185 #endif

```

Header files for Square Root Code:

File : exact_root.h

```

1 #ifndef EXACT_ROOT_HEADER
2 #define EXACT_ROOT_HEADER
3     #include <inttypes.h>
4
5     static const char *DIGITS = "0123456789";
6
7     char *root_digits_precise(char*, unsigned int);
8     uintmax_t uint_sqrt(uintmax_t);
9 #endif

```

File : bisect_root.h

```

1 #ifndef BISECT_ROOT_HEADER
2 #define BISECT_ROOT_HEADER
3
4     double bisect_sqrt(double, double);
5     double bisect_sqrt_it(double, unsigned int);
6     double ipow(double, unsigned int);
7     double bisect_nRoot(double, double, unsigned int);
8     void mpfr_bisect_sqrt(mpfr_t, mpfr_t, mpfr_t);
9     void mpfr_bisect_nRoot(mpfr_t, mpfr_t, mpfr_t, unsigned int);
10 #endif

```

File : newton_root.h

```

1 #ifndef NEWTON_ROOT_HEADER
2 #define NEWTON_ROOT_HEADER
3
4     double newton_sqrt_v1(double, double);
5     double newton_sqrt_v2(double, double);
6     double newton_sqrt_v3(double, double);
7     double newton_sqrt_v3_it(double, unsigned int);
8     void mpfr_newton_sqrt_v3(mpfr_t, mpfr_t, mpfr_t);
9 #endif

```

File : newton_inv_sqrt.h

```

1 #ifndef NEWTON_INV_SQRT_HEADER
2 #define NEWTON_INV_SQRT_HEADER

```

```

3
4 double newton_inv_sqrt(double, double);
5 double newton_inv_sqrt_it(double, unsigned int);
6 void mpfr_newton_inv_sqrt(mpfr_t, mpfr_t, mpfr_t);
7 #endif

```

A.3 Trigonometric Code

Code for Geometric Trigonometric Functions:

File : geometric_trig.c

```

1 #include <stdio.h>
2 #include <assert.h>
3 #include <math.h>
4 #include <gmp.h>
5 #include <mpfr.h>
6
7 #include "geometric_trig.h"
8 #include "trig_utilities.h"
9 #include "utilities.h"
10
11 #define INIT_CONSTANTS in = fopen(PI_INFILE, "r"); \
12     mpfr_init(MPFR_PI); \
13     mpfr_inp_str(MPFR_PI, in, 10, MPFR_RNDN); \
14     fclose(in); \
15     mpfr_init(MPFR_TWO_PI); \
16     mpfr_init(MPFR_HALF_PI); \
17     mpfr_div_ui(MPFR_HALF_PI, MPFR_PI, 2, MPFR_RNDN); \
18     mpfr_mul_ui(MPFR_TWO_PI, MPFR_PI, 2, MPFR_RNDN);
19
20 mpfr_t MPFR_PI, MPFR_HALF_PI, MPFR_TWO_PI;
21
22 double geometric_cos_bounded(double x, unsigned int n)
23 {
24     //Ensures that x is in the range [0, HALF_PI) and raises an error
25     // message if this is not the case.
26     assert(x >= 0 && x <= HALF_PI);
27
28     //Sets the first chord length that will be the basis of our induction
29     double h = (x*x)/pow(4, n);
30
31     //Performs the induction steps
32     for(int i = 0; i < n; i++)
33         h = h*(4-h);
34     //Returns the approximation of cos(x)
35     return 1 - h/2;
36 }
37
38 void mpfr_geometric_cos_bounded(mpfr_t R, mpfr_t x, unsigned int n)
39 {
40     assert(mpfr_cmp_ui(x, 0) >= 0 && mpfr_cmp(x, MPFR_HALF_PI) <= 0);
41
42     mpfr_t h, t;
43     mpz_t k;
44
45     mpfr_init(t);
46

```

```

47 mpfr_init(h);
48 mpfr_mul(h, x, x, MPFR_RNDN);
49 mpz_init(k);
50 mpz_ui_pow_ui(k, 4, n);
51 mpfr_div_z(h, h, k, MPFR_RNDN);
52
53
54 for(int i = 0; i < n; i++)
55 {
56     mpfr_ui_sub(t, 4, h, MPFR_RNDN);
57     mpfr_mul(h, h, t, MPFR_RNDN);
58 }
59
60 mpfr_div_ui(h, h, 2, MPFR_RNDN);
61 mpfr_ui_sub(R, 1, h, MPFR_RNDN);
62 }
63
64
65 double geometric_cos(double x, unsigned int n)
66 {
67     // We have two cases to consider,  $x \geq 0$  and  $x < 0$ 
68     if(x >= 0)
69     {
70         //Ensures x is in the range  $[0, \text{TWO\_PI})$  as ;
71         //  $\cos(x + \text{TWO\_PI}) = \cos(x)$ 
72         while(x >= TWO_PI)
73             x -= TWO_PI;
74
75         //Calculates the correct modification of x to accurately
76         // calculate  $\cos(x)$  when it is reduced to the range  $[0, \text{HALF\_PI})$ 
77         if(x >= PI)
78             if(x - PI >= HALF_PI)
79                 return geometric_cos_bounded(TWO_PI - x, n);
80             else
81                 return -1 * geometric_cos_bounded(x - PI, n);
82         else
83             if(x >= HALF_PI)
84                 return -1 * geometric_cos_bounded(PI - x, n);
85             else
86                 return geometric_cos_bounded(x, n);
87     }
88
89     // $\cos(x) = \cos(-x)$  in the second case
90     return geometric_cos(-x, n);
91 }
92
93 void mpfr_geometric_cos(mpfr_t R, mpfr_t x, unsigned int n)
94 {
95     mpfr_t y, t;
96     mpfr_init_set(y, x, MPFR_RNDN);
97     mpfr_init(t);
98
99     if(mpfr_cmp_ui(y, 0) >= 0)
100     {
101         while(mpfr_cmp(y, MPFR_TWO_PI) >= 0)
102             mpfr_sub(y, y, MPFR_TWO_PI, MPFR_RNDN);
103
104         if(mpfr_cmp(y, MPFR_PI) >= 0)

```



```

105     {
106         mpfr_sub(t, y, MPFR_PI, MPFR_RNDN);
107         if(mpfr_cmp(t, MPFR_HALF_PI) >= 0)
108         {
109             mpfr_sub(y, MPFR_TWO_PI, y, MPFR_RNDN);
110             mpfr_geometric_cos_bounded(R, y, n);
111         }
112         else
113         {
114             mpfr_sub(y, y, MPFR_PI, MPFR_RNDN);
115             mpfr_geometric_cos_bounded(R, y, n);
116             mpfr_neg(R, R, MPFR_RNDN);
117         }
118     }
119     else
120     {
121         if(mpfr_cmp(y, MPFR_HALF_PI) >= 0)
122         {
123             mpfr_sub(y, MPFR_PI, y, MPFR_RNDN);
124             mpfr_geometric_cos_bounded(R, y, n);
125             mpfr_neg(R, R, MPFR_RNDN);
126         }
127         else
128         {
129             mpfr_geometric_cos_bounded(R, y, n);
130         }
131     }
132 }
133 else
134 {
135     mpfr_neg(y, y, MPFR_RNDN);
136     mpfr_geometric_cos(R, y, n);
137 }
138 }
139
140 //sin(x) = cos(x - HALF_PI)
141 double geometric_sin(double x, unsigned int n)
142 {
143     return geometric_cos(x - HALF_PI, n);
144 }
145
146 void mpfr_geometric_sin(mpfr_t R, mpfr_t x, unsigned int n)
147 {
148     mpfr_t y;
149     mpfr_init(y);
150     mpfr_sub(y, x, MPFR_HALF_PI, MPFR_RNDN);
151     mpfr_geometric_cos(R, y, n);
152 }
153
154 //tan(x) = sin(x)/cos(x)
155 double geometric_tan(double x, unsigned int n)
156 {
157     return geometric_sin(x, n)/geometric_cos(x, n);
158 }
159
160 void mpfr_geometric_tan(mpfr_t R, mpfr_t x, unsigned int n)
161 {
162     mpfr_t S, C;

```

```

163
164     mpfr_init(S);
165     mpfr_init(C);
166     mpfr_geometric_sin(S, x, n);
167     mpfr_geometric_cos(C, x, n);
168
169     mpfr_div(R, S, C, MPFR_RNDN);
170     if(mpfr_cmp_ui(R, 1000000) > 0)
171         mpfr_set_inf(R, 1);
172     else if(mpfr_cmp_si(R, -1000000) < 0)
173         mpfr_set_inf(R, -1);
174 }
175
176 #ifdef COMPILE_MAIN
177 int main(int argc, char **argv)
178 {
179     double x, y;
180     unsigned int n, p, D;
181     mpfr_t R, X;
182     char sf[50];
183     FILE *in;
184
185     if(argc > 1)
186     {
187         switch(argv[1][0])
188         {
189             case 'a':
190                 if(argc == 5 &&
191                     sscanf(argv[2], "%lf", &x) == 1 &&
192                     sscanf(argv[3], "%u", &n) == 1 &&
193                     sscanf(argv[4], "%u", &D) == 1)
194                     printf("Cos(%.1f) = %.1f\n",
195                         d(D), x, D, geometric_cos(x, n));
196                 else
197                     printf("Usage: %s a <x=value for Cos(x)> <n> "
198                         "<D=Number of digits to display>\n",
199
200                         argv[0]);
201                 break;
202
203             case 'b':
204                 if(argc == 5 &&
205                     sscanf(argv[2], "%lf", &x) == 1 &&
206                     sscanf(argv[3], "%u", &n) == 1 &&
207                     sscanf(argv[4], "%u", &D) == 1)
208                     printf("Sin(%.1f) = %.1f\n",
209                         d(D), x, D, geometric_sin(x, n));
210                 else
211                     printf("Usage: %s b <x=value for Sin(x)> <n> "
212                         "<D=Number of digits to display>\n",
213
214                         argv[0]);
215                 break;
216
217             case 'c':
218                 if(argc == 5 &&
219                     sscanf(argv[2], "%lf", &x) == 1 &&
220                     sscanf(argv[3], "%u", &n) == 1 &&

```

```

221     sscanf(argv[4], "%u", &D) == 1)
222     printf("Tan(%.*lf) = %.*lf\n",
223           d(D), x, D, geometric_tan(x, n));
224 else
225     printf("Usage: %s a <x=value for Tan(x)> <n> "
226           "<D=Number of digits to display>\n",
227           argv[0]);
228
229 break;
230
231 case 'd':
232     if(argc == 6 &&
233         sscanf(argv[3], "%u", &D) == 1 &&
234         sscanf(argv[4], "%u", &n) == 1 &&
235         sscanf(argv[5], "%u", &p) == 1)
236     {
237         mpfr_set_default_prec(p);
238         INIT_CONSTANTS
239         if (mpfr_init_set_str(X, argv[2], 10, MPFR_RNDN) == 0)
240         {
241             mpfr_init(R);
242
243             sprintf(sf, "cos(%%.%%uRNf) =~\n\t%%.%%uRNf\n",
244                   d(D), D);
245
246             mpfr_geometric_cos(R, X, n);
247             mpfr_printf(sf, X, R);
248         }
249     else
250         printf("Usage: %s d <x=value for Cos(x)> "
251               "<D=Number of digits to display> "
252               "<n> <p=bits of precision to use>\n",
253               argv[0]);
254     }
255 else
256     printf("Usage: %s d <x=value for Cos(x)> "
257           "<D=Number of digits to display> "
258           "<n> <p=bits of precision to use>\n",
259           argv[0]);
260 break;
261
262 case 'e':
263     if(argc == 6 &&
264         sscanf(argv[3], "%u", &D) == 1 &&
265         sscanf(argv[4], "%u", &n) == 1 &&
266         sscanf(argv[5], "%u", &p) == 1)
267     {
268         mpfr_set_default_prec(p);
269         INIT_CONSTANTS
270         if (mpfr_init_set_str(X, argv[2], 10, MPFR_RNDN) == 0)
271         {
272             mpfr_init(R);
273
274             sprintf(sf, "sin(%%.%%uRNf) =~\n\t%%.%%uRNf\n",
275                   d(D), D);
276
277             mpfr_geometric_sin(R, X, n);
278             mpfr_printf(sf, X, R);

```

```

279     }
280     else
281         printf(" Usage: %s d <x=value for Sin(x)> "
282             "<D=Number of digits to display> "
283             "<n> <p=bits of precision to use>\n",
284             argv[0]);
285     }
286     else
287         printf(" Usage: %s d <x=value for Sin(x)> "
288             "<D=Number of digits to display> "
289             "<n> <p=bits of precision to use>\n",
290             argv[0]);
291     break;
292
293     case 'f':
294         if( argc == 6 &&
295             sscanf(argv[3], "%u", &D) == 1 &&
296             sscanf(argv[4], "%u", &n) == 1 &&
297             sscanf(argv[5], "%u", &p) == 1)
298         {
299             mpfr_set_default_prec(p);
300             INIT_CONSTANTS
301             if (mpfr_init_set_str(X, argv[2], 10, MPFR_RNDN) == 0)
302             {
303                 mpfr_init(R);
304
305                 sprintf(sf, " tan(%%.%uRNf) =~\n\t%%.%uRNf\n",
306                     d(D), D);
307
308                 mpfr_geometric_tan(R, X, n);
309                 mpfr_printf(sf, X, R);
310             }
311         else
312             printf(" Usage: %s d <x=value for Tan(x)> "
313                 "<D=Number of digits to display> "
314                 "<n> <p=bits of precision to use>\n",
315                 argv[0]);
316     }
317     else
318         printf(" Usage: %s d <x=value for Tan(x)> "
319             "<D=Number of digits to display> "
320             "<n> <p=bits of precision to use>\n",
321             argv[0]);
322     break;
323
324     default:
325         printf(" Usage: %s <a/b/c/d/e/f> <arguments>\n", argv[0]);
326     }
327 }
328 else
329     printf(" Usage: %s <a/b/c/d/e/f> <arguments>\n", argv[0]);
330 }
331 #endif

```

Code for Geometric Inverse Trigonometric Functions:

File : geometric_inv_trig.c

```
1 | #include <assert.h>
```

```

2 #include <stdio.h>
3 #include <math.h>
4 #include <gmp.h>
5 #include <mpfr.h>
6
7 #include "geometric_inv_trig.h"
8 #include "trig_utilities.h"
9 #include "utilities.h"
10
11 #define INIT_CONSTANTS in = fopen(PI_INFILE, "r"); \
12     mpfr_init(MPFR_PI); \
13     mpfr_inp_str(MPFR_PI, in, 10, MPFR_RNDN); \
14     fclose(in); \
15     mpfr_init(MPFR_HALF_PI); \
16     mpfr_div_ui(MPFR_HALF_PI, MPFR_PI, 2, MPFR_RNDN);
17
18 mpfr_t MPFR_PI, MPFR_HALF_PI;
19
20 double geometric_acos_bounded(double x, unsigned int n)
21 {
22     //Ensures the given value is a valid cosine value
23     assert(x >= 0 && x <= 1);
24
25     //Reversing the last line of geometric_cos_bounded
26     double h = 2-2*x;
27
28     //Reverses the iterative process oof geometric_cos_bounded
29     for(int i = 0; i < n; i++)
30         h = 2 - sqrt(4 - h);
31
32     //Reverses the initialisation proceduce in geometric_cos_bounded
33     h *= pow(4, n);
34     return sqrt(h);
35 }
36
37 void mpfr_geometric_acos_bounded(mpfr_t R, mpfr_t x, unsigned int n)
38 {
39     assert(mpfr_cmp_ui(x, 0) >= 0 && mpfr_cmp_ui(x, 1) <= 0);
40
41     mpfr_t h;
42
43     mpfr_init(h);
44     mpfr_ui_sub(h, 1, x, MPFR_RNDN);
45     mpfr_mul_ui(h, h, 2, MPFR_RNDN);
46
47     for(int i = 0; i < n; i++)
48     {
49         mpfr_ui_sub(h, 4, h, MPFR_RNDN);
50         mpfr_sqrt(h, h, MPFR_RNDN);
51         mpfr_ui_sub(h, 2, h, MPFR_RNDN);
52     }
53
54     mpfr_ui_pow_ui(R, 4, n, MPFR_RNDN);
55     mpfr_mul(R, h, R, MPFR_RNDN);
56     mpfr_sqrt(R, R, MPFR_RNDN);
57 }
58
59 double geometric_acos(double x, unsigned int n)

```

```

60 {
61     assert(x >= -1 && x <= 1);
62     return x >= 0 ? geometric_acos_bounded(x,n)
63         : PI - geometric_acos_bounded(-x,n);
64 }
65
66 void mpfr_geometric_acos(mpfr_t R, mpfr_t x, unsigned int n)
67 {
68     assert(mpfr_cmp_si(x, -1) >= 0 && mpfr_cmp_si(x, 1) <= 0);
69     mpfr_t y;
70
71     if(mpfr_cmp_ui(x, 0) < 0)
72     {
73         mpfr_init_set(y, x, MPFR_RNDN);
74         mpfr_neg(y, x, MPFR_RNDN);
75         mpfr_geometric_acos_bounded(R,y,n);
76         mpfr_sub(R, MPFR_PI, R, MPFR_RNDN);
77     }
78     else
79         mpfr_geometric_acos_bounded(R,x,n);
80 }
81
82 double geometric_asin(double x, unsigned int n)
83 {
84     assert(x >= -1 && x <= 1);
85     return HALF_PI - geometric_acos(x, n);
86 }
87
88 void mpfr_geometric_asin(mpfr_t R, mpfr_t x, unsigned int n)
89 {
90     assert(mpfr_cmp_si(x, -1) >= 0 && mpfr_cmp_si(x, 1) <= 0);
91
92     mpfr_geometric_acos(R, x, n);
93     mpfr_sub(R, MPFR_HALF_PI, R, MPFR_RNDN);
94 }
95
96 double geometric_atan(double x, unsigned int n)
97 {
98     return geometric_asin(x/sqrt(x*x + 1), n);
99 }
100
101 void mpfr_geometric_atan(mpfr_t R, mpfr_t x, unsigned int n)
102 {
103     mpfr_t y;
104
105     mpfr_init(y);
106     mpfr_mul(y, x, x, MPFR_RNDN);
107     mpfr_add_ui(y, y, 1, MPFR_RNDN);
108     mpfr_sqrt(y, y, MPFR_RNDN);
109     mpfr_div(y, x, y, MPFR_RNDN);
110
111     mpfr_geometric_asin(R, y, n);
112 }
113
114 #ifdef COMPILE_MAIN
115 int main(int argc, char **argv)
116 {
117     double x, y;

```

```

118 unsigned int n, p, D;
119 mpfr_t R, X;
120 char sf[50];
121 FILE *in;
122
123 if(argc > 1)
124 {
125     switch(argv[1][0])
126     {
127         case 'a':
128             if(argc == 5 &&
129                 sscanf(argv[2], "%lf", &x) == 1 &&
130                 sscanf(argv[3], "%u", &n) == 1 &&
131                 sscanf(argv[4], "%u", &D) == 1)
132                 printf("arcCos(%.*lf) = %.*lf\n",
133                     d(D), x, D, geometric_acos(x, n));
134             else
135                 printf("Usage: %s a <x=value for arcCos(x)> <n> "
136                     "<D=Number of digits to display>\n",
137                     argv[0]);
138             break;
139
140         case 'b':
141             if(argc == 5 &&
142                 sscanf(argv[2], "%lf", &x) == 1 &&
143                 sscanf(argv[3], "%u", &n) == 1 &&
144                 sscanf(argv[4], "%u", &D) == 1)
145                 printf("arcSin(%.*lf) = %.*lf\n",
146                     d(D), x, D, geometric_asin(x, n));
147             else
148                 printf("Usage: %s b <x=value for arcSin(x)> <n> "
149                     "<D=Number of digits to display>\n",
150                     argv[0]);
151             break;
152
153         case 'c':
154             if(argc == 5 &&
155                 sscanf(argv[2], "%lf", &x) == 1 &&
156                 sscanf(argv[3], "%u", &n) == 1 &&
157                 sscanf(argv[4], "%u", &D) == 1)
158                 printf("Tan(%.*lf) = %.*lf\n",
159                     d(D), x, D, geometric_atan(x, n));
160             else
161                 printf("Usage: %s a <x=value for arcTan(x)> <n> "
162                     "<D=Number of digits to display>\n",
163                     argv[0]);
164             break;
165
166         case 'd':
167             if(argc == 6 &&
168                 sscanf(argv[3], "%u", &D) == 1 &&
169                 sscanf(argv[4], "%u", &n) == 1 &&
170                 sscanf(argv[5], "%u", &p) == 1)
171             {
172                 mpfr_set_default_prec(p);
173                 INIT_CONSTANTS
174                 if (mpfr_init_set_str(X, argv[2], 10, MPFR_RNDN) == 0)
175

```

```

176         mpfr_init(R);
177
178         sprintf(sf, " arcCos(%%.%uRNf) =~\n\t%%%.%uRNf\n",
179             d(D), D);
180
181         mpfr_geometric_acos(R, X, n);
182         mpfr_printf(sf, X, R);
183     }
184     else
185         printf(" Usage: %s d <x=value for arcCos(x)> "
186             "<D=Number of digits to display> "
187             "<n> <p=bits of precision to use>\n",
188             argv[0]);
189 }
190 else
191     printf(" Usage: %s d <x=value for arcCos(x)> "
192         "<D=Number of digits to display> "
193         "<n> <p=bits of precision to use>\n",
194         argv[0]);
195 break;
196
197 case 'e':
198     if(argc == 6 &&
199         sscanf(argv[3], "%u", &D) == 1 &&
200         sscanf(argv[4], "%u", &n) == 1 &&
201         sscanf(argv[5], "%u", &p) == 1)
202     {
203         mpfr_set_default_prec(p);
204         INIT_CONSTANTS
205         if (mpfr_init_set_str(X, argv[2], 10, MPFR_RNDN) == 0)
206         {
207             mpfr_init(R);
208
209             sprintf(sf, " arcSin(%%.%uRNf) =~\n\t%%%.%uRNf\n",
210                 d(D), D);
211
212             mpfr_geometric_asin(R, X, n);
213             mpfr_printf(sf, X, R);
214         }
215     else
216         printf(" Usage: %s d <x=value for arcSin(x)> "
217             "<D=Number of digits to display> "
218             "<n> <p=bits of precision to use>\n",
219             argv[0]);
220 }
221 else
222     printf(" Usage: %s d <x=value for arcSin(x)> "
223         "<D=Number of digits to display> "
224         "<n> <p=bits of precision to use>\n",
225         argv[0]);
226 break;
227
228 case 'f':
229     if(argc == 6 &&
230         sscanf(argv[3], "%u", &D) == 1 &&
231         sscanf(argv[4], "%u", &n) == 1 &&
232         sscanf(argv[5], "%u", &p) == 1)
233     {

```



```

234     mpfr_set_default_prec(p);
235     INIT_CONSTANTS
236     if (mpfr_init_set_str(X, argv[2], 10, MPFR_RNDN) == 0)
237     {
238         mpfr_init(R);
239
240         sprintf(sf, " arcTan(%%.%%uRNf) =~\n\t%.%%uRNf\n",
241             d(D), D);
242
243         mpfr_geometric_atan(R, X, n);
244         mpfr_printf(sf, X, R);
245     }
246     else
247         printf(" Usage: %s d <x=value for arcTan(x)> "
248             "<D=Number of digits to display> "
249             "<n> <p=bits of precision to use>\n",
250             argv[0]);
251 }
252 else
253     printf(" Usage: %s d <x=value for arcTan(x)> "
254         "<D=Number of digits to display> "
255         "<n> <p=bits of precision to use>\n",
256         argv[0]);
257 break;
258
259 default:
260     printf(" Usage: %s <a/b/c/d/e/f> <arguments>\n", argv[0]);
261 }
262 }
263 else
264     printf(" Usage: %s <a/b/c/d/e/f> <arguments>\n", argv[0]);
265 }
266 #endif

```

Code for Taylor Trigonometric Functions:

File : taylor_trig.c

```

1  #include <stdio.h>
2  #include <assert.h>
3  #include <gmp.h>
4  #include <mpfr.h>
5
6  #include "taylor_trig.h"
7  #include "trig_utilities.h"
8  #include "utilities.h"
9
10 #define INIT_CONSTANTS in = fopen(PI_INFILE, "r"); \
11     mpfr_init(MPFR_PI); \
12     mpfr_inp_str(MPFR_PI, in, 10, MPFR_RNDN); \
13     fclose(in); \
14     mpfr_init(MPFR_TWO_PI); \
15     mpfr_init(MPFR_HALF_PI); \
16     mpfr_div_ui(MPFR_HALF_PI, MPFR_PI, 2, MPFR_RNDN); \
17     mpfr_mul_ui(MPFR_TWO_PI, MPFR_PI, 2, MPFR_RNDN);
18
19 mpfr_t MPFR_PI, MPFR_HALF_PI, MPFR_TWO_PI;
20
21 double taylor_cos_bounded(double x, unsigned int N)

```

```

22 {
23     assert(x >= 0 && x <= HALF_PI);
24     double c = 1, x_2 = x*x, a = 1, b = 1;
25     for(int n = 1; n < N; n++)
26     {
27         a /= (2*n - 1)*(2*(n++));
28         b *= x_2;
29         c -= a*b;
30         a /= (2*n - 1)*(2*(n));
31         b *= x_2;
32         c += a*b;
33     }
34     return c;
35 }
36
37 void mpfr_taylor_cos_bounded(mpfr_t R, mpfr_t x, unsigned int N)
38 {
39     assert(mpfr_cmp_ui(x, 0) >= 0 && mpfr_cmp(x, MPFR_HALF_PI) <= 0);
40     mpfr_t t, x_2;
41
42     mpfr_init_set_ui(R, 1, MPFR_RNDN);
43     mpfr_init_set_ui(t, 1, MPFR_RNDN);
44     mpfr_init(x_2);
45     mpfr_mul(x_2, x, x, MPFR_RNDN);
46
47     for(int n = 1; n < N; n++)
48     {
49         mpfr_div_ui(t, t, (2*n-1)*(2*(n++)), MPFR_RNDN);
50         mpfr_mul(t, t, x_2, MPFR_RNDN);
51         mpfr_sub(R, R, t, MPFR_RNDN);
52         mpfr_div_ui(t, t, (2*n-1)*(2*n), MPFR_RNDN);
53         mpfr_mul(t, t, x_2, MPFR_RNDN);
54         mpfr_add(R, R, t, MPFR_RNDN);
55     }
56 }
57
58 double taylor_sin_bounded(double x, unsigned int N)
59 {
60     assert(x >= 0 && x <= HALF_PI);
61     double s = x, x_2 = x*x, a = 1, b = x;
62     for(int n = 1; n < N; n++)
63     {
64         a /= (2*n + 1)*(2*(n++));
65         b *= x_2;
66         s -= a*b;
67         a /= (2*n + 1)*(2*n);
68         b *= x_2;
69         s += a*b;
70     }
71     return s;
72 }
73
74 void mpfr_taylor_sin_bounded(mpfr_t R, mpfr_t x, unsigned int N)
75 {
76     mpfr_printf("%.20RF\n", x);
77     assert(mpfr_cmp_ui(x, 0) >= 0 && mpfr_cmp(x, MPFR_HALF_PI) <= 0);
78     mpfr_t t, x_2;
79

```

```

80 mpfr_init_set(R, x, MPFR_RNDN);
81 mpfr_init_set(t, x, MPFR_RNDN);
82 mpfr_init(x_2);
83 mpfr_mul(x_2, x, x, MPFR_RNDN);
84
85 for(int n = 1; n < N; n++)
86 {
87     mpfr_div_ui(t, t, (2*n+1)*(2*(n++)), MPFR_RNDN);
88     mpfr_mul(t, t, x_2, MPFR_RNDN);
89     mpfr_sub(R, R, t, MPFR_RNDN);
90     mpfr_div_ui(t, t, (2*n+1)*(2*n), MPFR_RNDN);
91     mpfr_mul(t, t, x_2, MPFR_RNDN);
92     mpfr_add(R, R, t, MPFR_RNDN);
93 }
94 }
95
96 double taylor_cos(double x, unsigned int N)
97 {
98     if(x >= 0)
99     {
100         while(x >= TWO_PI)
101             x -= TWO_PI;
102
103         if(x >= PI)
104             if(x - PI >= HALF_PI)
105                 return taylor_cos_bounded(TWO_PI - x, N);
106             else
107                 return -1 * taylor_cos_bounded(x - PI, N);
108         else
109             if(x >= HALF_PI)
110                 return -1 * taylor_cos_bounded(PI - x, N);
111             else
112                 return taylor_cos_bounded(x, N);
113     }
114
115     return taylor_cos(-x, N);
116 }
117
118 void mpfr_taylor_cos(mpfr_t R, mpfr_t x, unsigned int N)
119 {
120     mpfr_t y, t;
121     mpfr_init_set(y, x, MPFR_RNDN);
122     mpfr_init(t);
123     if(mpfr_cmp_ui(y, 0) >= 0)
124     {
125         while(mpfr_cmp(y, MPFR_TWO_PI) >= 0)
126             mpfr_sub(y, y, MPFR_TWO_PI, MPFR_RNDN);
127
128         if(mpfr_cmp(y, MPFR_PI) >= 0)
129         {
130             mpfr_sub(t, y, MPFR_PI, MPFR_RNDN);
131             if(mpfr_cmp(t, MPFR_HALF_PI) >= 0)
132             {
133                 mpfr_sub(y, MPFR_TWO_PI, y, MPFR_RNDN);
134                 mpfr_taylor_cos_bounded(R, y, N);
135             }
136             else
137             {

```

```

138     mpfr_sub(y, y, MPFR_PI, MPFR_RNDN);
139     mpfr_taylor_cos_bounded(R, y, N);
140     mpfr_neg(R, R, MPFR_RNDN);
141 }
142 }
143 else
144 {
145     if(mpfr_cmp(y, MPFR_HALF_PI) >= 0)
146     {
147         mpfr_sub(y, MPFR_PI, y, MPFR_RNDN);
148         mpfr_taylor_cos_bounded(R, y, N);
149         mpfr_neg(R, R, MPFR_RNDN);
150     }
151     else
152     {
153         mpfr_taylor_cos_bounded(R, y, N);
154     }
155 }
156 }
157 else
158 {
159     mpfr_neg(y, y, MPFR_RNDN);
160     mpfr_taylor_cos_bounded(R, y, N);
161 }
162 }
163
164 double taylor_sin(double x, unsigned int N)
165 {
166     if(x >= 0)
167     {
168         while(x >= TWO_PI)
169             x -= TWO_PI;
170
171         if(x >= PI)
172             if(x - PI >= HALF_PI)
173                 return -1 * taylor_sin_bounded(TWO_PI - x, N);
174             else
175                 return -1 * taylor_sin_bounded(x - PI, N);
176         else
177             if(x >= HALF_PI)
178                 return taylor_sin_bounded(PI - x, N);
179             else
180                 return taylor_sin_bounded(x, N);
181     }
182
183     return -1 * taylor_sin(-x, N);
184 }
185
186 void mpfr_taylor_sin(mpfr_t R, mpfr_t x, unsigned int N)
187 {
188     mpfr_t y, t;
189     mpfr_init_set(y, x, MPFR_RNDN);
190     mpfr_init(t);
191
192     if(mpfr_cmp_ui(y, 0) >= 0)
193     {
194         while(mpfr_cmp(y, MPFR_TWO_PI) >= 0)
195             mpfr_sub(y, y, MPFR_TWO_PI, MPFR_RNDN);

```

```

196
197     if(mpfr_cmp(y, MPFR_PI) >= 0)
198     {
199         mpfr_sub(t, y, MPFR_PI, MPFR_RNDN);
200         if(mpfr_cmp(t, MPFR_PI) >= 0)
201         {
202             mpfr_sub(y, MPFR_TWO_PI, y, MPFR_RNDN);
203             mpfr_taylor_sin_bounded(R, y, N);
204             mpfr_neg(R, R, MPFR_RNDN);
205         }
206         else
207         {
208             mpfr_sub(y, y, MPFR_PI, MPFR_RNDN);
209             mpfr_taylor_sin_bounded(R, y, N);
210             mpfr_neg(R, R, MPFR_RNDN);
211         }
212     }
213     else
214     {
215         if(mpfr_cmp(y, MPFR_HALF_PI) >= 0)
216         {
217             mpfr_sub(y, MPFR_PI, y, MPFR_RNDN);
218             mpfr_taylor_sin_bounded(R, y, N);
219         }
220         else
221         {
222             mpfr_taylor_sin_bounded(R, y, N);
223         }
224     }
225 }
226 else
227 {
228     mpfr_neg(y, y, MPFR_RNDN);
229     mpfr_taylor_sin(R, y, N);
230     mpfr_neg(R, R, MPFR_RNDN);
231 }
232 }
233
234
235 double taylor_tan(double x, unsigned int N)
236 {
237     return taylor_sin(x,N)/taylor_cos(x,N);
238 }
239
240 void mpfr_taylor_tan(mpfr_t R, mpfr_t x, unsigned int N)
241 {
242     mpfr_t S, C;
243     mpfr_init(S);
244     mpfr_init(C);
245     mpfr_taylor_sin(S, x, N);
246     mpfr_taylor_cos(C, x, N);
247     mpfr_div(R, S, C, N);
248 }
249
250 #ifdef COMPILE_MAIN
251 int main(int argc, char **argv)
252 {
253     double x, y;

```

```

254 unsigned int n, p, D;
255 mpfr_t R, X;
256 char sf[50];
257 FILE *in;
258
259 if(argc > 1)
260 {
261     switch(argv[1][0])
262     {
263         case 'a':
264             if(argc == 5 &&
265                 sscanf(argv[2], "%lf", &x) == 1 &&
266                 sscanf(argv[3], "%u", &n) == 1 &&
267                 sscanf(argv[4], "%u", &D) == 1)
268                 printf("Cos(%.1f) = %.1f\n",
269                     d(D), x, D, taylor_cos(x, n));
270             else
271                 printf("Usage: %s a <x=value for Cos(x)> <n> "
272                     "<D=Number of digits to display>\n",
273                     argv[0]);
274             break;
275
276         case 'b':
277             if(argc == 5 &&
278                 sscanf(argv[2], "%lf", &x) == 1 &&
279                 sscanf(argv[3], "%u", &n) == 1 &&
280                 sscanf(argv[4], "%u", &D) == 1)
281                 printf("Sin(%.1f) = %.1f\n",
282                     d(D), x, D, taylor_sin(x, n));
283             else
284                 printf("Usage: %s b <x=value for Sin(x)> <n> "
285                     "<D=Number of digits to display>\n",
286                     argv[0]);
287             break;
288
289         case 'c':
290             if(argc == 5 &&
291                 sscanf(argv[2], "%lf", &x) == 1 &&
292                 sscanf(argv[3], "%u", &n) == 1 &&
293                 sscanf(argv[4], "%u", &D) == 1)
294                 printf("Tan(%.1f) = %.1f\n",
295                     d(D), x, D, taylor_tan(x, n));
296             else
297                 printf("Usage: %s a <x=value for Tan(x)> <n> "
298                     "<D=Number of digits to display>\n",
299                     argv[0]);
300             break;
301
302         case 'd':
303             if(argc == 6 &&
304                 sscanf(argv[3], "%u", &D) == 1 &&
305                 sscanf(argv[4], "%u", &n) == 1 &&
306                 sscanf(argv[5], "%u", &p) == 1)
307             {
308                 mpfr_set_default_prec(p);
309             }
310     }
311 }

```

```

312     INIT_CONSTANTS
313     if (mpfr_init_set_str(X, argv[2], 10, MPFR_RNDN) == 0)
314     {
315         mpfr_init(R);
316
317         sprintf(sf, "cos(%%.%%uRNf) =~\n\tt%%.%%uRNf\n",
318             d(D), D);
319
320         mpfr_taylor_cos(R, X, n);
321         mpfr_printf(sf, X, R);
322     }
323     else
324         printf("Usage: %s d <x=value for Cos(x)> "
325             "<D=Number of digits to display> "
326             "<n> <p=bits of precision to use>\n",
327             argv[0]);
328 }
329 else
330     printf("Usage: %s d <x=value for Cos(x)> "
331         "<D=Number of digits to display> "
332         "<n> <p=bits of precision to use>\n",
333         argv[0]);
334 break;
335
336 case 'e':
337     if (argc == 6 &&
338         sscanf(argv[3], "%u", &D) == 1 &&
339         sscanf(argv[4], "%u", &n) == 1 &&
340         sscanf(argv[5], "%u", &p) == 1)
341     {
342         mpfr_set_default_prec(p);
343         INIT_CONSTANTS
344         if (mpfr_init_set_str(X, argv[2], 10, MPFR_RNDN) == 0)
345         {
346             mpfr_init(R);
347
348             sprintf(sf, "sin(%%.%%uRNf) =~\n\tt%%.%%uRNf\n",
349                 d(D), D);
350
351             mpfr_taylor_sin(R, X, n);
352             mpfr_printf(sf, X, R);
353         }
354         else
355             printf("Usage: %s d <x=value for Sin(x)> "
356                 "<D=Number of digits to display> "
357                 "<n> <p=bits of precision to use>\n",
358                 argv[0]);
359     }
360     else
361         printf("Usage: %s d <x=value for Sin(x)> "
362             "<D=Number of digits to display> "
363             "<n> <p=bits of precision to use>\n",
364             argv[0]);
365 break;
366
367 case 'f':
368     if (argc == 6 &&
369         sscanf(argv[3], "%u", &D) == 1 &&

```

```

370     sscanf(argv[4], "%u", &n) == 1 &&
371     sscanf(argv[5], "%u", &p) == 1)
372 {
373     mpfr_set_default_prec(p);
374     INIT_CONSTANTS
375     if (mpfr_init_set_str(X, argv[2], 10, MPFR_RNDN) == 0)
376     {
377         mpfr_init(R);
378
379         sprintf(sf, "tan(%%.%uRNf) = ~\n\t%%.%uRNf\n",
380             d(D), D);
381
382         mpfr_taylor_tan(R, X, n);
383         mpfr_printf(sf, X, R);
384     }
385     else
386         printf("Usage: %s d <x=value for Tan(x)> "
387             "<D=Number of digits to display> "
388             "<n> <p=bits of precision to use>\n",
389             argv[0]);
390 }
391 else
392     printf("Usage: %s d <x=value for Tan(x)> "
393         "<D=Number of digits to display> "
394         "<n> <p=bits of precision to use>\n",
395         argv[0]);
396 break;
397
398 default:
399     printf("Usage: %s <a/b/c/d/e/f> <arguments>\n", argv[0]);
400 }
401 }
402 else
403     printf("Usage: %s <a/b/c/d/e/f> <arguments>\n", argv[0]);
404 }
405 #endif

```

Code for Taylor Inverse Trigonometric Functions:

File : taylor_inv_trig.c

```

1  #include <assert.h>
2  #include <stdio.h>
3  #include <math.h>
4  #include <gmp.h>
5  #include <mpfr.h>
6
7  #include "taylor_inv_trig.h"
8  #include "trig_utilities.h"
9  #include "utilities.h"
10
11 #define INIT_CONSTANTS in = fopen(PI_INFILE, "r"); \
12     mpfr_init(MPFR_PI); \
13     mpfr_inp_str(MPFR_PI, in, 10, MPFR_RNDN); \
14     fclose(in); \
15     mpfr_init(MPFR_HALF_PI); \
16     mpfr_div_ui(MPFR_HALF_PI, MPFR_PI, 2, MPFR_RNDN);
17
18 mpfr_t MPFR_PI, MPFR_HALF_PI;

```



```

19
20 double taylor_asin(double x, unsigned int N)
21 {
22     assert(x >= -1 && x <= 1);
23
24     if(x == 1)
25         return HALF_PI;
26     else if(x == -1)
27         return -1 * HALF_PI;
28
29     double s = x, x_2 = x*x, t = x;
30
31     for(int n = 1; n < N; n++)
32     {
33         t *= 2*n*(2*n - 1)*x_2;
34         t /= 4*n*n;
35         s += t/(2*n+1);
36     }
37     return s;
38 }
39
40 void mpfr_taylor_asin(mpfr_t R, mpfr_t x, unsigned int N)
41 {
42     assert(mpfr_cmp_si(x, -1) >= 0 && mpfr_cmp_si(x, 1) <= 0);
43
44     if(mpfr_cmp_si(x, 1) == 0)
45         mpfr_set(R, MPFR_HALF_PI, MPFR_RNDN);
46     else if(mpfr_cmp_si(x, -1) == 0)
47     {
48         mpfr_set(R, MPFR_HALF_PI, MPFR_RNDN);
49         mpfr_neg(R, R, MPFR_RNDN);
50     }
51     else
52     {
53         mpfr_t x_2, t, T;
54         mpfr_init(x_2);
55         mpfr_mul(x_2, x, x, MPFR_RNDN);
56         mpfr_init_set(t, x, MPFR_RNDN);
57         mpfr_init(T);
58         mpfr_set(R, x, MPFR_RNDN);
59         for(int n = 1; n < N; n++)
60         {
61             mpfr_mul_ui(t, t, 2*n*(2*n - 1), MPFR_RNDN);
62             mpfr_mul(t, t, x_2, MPFR_RNDN);
63             mpfr_div_ui(t, t, 4*n*n, MPFR_RNDN);
64             mpfr_div_ui(T, t, 2*n + 1, MPFR_RNDN);
65             mpfr_add(R, R, T, MPFR_RNDN);
66         }
67     }
68 }
69
70 double taylor_acos(double x, unsigned int N)
71 {
72     return HALF_PI - taylor_asin(x, N);
73 }
74
75 void mpfr_taylor_acos(mpfr_t R, mpfr_t x, unsigned int N)
76 {

```

```

77     mpfr_taylor_asin(R, x, N);
78     mpfr_sub(R, MPFR_HALF_PI, R, MPFR_RNDN);
79 }
80
81 double taylor_atan_bounded(double x, unsigned int N)
82 {
83     assert(x >= 0 && x <= 1);
84     double t = 0, x_2 = x*x, y = x;
85     for(int n = 0; n < N; n++)
86     {
87         t += y/(2*(n++) + 1);
88         y *= x_2;
89         t -= y/(2*n + 1);
90         y *= x_2;
91     }
92     return t;
93 }
94
95 void mpfr_taylor_atan_bounded(mpfr_t R, mpfr_t x, unsigned int N)
96 {
97     assert(mpfr_cmp_ui(x, 0) >= 0 && mpfr_cmp_ui(x, 1) <= 1);
98     mpfr_t x_2, y, a;
99     mpfr_init_set(y, x, MPFR_RNDN);
100    mpfr_init_set_ui(R, 0, MPFR_RNDN);
101    mpfr_init(x_2);
102    mpfr_mul(x_2, x, x, MPFR_RNDN);
103    mpfr_init(a);
104    for(int n = 0; n < N; n++)
105    {
106        mpfr_div_ui(a, y, 2*(n++) + 1, MPFR_RNDN);
107        mpfr_add(R, R, a, MPFR_RNDN);
108        mpfr_mul(y, y, x_2, MPFR_RNDN);
109        mpfr_div_ui(a, y, 2*n + 1, MPFR_RNDN);
110        mpfr_sub(R, R, a, MPFR_RNDN);
111        mpfr_mul(y, y, x_2, MPFR_RNDN);
112    }
113 }
114
115 double taylor_atan(double x, unsigned int N)
116 {
117     if(x < 0)
118         return -taylor_atan(-x, N);
119
120     if(x >= 1)
121         return HALF_PI/2 + taylor_atan_bounded((x - 1)/(x + 1), N);
122
123     return taylor_atan_bounded(x, N);
124 }
125
126 void mpfr_taylor_atan(mpfr_t R, mpfr_t x, unsigned int N)
127 {
128     mpfr_t y, pi_4, z;
129     mpfr_init(y);
130     mpfr_init(pi_4);
131     mpfr_init(z);
132
133     if(mpfr_cmp_ui(x, 0) < 0)
134     {

```

```

135     mpfr_neg(y, x, MPFR_RNDN);
136     mpfr_taylor_atan(R, y, N);
137     mpfr_neg(R, R, MPFR_RNDN);
138 }
139 else if (mpfr_cmp_ui(x, 1) >= 0)
140 {
141     mpfr_div_ui(pi_4, MPFR_HALF_PI, 2, MPFR_RNDN);
142     mpfr_add_ui(z, x, 1, MPFR_RNDN);
143     mpfr_sub_ui(y, x, 1, MPFR_RNDN);
144     mpfr_div(y, y, z, MPFR_RNDN);
145     mpfr_taylor_atan_bounded(R, y, N);
146     mpfr_add(R, R, pi_4, MPFR_RNDN);
147 }
148 else
149 {
150     mpfr_taylor_atan_bounded(R, x, N);
151 }
152 }
153
154 #ifdef COMPILE_MAIN
155 int main(int argc, char **argv)
156 {
157     double x, y;
158     unsigned int n, p, D;
159     mpfr_t R, X;
160     char sf[50];
161     FILE *in;
162
163     if (argc > 1)
164     {
165         switch (argv[1][0])
166         {
167             case 'a':
168                 if (argc == 5 &&
169                     sscanf(argv[2], "%lf", &x) == 1 &&
170                     sscanf(argv[3], "%u", &n) == 1 &&
171                     sscanf(argv[4], "%u", &D) == 1)
172                     printf("arcCos(%.1f) = %.1f\n",
173                         d(D), x, D, taylor_acos(x, n));
174                 else
175                     printf("Usage: %s a <x=value for arcCos(x)> <n> "
176                         "<D=Number of digits to display>\n",
177                         argv[0]);
178                 break;
179
180             case 'b':
181                 if (argc == 5 &&
182                     sscanf(argv[2], "%lf", &x) == 1 &&
183                     sscanf(argv[3], "%u", &n) == 1 &&
184                     sscanf(argv[4], "%u", &D) == 1)
185                     printf("arcSin(%.1f) = %.1f\n",
186                         d(D), x, D, taylor_asin(x, n));
187                 else
188                     printf("Usage: %s b <x=value for arcSin(x)> <n> "
189                         "<D=Number of digits to display>\n",
190                         argv[0]);
191                 break;
192

```

```

193 case 'c':
194     if (argc == 5 &&
195         sscanf(argv[2], "%lf", &x) == 1 &&
196         sscanf(argv[3], "%u", &n) == 1 &&
197         sscanf(argv[4], "%u", &D) == 1)
198         printf("Tan(%.*lf) = %.*lf\n",
199             d(D), x, D, taylor_atan(x, n));
200     else
201         printf("Usage: %s a <x=value for arcTan(x)> <n> "
202             "<D=Number of digits to display>\n",
203             argv[0]);
204     break;
205
206 case 'd':
207     if (argc == 6 &&
208         sscanf(argv[3], "%u", &D) == 1 &&
209         sscanf(argv[4], "%u", &n) == 1 &&
210         sscanf(argv[5], "%u", &p) == 1)
211     {
212         mpfr_set_default_prec(p);
213         INIT_CONSTANTS
214         if (mpfr_init_set_str(X, argv[2], 10, MPFR_RNDN) == 0)
215         {
216             mpfr_init(R);
217
218             sprintf(sf, " arcCos(%%.%uRNf) =~\n\t%.%uRNf\n",
219                 d(D), D);
220
221             mpfr_taylor_acos(R, X, n);
222             mpfr_printf(sf, X, R);
223         }
224     else
225         printf("Usage: %s d <x=value for arcCos(x)> "
226             "<D=Number of digits to display> "
227             "<n> <p=bits of precision to use>\n",
228             argv[0]);
229     }
230     else
231         printf("Usage: %s d <x=value for arcCos(x)> "
232             "<D=Number of digits to display> "
233             "<n> <p=bits of precision to use>\n",
234             argv[0]);
235     break;
236
237 case 'e':
238     if (argc == 6 &&
239         sscanf(argv[3], "%u", &D) == 1 &&
240         sscanf(argv[4], "%u", &n) == 1 &&
241         sscanf(argv[5], "%u", &p) == 1)
242     {
243         mpfr_set_default_prec(p);
244         INIT_CONSTANTS
245         if (mpfr_init_set_str(X, argv[2], 10, MPFR_RNDN) == 0)
246         {
247             mpfr_init(R);
248
249             sprintf(sf, " arcSin(%%.%uRNf) =~\n\t%.%uRNf\n",
250                 d(D), D);

```

```

251     mpfr_taylor_asin(R, X, n);
252     mpfr_printf(sf, X, R);
253 }
254 else
255     printf(" Usage: %s d <x=value for arcSin(x)> "
256           "<D=Number of digits to display> "
257           "<n> <p=bits of precision to use>\n",
258           argv[0]);
259 }
260 else
261     printf(" Usage: %s d <x=value for arcSin(x)> "
262           "<D=Number of digits to display> "
263           "<n> <p=bits of precision to use>\n",
264           argv[0]);
265 break;
266
267 case 'f':
268     if(argc == 6 &&
269        sscanf(argv[3], "%u", &D) == 1 &&
270        sscanf(argv[4], "%u", &n) == 1 &&
271        sscanf(argv[5], "%u", &p) == 1)
272     {
273         mpfr_set_default_prec(p);
274         INIT_CONSTANTS
275         if (mpfr_init_set_str(X, argv[2], 10, MPFR_RNDN) == 0)
276         {
277             mpfr_init(R);
278
279             sprintf(sf, " arcTan(%u.%uRNf) =~\n\t%u.%uRNf\n",
280                    d(D), D);
281
282             mpfr_taylor_atan(R, X, n);
283             mpfr_printf(sf, X, R);
284         }
285     }
286     else
287         printf(" Usage: %s d <x=value for arcTan(x)> "
288               "<D=Number of digits to display> "
289               "<n> <p=bits of precision to use>\n",
290               argv[0]);
291 }
292 else
293     printf(" Usage: %s d <x=value for arcTan(x)> "
294           "<D=Number of digits to display> "
295           "<n> <p=bits of precision to use>\n",
296           argv[0]);
297 break;
298
299 default:
300     printf(" Usage: %s <a/b/c/d/e/f> <arguments>\n", argv[0]);
301 }
302 }
303 else
304     printf(" Usage: %s <a/b/c/d/e/f> <arguments>\n", argv[0]);
305 }
306 #endif

```

Code for CORDIC Functions:

```

1 #include <stdio.h>
2 #include <stdlib.h>
3 #include <math.h>
4 #include <assert.h>
5
6 #include "cordic_trig.h"
7 #include "utilities.h"
8
9 const cordic_fixed_t TRIG_ANGLES[] = FIXED_ANGLES;
10 const cordic_fixed_t TRIG_K_VALUES[] = FIXED_K_VALUES;
11
12 double *cordic_trig(const double theta, const unsigned int iter)
13 {
14     assert(-HALF_PI <= theta && theta <= HALF_PI);
15     unsigned int n = (iter > MAX_ITER || iter == 0) ? MAX_ITER : iter;
16
17     cordic_fixed_t x = TRIG_K_VALUES[n-1], y = 0, t;
18     cordic_fixed_t beta = double_to_fixed(theta);
19
20     double *result = malloc(2*sizeof(*result));
21
22     if(beta == 0)
23     {
24         x = FIXED_ONE;
25         beta = 0;
26     }
27     else if(beta == FIXED_HALF_PI)
28     {
29         x = 0;
30         y = FIXED_ONE;
31         beta = 0;
32     }
33     else if(beta == -FIXED_HALF_PI)
34     {
35         x = 0;
36         y = -FIXED_ONE;
37         beta = 0;
38     }
39
40     if(beta)
41     {
42         for(int i = 0; i < n; i++)
43         {
44             t = x;
45             if(beta >= 0)
46             {
47                 x = x - (y >> i);
48                 y = y + (t >> i);
49                 beta -= TRIG_ANGLES[i];
50             }
51             else
52             {
53                 x = x + (y >> i);
54                 y = y - (t >> i);
55                 beta += TRIG_ANGLES[i];
56             }
57         }
58     }
59 }

```

```

58     }
59
60     result[0] = fixed_to_double(x);
61     result[1] = fixed_to_double(y);
62     return result;
63 }
64
65 double cordic_atan_bounded(const double z, const unsigned int iter)
66 {
67     assert(0 <= z && z <= 1);
68     unsigned int n = (iter > MAX_ITER || iter == 0) ? MAX_ITER : iter;
69
70     cordic_fixed_t x = FIXED_ONE >> 1, y = double_to_fixed(z) >> 1, t;
71     cordic_fixed_t beta = 0;
72
73     if(y == 0)
74         return 0;
75
76     for(int i = 0; i < n; i++)
77     {
78         t = x;
79         if(y < 0)
80         {
81             x = x - (y >> i);
82             y = y + (t >> i);
83             beta -= TRIG_ANGLES[i];
84         }
85         else
86         {
87             x = x + (y >> i);
88             y = y - (t >> i);
89             beta += TRIG_ANGLES[i];
90         }
91     }
92
93     return fixed_to_double(beta);
94 }
95
96 double cordic_cos(double x, unsigned int n)
97 {
98     double *r;
99
100    if(x >= 0)
101    {
102        while(x >= TWO_PI)
103            x -= TWO_PI;
104
105        if(x >= PI)
106            if(x - PI >= HALF_PI)
107                r = cordic_trig(TWO_PI - x, n);
108            else
109            {
110                r = cordic_trig(x - PI, n);
111                r[0] = -r[0];
112            }
113        else
114            if(x >= HALF_PI)
115            {

```

```

116     r = cordic_trig(PI - x, n);
117     r[0] = -r[0];
118 }
119 else
120     r = cordic_trig(x, n);
121 return r[0];
122 }
123 return cordic_cos(-x, n);
124 }
125
126 double cordic_sin(double x, unsigned int n)
127 {
128     return cordic_cos(x - HALF_PI, n);
129 }
130
131 double cordic_tan(double x, unsigned int n)
132 {
133     double *r;
134     if(x >= 0)
135     {
136         while(x >= PI)
137             x -= PI;
138
139         if(x >= HALF_PI)
140         {
141             r = cordic_trig(PI - x, n);
142             return -1 * r[1]/r[0];
143         }
144
145         r = cordic_trig(x, n);
146         return r[1]/r[0];
147     }
148     return -1 * cordic_tan(-x, n);
149 }
150
151 double cordic_acos(double x, unsigned int n)
152 {
153     assert(-1 <= x && x <= 1);
154     return x == 0 ? HALF_PI
155            : x > 0 ? cordic_atan(sqrt(1 - x*x)/x, n)
156                   : HALF_PI + cordic_asin(-x, n);
157 }
158
159 double cordic_asin(double x, unsigned int n)
160 {
161     assert(-1 <= x && x <= 1);
162     return x == 1 ? HALF_PI
163            : x == -1 ? -HALF_PI
164                   : cordic_atan(x/sqrt(1 - x*x), n);
165 }
166
167 double cordic_atan(double x, unsigned int n)
168 {
169     if(x < 0)
170         return -cordic_atan(-x, n);
171
172     if(x >= 1)
173         return HALF_PI/2 + cordic_atan_bounded((x-1)/(x+1), n);

```



```

174
175     return cordic_atan_bounded(x, n);
176 }
177
178 #ifdef COMPILE_MAIN
179 int main(int argc, char **argv)
180 {
181     double x;
182     unsigned int n, D;
183
184     if(argc > 1)
185     {
186         switch(argv[1][0])
187         {
188             case 'a':
189                 if(argc == 5 &&
190                     sscanf(argv[2], "%lf", &x) == 1 &&
191                     sscanf(argv[3], "%u", &n) == 1 &&
192                     sscanf(argv[4], "%u", &D) == 1)
193                     printf("Cos(%.1f) = %.1f\n",
194                         d(D), x, D, cordic_cos(x, n));
195                 else
196                     printf("Uasge: %s a <x=value for Cos(x)> <n> "
197                         "<D=Number of digits to display>\n",
198                         argv[0]);
199                 break;
200
201             case 'b':
202                 if(argc == 5 &&
203                     sscanf(argv[2], "%lf", &x) == 1 &&
204                     sscanf(argv[3], "%u", &n) == 1 &&
205                     sscanf(argv[4], "%u", &D) == 1)
206                     printf("Sin(%.1f) = %.1f\n",
207                         d(D), x, D, cordic_sin(x, n));
208                 else
209                     printf("Uasge: %s a <x=value for Sin(x)> <n> "
210                         "<D=Number of digits to display>\n",
211                         argv[0]);
212                 break;
213
214             case 'c':
215                 if(argc == 5 &&
216                     sscanf(argv[2], "%lf", &x) == 1 &&
217                     sscanf(argv[3], "%u", &n) == 1 &&
218                     sscanf(argv[4], "%u", &D) == 1)
219                     printf("Tan(%.1f) = %.1f\n",
220                         d(D), x, D, cordic_tan(x, n));
221                 else
222                     printf("Uasge: %s a <x=value for Tan(x)> <n> "
223                         "<D=Number of digits to display>\n",
224                         argv[0]);
225                 break;
226
227             case 'd':
228                 if(argc == 5 &&
229                     sscanf(argv[2], "%lf", &x) == 1 &&
230                     sscanf(argv[3], "%u", &n) == 1 &&
231                     sscanf(argv[4], "%u", &D) == 1)

```

```

232     printf("aTan(%.1f) = %.1f\n",
233           d(D), x, D, cordic_atan(x, n));
234     else
235         printf("Usage: %s d <x=value for aTan(x)> <n> "
236               "<D=Number of digits to display>\n",
237               argv[0]);
238     break;
239
240     case 'e':
241         if(argc == 5 &&
242            sscanf(argv[2], "%.1f", &x) == 1 &&
243            sscanf(argv[3], "%u", &n) == 1 &&
244            sscanf(argv[4], "%u", &D) == 1)
245             printf("aCos(%.1f) = %.1f\n",
246                   d(D), x, D, cordic_acos(x, n));
247         else
248             printf("Usage: %s d <x=value for aCos(x)> <n> "
249                   "<D=Number of digits to display>\n",
250                   argv[0]);
251         break;
252
253     case 'f':
254         if(argc == 5 &&
255            sscanf(argv[2], "%.1f", &x) == 1 &&
256            sscanf(argv[3], "%u", &n) == 1 &&
257            sscanf(argv[4], "%u", &D) == 1)
258             printf("aSin(%.1f) = %.1f\n",
259                   d(D), x, D, cordic_asin(x, n));
260         else
261             printf("Usage: %s d <x=value for aSin(x)> <n> "
262                   "<D=Number of digits to display>\n",
263                   argv[0]);
264         break;
265
266     default:
267         printf("Usage: %s <a/b/c/d/e/f> <arguments>\n", argv[0]);
268 }
269 }
270 else
271     printf("Usage: %s <a/b/c/d/e/f> <arguments>\n", argv[0]);
272 }
273 #endif

```

Header files for Trigonometric Functions:

File : geometric_trig.h

```

1  #ifndef GEOMETRIC_TRIG_HEADER
2  #define GEOMETRIC_TRIG_HEADER
3
4  double geometric_cos_bounded(double, unsigned int);
5  void mpfr_geometric_cos_bounded(mpfr_t, mpfr_t, unsigned int);
6
7  double geometric_cos(double, unsigned int);
8  double geometric_sin(double, unsigned int);
9  double geometric_tan(double, unsigned int);
10
11 void mpfr_geometric_cos(mpfr_t, mpfr_t, unsigned int);
12 void mpfr_geometric_sin(mpfr_t, mpfr_t, unsigned int);

```

```

13 | void mpfr_geometric_tan(mpfr_t, mpfr_t, unsigned int);
14 |
15 | #endif

```

File : geometric_inv_trig.h

```

1 | #ifndef GEOMETRIC_INV_TRIG_HEADER
2 | #define GEOMETRIC_INV_TRIG_HEADER
3 |
4 | double geometric_acos_bounded(double, unsigned int);
5 | void mpfr_geometric_acos_bounded(mpfr_t, mpfr_t, unsigned int);
6 |
7 | double geometric_acos(double, unsigned int);
8 | double geometric_asin(double, unsigned int);
9 | double geometric_atan(double, unsigned int);
10 |
11 | void mpfr_geometric_acos(mpfr_t, mpfr_t, unsigned int);
12 | void mpfr_geometric_asin(mpfr_t, mpfr_t, unsigned int);
13 | void mpfr_geometric_atan(mpfr_t, mpfr_t, unsigned int);
14 |
15 | #endif

```

File : taylor_trig.h

```

1 | #ifndef TAYLOR_TRIG_HEADER
2 | #define TAYLOR_TRIG_HEADER
3 |
4 | double taylor_cos_bounded(double, unsigned int);
5 | double taylor_sin_bounded(double, unsigned int);
6 | double taylor_sin(double, unsigned int);
7 | double taylor_cos(double, unsigned int);
8 | double taylor_tan(double, unsigned int);
9 |
10 | void mpfr_taylor_cos_bounded(mpfr_t, mpfr_t, unsigned int);
11 | void mpfr_taylor_sin_bounded(mpfr_t, mpfr_t, unsigned int);
12 | void mpfr_taylor_sin(mpfr_t, mpfr_t, unsigned int);
13 | void mpfr_taylor_cos(mpfr_t, mpfr_t, unsigned int);
14 | void mpfr_taylor_tan(mpfr_t, mpfr_t, unsigned int);
15 |
16 | #endif

```

File : taylor_inv_trig.h

```

1 | #ifndef TAYLOR_INV_TRIG_HEADER
2 | #define TAYLOR_INV_TRIG_HEADER
3 |
4 | double taylor_asin(double, unsigned int);
5 | double taylor_acos(double, unsigned int);
6 | double taylor_atan_bounded(double, unsigned int);
7 | double taylor_atan(double, unsigned int);
8 |
9 | void mpfr_taylor_asin(mpfr_t, mpfr_t, unsigned int);
10 | void mpfr_taylor_acos(mpfr_t, mpfr_t, unsigned int);
11 | void mpfr_taylor_atan_bounded(mpfr_t, mpfr_t, unsigned int);
12 | void mpfr_taylor_atan(mpfr_t, mpfr_t, unsigned int);
13 | #endif

```

File : cordic_trig.h

```

1 #ifndef CORDIC_TRIG_HEADER
2 #define CORDIC_TRIG_HEADER
3
4 #include "trig_fixed.h"
5 #include "trig_utilities.h"
6
7 typedef TRIG_FIXED_TYPE cordic_fixed_t;
8
9 #if BITS == 64
10 #define FIXED_ONE 0x4000000000000000
11 #define NEG_CONSTANT 0x8000000000000000
12 #define FIXED_HALF_PI 0x6487ed5110b4611a
13 #define FIXED_ANGLES {0x3243f6a8885a308d, 0x1dac670561bb4f68, \
14     0x0fadba9c96406eb1, 0x07f56ea6ab0bdb71, \
15     0x03feab76e59fbd38, 0x01ffd55bba97624a, \
16     0x00fffaaaddb94d5, 0x007fff5556eeea5c, \
17     0x003fffeaaab7776e, 0x001ffffd5555bbbb, \
18     0x000fffffaaaaadd, 0x0007ffff555556e, \
19     0x0003fffffeaaaaab, 0x0001fffffd55555, \
20     0x0000ffffffffffaaaa, 0x00007fffffd5555, \
21     0x00003ffffffffffeaaa, 0x00001ffffffffffd55, \
22     0x00000fffffffffffaa, 0x000007ffffffffff5, \
23     0x000003fffffffffffe, 0x000001ffffffffff, \
24     0x0000010000000000, 0x0000008000000000, \
25     0x0000004000000000, 0x0000002000000000, \
26     0x0000001000000000, 0x0000000800000000, \
27     0x0000000400000000, 0x0000000200000000, \
28     0x0000000100000000, 0x0000000080000000, \
29     0x0000000040000000, 0x0000000020000000, \
30     0x0000000010000000, 0x0000000008000000, \
31     0x0000000004000000, 0x0000000002000000, \
32     0x0000000001000000, 0x0000000000800000, \
33     0x0000000000400000, 0x0000000000200000, \
34     0x0000000000100000, 0x0000000000080000, \
35     0x0000000000040000, 0x0000000000020000, \
36     0x0000000000010000, 0x0000000000008000, \
37     0x0000000000004000, 0x0000000000002000, \
38     0x0000000000001000, 0x0000000000000800, \
39     0x0000000000000400, 0x0000000000000200, \
40     0x0000000000000100, 0x0000000000000080, \
41     0x0000000000000040, 0x0000000000000020, \
42     0x0000000000000010, 0x0000000000000008, \
43     0x0000000000000004, 0x0000000000000002, \
44     0x0000000000000001}
45 #define FIXED_K_VALUES {0x2d413cccf779921, 0x287a26c490921db6, \
46     0x2744c374daf46d2f, 0x26f72283bd67fbda, \
47     0x26e3b58305ddeb19, 0x26ded9f57b2c3e7a, \
48     0x26dda30d3e4fd185, 0x26dd5552e1641def, \
49     0x26dd41e4454da117, 0x26dd3d089dfa47c8, \
50     0x26dd3bd1b42095ce, 0x26dd3b83f9a9db95, \
51     0x26dd3b708b0c282b, 0x26dd3b6baf64bb03, \
52     0x26dd3b6a787adfb4, 0x26dd3b6a2ac068e0, \
53     0x26dd3b6a1751cb2b, 0x26dd3b6a127623be, \
54     0x26dd3b6a113f39e3, 0x26dd3b6a10f17f6c, \
55     0x26dd3b6a10de10ce, 0x26dd3b6a10d93527, \
56     0x26dd3b6a10d7fe3d, 0x26dd3b6a10d7b082, \
57     0x26dd3b6a10d79d14, 0x26dd3b6a10d79838, \
58     0x26dd3b6a10d79701, 0x26dd3b6a10d796b3, \

```

```

59         0x26dd3b6a10d796a0, 0x26dd3b6a10d7969b, \
60         0x26dd3b6a10d7969a, 0x26dd3b6a10d7969a, \
61         0x26dd3b6a10d7969a, 0x26dd3b6a10d79699, \
62         0x26dd3b6a10d79699, 0x26dd3b6a10d79699, \
63         0x26dd3b6a10d79699, 0x26dd3b6a10d79699, \
64         0x26dd3b6a10d79699, 0x26dd3b6a10d79699, \
65         0x26dd3b6a10d79699, 0x26dd3b6a10d79699, \
66         0x26dd3b6a10d79699, 0x26dd3b6a10d79699, \
67         0x26dd3b6a10d79699, 0x26dd3b6a10d79699, \
68         0x26dd3b6a10d79699, 0x26dd3b6a10d79699, \
69         0x26dd3b6a10d79699, 0x26dd3b6a10d79699, \
70         0x26dd3b6a10d79699, 0x26dd3b6a10d79699, \
71         0x26dd3b6a10d79699, 0x26dd3b6a10d79699, \
72         0x26dd3b6a10d79699, 0x26dd3b6a10d79699, \
73         0x26dd3b6a10d79699, 0x26dd3b6a10d79699, \
74         0x26dd3b6a10d79699, 0x26dd3b6a10d79699, \
75         0x26dd3b6a10d79699, 0x26dd3b6a10d79699, \
76         0x26dd3b6a10d79699}
77     #define MAX_ITER 63
78 #elif BITS == 32
79     #define FIXED_ONE 0x40000000
80     #define NEG_CONSTANT 0x80000000
81     #define FIXED_HALF_PI 0x6487ed51
82     #define FIXED_ANGLES {0x3243f6a8, 0x1dac6705, 0x0fadbafc, \
83         0x07f56ea6, 0x03feab76, 0x01ffd55b, \
84         0x00fffaaa, 0x007fff55, 0x003fffea, \
85         0x001ffffd, 0x00100000, 0x00080000, \
86         0x00040000, 0x00020000, 0x00010000, \
87         0x00008000, 0x00004000, 0x00002000, \
88         0x00001000, 0x00000800, 0x00000400, \
89         0x00000200, 0x00000100, 0x00000080, \
90         0x00000040, 0x00000020, 0x00000010, \
91         0x00000008, 0x00000004, 0x00000002, \
92         0x00000001}
93     #define FIXED_K_VALUES {0x2d413ccc, 0x287a26c4, 0x2744c374, \
94         0x26f72283, 0x26e3b583, 0x26ded9f5, \
95         0x26dda30d, 0x26dd5552, 0x26dd41e4, \
96         0x26dd3d08, 0x26dd3bd1, 0x26dd3b83, \
97         0x26dd3b70, 0x26dd3b6b, 0x26dd3b6a, \
98         0x26dd3b6a, 0x26dd3b6a, 0x26dd3b6a, \
99         0x26dd3b6a, 0x26dd3b6a, 0x26dd3b6a, \
100        0x26dd3b6a, 0x26dd3b6a, 0x26dd3b6a, \
101        0x26dd3b6a, 0x26dd3b6a, 0x26dd3b6a, \
102        0x26dd3b6a, 0x26dd3b6a, 0x26dd3b6a, \
103        0x26dd3b6a}
104     #define max_iter 30
105 #elif BITS == 16
106     #define FIXED_ONE 0x4000
107     #define NEG_CONSTANT 0x8000
108     #define FIXED_HALF_PI 0x6487
109     #define FIXED_ANGLES {0x3243, 0x1dac, 0x0fad, 0x07f5, 0x03fe, \
110         0x01ff, 0x0100, 0x0080, 0x0040, 0x0020, \
111         0x0010, 0x0008, 0x0004, 0x0002, 0x0001}
112     #define FIXED_K_VALUES {0x2d41, 0x287a, 0x2744, 0x26f7, 0x26e3, \
113         0x26de, 0x26dd, 0x26dd, 0x26dd, 0x26dd, \
114         0x26dd, 0x26dd, 0x26dd, 0x26dd, 0x26dd}
115     #define MAX_ITER 15
116 #elif BITS == 8

```

```

117 #define FIXED_ONE 0x40
118 #define NEG_CONSTANT 0x80
119 #define FIXED_HALF_PI 0x64
120 #define FIXED_ANGLES {0x32, 0x1d, 0x10, 0x08, 0x04, 0x02, 0x01}
121 #define FIXED_K_VALUES {0x2d, 0x28, 0x27, 0x26, 0x26, 0x26, 0x26}
122 #define MAX_ITER 7
123 #else
124 #error "you shouldn't be able to get here; you done messed up"
125 #endif
126
127 double *cordic_trig(const double, const unsigned int);
128
129 double cordic_cos(double, unsigned int);
130 double cordic_sin(double, unsigned int);
131 double cordic_tan(double, unsigned int);
132
133 double cordic_atan_bounde(const double, const unsigned int);
134 double cordic_atan(double, unsigned int);
135 double cordic_acos(double, unsigned int);
136 double cordic_asin(double, unsigned int);
137
138 #endif

```

A.4 Exponential and Logarithm Code

Code for Integer Exponentiation:

File : int_exp.c

```

1 #include <stdio.h>
2 #include <gmp.h>
3 #include <mpfr.h>
4
5 #include "int_exp.h"
6 #include "utilities.h"
7
8 double naive_int_exp(const double x, const int a)
9 {
10     if(a < 0)
11         return 1/naive_int_exp(x, -a);
12     double z = 1;
13     int n = a;
14     while(n--)
15         z *= x;
16     return z;
17 }
18
19 double squaring_int_exp(const double x, const int a)
20 {
21     if(a < 0)
22         return 1/squaring_int_exp(x, -a);
23     double y = x, z = 1;
24     int n = a;
25     while(n)
26     {
27         if(n%2)
28         {
29             z *= y;

```

```

30     --n;
31 }
32 y *= y;
33 n >>= 1;
34 }
35 return z;
36 }
37
38 void mpfr_naive_int_exp(mpfr_t R, mpfr_t x, mpz_t a)
39 {
40     if(mpz_cmp_ui(a, 0) < 0)
41     {
42         mpz_t b;
43         mpz_init_set(b, a);
44         mpz_neg(b, b);
45         mpfr_naive_int_exp(R, x, b);
46         mpfr_ui_div(R, 1, R, MPFR_RNDN);
47     }
48     else
49     {
50         mpz_t n;
51         mpz_init_set(n, a);
52         mpfr_set_ui(R, 1, MPFR_RNDN);
53         while(mpz_cmp_ui(n, 0) > 0)
54         {
55             mpfr_mul(R, R, x, MPFR_RNDN);
56             mpz_sub_ui(n, n, 1);
57         }
58     }
59 }
60
61 void mpfr_squaring_int_exp(mpfr_t R, mpfr_t x, mpz_t a)
62 {
63     if(mpz_cmp_ui(a, 0) < 0)
64     {
65         mpz_t b;
66         mpz_init_set(b, a);
67         mpz_neg(b, b);
68         mpfr_squaring_int_exp(R, x, b);
69         mpfr_ui_div(R, 1, R, MPFR_RNDN);
70     }
71     else
72     {
73         mpfr_t y;
74         mpz_t n;
75         mpfr_init_set(y, x, MPFR_RNDN);
76         mpfr_set_ui(R, 1, MPFR_RNDN);
77         mpz_init_set(n, a);
78         while(mpz_cmp_ui(n, 0) > 0)
79         {
80             if(mpz_odd_p(n))
81             {
82                 mpfr_mul(R, R, y, MPFR_RNDN);
83                 mpz_sub_ui(n, n, 1);
84             }
85             mpfr_mul(y, y, y, MPFR_RNDN);
86             mpz_div_ui(n, n, 2);
87         }

```

```

88     }
89 }
90
91 #ifdef COMPILE_MAIN
92 int main(int argc, char **argv)
93 {
94     double x;
95     unsigned int n, D, p;
96     mpfr_t X, R;
97     mpz_t N;
98     char sf[50];
99
100    if(argc > 1)
101    {
102        switch(argv[1][0])
103        {
104            case 'a':
105                if(argc == 5 &&
106                    sscanf(argv[2], "%lf", &x) == 1 &&
107                    sscanf(argv[3], "%u", &n) == 1 &&
108                    sscanf(argv[4], "%u", &D) == 1)
109                    printf("(%.*lf)^(%d) = %.*lf (Naive)\n",
110                        d(D), x, n, D, naive_int_exp(x, n));
111                else
112                    printf("Usage: %s a <x=Base for exp> "
113                        "<n=Exponent for exp> "
114                        "<D=Number of digits to display>\n",
115                        argv[0]);
116                break;
117
118            case 'b':
119                if(argc == 5 &&
120                    sscanf(argv[2], "%lf", &x) == 1 &&
121                    sscanf(argv[3], "%u", &n) == 1 &&
122                    sscanf(argv[4], "%u", &D) == 1)
123                    printf("(%.*lf)^(%d) = %.*lf (Squaring)\n",
124                        d(D), x, n, D, squaring_int_exp(x, n));
125                else
126                    printf("Usage: %s b <x=Base for exp> "
127                        "<n=Exponent for exp> "
128                        "<D=Number of digits to display>\n",
129                        argv[0]);
130                break;
131
132            case 'c':
133                if (argc == 6 &&
134                    sscanf(argv[4], "%u", &D) == 1 &&
135                    sscanf(argv[5], "%u", &p) == 1)
136                {
137                    mpfr_set_default_prec(p);
138
139                    if (mpfr_init_set_str(X, argv[2], 10, MPFR_RNDN)==0 &&
140                        mpz_init_set_str(N, argv[3], 10) == 0)
141                    {
142                        mpfr_init(R);
143
144                        sprintf(sf, "(%.%.%uRnf)^%Zd =~\t(Naive)"
145                            "\n\t%.%.%uRnf\n", d(D), D);

```



```

146     mpfr_naive_int_exp(R, X, N);
147     mpfr_printf(sf, X, N, R);
148 }
149 else
150     printf(" Usage: %s c <X=Base for exp> "
151           "<N=Exponent for exp> "
152           "<D=Number of digitsto calculate to> "
153           "<p=bits of precision>\n", argv[0]);
154 }
155 else
156     printf(" Usage: %s c <X=Base for exp> "
157           "<N=Exponent for exp> "
158           "<D=Number of digitsto calculate to> "
159           "<p=bits of precision>\n", argv[0]);
160 break;
161
162 case 'd':
163     if (argc == 6 &&
164         sscanf(argv[4], "%u", &D) == 1 &&
165         sscanf(argv[5], "%u", &p) == 1)
166     {
167         mpfr_set_default_prec(p);
168
169         if (mpfr_init_set_str(X, argv[2], 10, MPFR_RNDN)==0 &&
170             mpz_init_set_str(N, argv[3], 10) == 0)
171         {
172             mpfr_init(R);
173
174             sprintf(sf, "(%%.%%uRNf)^%%Zd =~\t(Squaring)"
175                   "\n\t%%.%%uRNf\n", d(D), D);
176
177             mpfr_squaring_int_exp(R, X, N);
178             mpfr_printf(sf, X, N, R);
179         }
180     }
181     else
182         printf(" Usage: %s d <X=Base for exp> "
183               "<N=Exponent for exp> "
184               "<D=Number of digitsto calculate to> "
185               "<p=bits of precision>\n", argv[0]);
186 }
187 else
188     printf(" Usage: %s d <X=Base for exp> "
189           "<N=Exponent for exp> "
190           "<D=Number of digitsto calculate to> "
191           "<p=bits of precision>\n", argv[0]);
192 break;
193
194 default:
195     printf(" Usage: %s <a/b/c/d> <arguments>\n", argv[0]);
196 }
197 }
198 else
199     printf(" Usage: %s <a/b/c/d> <arguments>\n", argv[0]);
200 }
201 #endif

```

Code for Taylor Exponentials and Logarithms:

```

1 #include <stdio.h>
2 #include <assert.h>
3 #include <gmp.h>
4 #include <mpfr.h>
5 #include <math.h>
6
7 #include "int_exp.h"
8 #include "log_exp_utilities.h"
9 #include "utilities.h"
10 #include "taylor_exp_log.h"
11
12 #define INIT_CONSTANTS in = fopen(NAT_LOG_2.INFILE, "r"); \
13     mpfr_init(MPFR_NAT_LOG_2); \
14     mpfr_inp_str(MPFR_NAT_LOG_2, in, 10, MPFR_RNDN); \
15     fclose(in);
16
17 mpfr_t MPFR_NAT_LOG_2;
18
19 double naive_exp(double x, unsigned int n)
20 {
21     double z = 1 + x/n;
22     return z > 0 ? squaring_int_exp(z, n)
23         : n%2 ? -squaring_int_exp(-z, n)
24         : squaring_int_exp(-z, n);
25 }
26
27 void mpfr_naive_exp(mpfr_t R, mpfr_t x, mpz_t n)
28 {
29     assert(mpz_cmp_ui(n, 0) >= 0);
30
31     mpfr_t z;
32
33     mpfr_init_set(z, x, MPFR_RNDN);
34     mpfr_div_z(z, z, n, MPFR_RNDN);
35     mpfr_add_ui(z, z, 1, MPFR_RNDN);
36
37     if(mpfr_cmp_ui(z, 0) > 0)
38         mpfr_squaring_int_exp(R, z, n);
39     else
40     {
41         mpfr_neg(z, z, MPFR_RNDN);
42         mpfr_squaring_int_exp(R, z, n);
43
44         if(mpz_odd_p(n))
45             mpfr_neg(R, R, MPFR_RNDN);
46     }
47 }
48
49 double taylor_exp(double x, unsigned int n)
50 {
51     double t, z;
52     t = 1;
53     z = 1;
54
55     for(int k = 1; k < n; ++k)
56     {
57         t *= x;

```

```

58     t /= k;
59     z += t;
60 }
61
62     return z;
63 }
64
65 void mpfr_taylor_exp(mpfr_t R, mpfr_t x, mpz_t n)
66 {
67     assert(mpz_cmp_ui(n, 0) >= 0);
68
69     mpfr_t t;
70     mpz_t k;
71
72     mpfr_init_set_ui(t, 1, MPFR_RNDN);
73     mpfr_set_ui(R, 1, MPFR_RNDN);
74
75     for(mpz_init_set_ui(k, 1); mpz_cmp(k, n) < 0; mpz_add_ui(k, k, 1))
76     {
77         mpfr_mul(t, t, x, MPFR_RNDN);
78         mpfr_div_z(t, t, k, MPFR_RNDN);
79         mpfr_add(R, R, t, MPFR_RNDN);
80     }
81 }
82
83 double taylor_nat_log(double x, unsigned int n)
84 {
85     assert(n > 0);
86     assert(x > 0);
87
88     double a, t, z;
89     int b;
90
91     a = frexp(x, &b);
92     a = 1 - a;
93
94     t = a;
95     z = a;
96
97     for(int k = 2; k < n; ++k)
98     {
99         t *= a;
100         z += t/k;
101     }
102
103     return b*NAT_LOG_2 - z;
104 }
105
106 void mpfr_taylor_nat_log(mpfr_t R, mpfr_t x, mpz_t n)
107 {
108     assert(mpz_cmp_ui(n, 0) > 0);
109     assert(mpfr_cmp_ui(x, 0) > 0);
110
111     mpfr_t a, t, tt;
112     mpfr_exp_t b;
113     mpz_t k;
114     unsigned int f = 1000, F = 1000;
115

```

```

116 mpfr_init(a);
117 mpfr_frexp(&b, a, x, MPFR_RNDN);
118 mpfr_ui_sub(a, 1, a, MPFR_RNDN);
119
120 mpfr_init_set(t, a, MPFR_RNDN);
121 mpfr_init(tt);
122 mpfr_set(R, a, MPFR_RNDN);
123
124 for(mpz_init_set_ui(k, 2); mpz_cmp(k, n) < 0; mpz_add_ui(k, k, 1))
125 {
126     mpfr_mul(t, t, a, MPFR_RNDN);
127     mpfr_div_z(tt, t, k, MPFR_RNDN);
128     mpfr_add(R, R, tt, MPFR_RNDN);
129 }
130
131 mpfr_mul_si(a, MPFR_NAT_LOG_2, b, MPFR_RNDN);
132 mpfr_sub(R, a, R, MPFR_RNDN);
133 }
134
135 double taylor_log(double x, double y, unsigned int n)
136 {
137     assert(x > 0);
138     assert(y > 0);
139     assert(n > 0);
140
141     return taylor_nat_log(y, n)/taylor_nat_log(x, n);
142 }
143
144 double taylor_pow(double x, double y, unsigned int n)
145 {
146     assert(x > 0);
147     assert(n > 0);
148
149     return taylor_exp(y*taylor_nat_log(x, n), n);
150 }
151
152 void mpfr_taylor_log(mpfr_t R, mpfr_t x, mpfr_t y, mpz_t n)
153 {
154     assert(mpfr_cmp_ui(x, 0) > 0);
155     assert(mpfr_cmp_ui(y, 0) > 0);
156     assert(mpz_cmp_ui(n, 0) >= 0);
157
158     mpfr_t A;
159     mpfr_init(A);
160
161     mpfr_taylor_nat_log(A, x, n);
162     mpfr_taylor_nat_log(R, y, n);
163
164     mpfr_div(R, R, A, MPFR_RNDN);
165 }
166
167 void mpfr_taylor_pow(mpfr_t R, mpfr_t x, mpfr_t y, mpz_t n)
168 {
169     assert(mpfr_cmp_ui(x, 0) > 0);
170     assert(mpz_cmp_ui(n, 0) > 0);
171
172     mpfr_t A;
173     mpfr_init(A);

```

```

174
175     mpfr_taylor_nat_log(A, x, n);
176     mpfr_mul(A, y, A, MPFR_RNDN);
177     mpfr_taylor_exp(R, A, n);
178 }
179
180 #ifdef COMPILE_MAIN
181 int main(int argc, char **argv)
182 {
183     double x, y;
184     unsigned int n, D, p;
185     mpfr_t X, Y, R;
186     mpz_t N;
187     char sf[50];
188     FILE *in;
189
190     if(argc > 1)
191     {
192         switch(argv[1][0])
193         {
194             case 'a':
195                 if(argc == 5 &&
196                     sscanf(argv[2], "%lf", &x) == 1 &&
197                     sscanf(argv[3], "%u", &n) == 1 &&
198                     sscanf(argv[4], "%u", &D) == 1)
199                     printf("exp(%.1f) ~= %.1f (naive)\n",
200                         d(D), x, D, naive_exp(x, n));
201                 else
202                     printf("Usage: %s a <x=Value for exp(x)> <n> "
203                         "<D=digits to display>\n", argv[0]);
204                 break;
205
206             case 'b':
207                 if(argc == 5 &&
208                     sscanf(argv[2], "%lf", &x) == 1 &&
209                     sscanf(argv[3], "%u", &n) == 1 &&
210                     sscanf(argv[4], "%u", &D) == 1)
211                     printf("exp(%.1f) ~= %.1f\n",
212                         d(D), x, D, taylor_exp(x, n));
213                 else
214                     printf("Usage: %s b <x=Value for exp(x)> <n> "
215                         "<D=digits to display>\n", argv[0]);
216                 break;
217
218             case 'c':
219                 if(argc == 5 &&
220                     sscanf(argv[2], "%lf", &x) == 1 &&
221                     sscanf(argv[3], "%u", &n) == 1 &&
222                     sscanf(argv[4], "%u", &D) == 1)
223                     printf("ln(%.1f) ~= %.1f\n",
224                         d(D), x, D, taylor_nat_log(x, n));
225                 else
226                     printf("Usage: %s c <x=Value for ln(x)> <n> "
227                         "<D=digits to display>\n", argv[0]);
228                 break;
229
230             case 'd':
231                 if(argc == 6 &&

```

```

232     sscanf(argv[2], "%lf", &x) == 1 &&
233     sscanf(argv[3], "%lf", &y) == 1 &&
234     sscanf(argv[4], "%u", &n) == 1 &&
235     sscanf(argv[5], "%u", &D) == 1)
236     printf("pow(%.1f, %.1f) ~= %.1f\n",
237           d(D), x, d(D), y, D, taylor_pow(x, y, n));
238 else
239     printf("Usage: %s d <x=Value for pow(x,y)> "
240           "<y=Value for pow(x,y)> <n> "
241           "<D=digits to display>\n", argv[0]);
242 break;
243
244 case 'e':
245     if(argc == 6 &&
246         sscanf(argv[2], "%lf", &x) == 1 &&
247         sscanf(argv[3], "%lf", &y) == 1 &&
248         sscanf(argv[4], "%u", &n) == 1 &&
249         sscanf(argv[5], "%u", &D) == 1)
250         printf("log(%.1f, %.1f) ~= %.1f\n",
251               d(D), x, d(D), y, D, taylor_log(x, y, n));
252     else
253         printf("Usage: %s e <x=Value for log(x, y)> "
254               "<y=Value for log(x, y)> <n> "
255               "<D=digits to display>\n", argv[0]);
256 break;
257
258 case 'f':
259     if(argc == 6 &&
260         sscanf(argv[4], "%u", &D) == 1 &&
261         sscanf(argv[5], "%u", &p) == 1)
262     {
263         mpfr_set_default_prec(p);
264         INIT_CONSTANTS
265
266         if(mpfr_init_set_str(X, argv[2], 10, MPFR_RNDN) == 0 &&
267             mpz_init_set_str(N, argv[3], 10) == 0)
268         {
269             mpfr_init(R);
270
271             sprintf(sf, "exp(%%.1f) = ~\t(naive)"
272                   "\n\t%.1f\n", d(D), D);
273
274             mpfr_naive_exp(R, X, N);
275             mpfr_printf(sf, X, R);
276         }
277     else
278         printf("Usage: %s f <X=value for exp(X)> <N> "
279               "<D=Digits to display> "
280               "<p=bits of precision>\n", argv[0]);
281 }
282 else
283     printf("Usage: %s f <X=value for exp(X)> <N> "
284           "<D=Digits to display> "
285           "<p=bits of precision>\n", argv[0]);
286 break;
287
288 case 'g':
289     if(argc == 6 &&

```

```

290     sscanf(argv[4], "%u", &D) == 1 &&
291     sscanf(argv[5], "%u", &p) == 1)
292 {
293     mpfr_set_default_prec(p);
294     INIT_CONSTANTS
295
296     if(mpfr_init_set_str(X, argv[2], 10, MPFR_RNDN) == 0 &&
297         mpz_init_set_str(N, argv[3], 10) == 0)
298     {
299         mpfr_init(R);
300
301         sprintf(sf, "exp(%%.%%uRNf) =~"
302             "\n\t%.%%uRNf\n", d(D), D);
303
304         mpfr_taylor_exp(R, X, N);
305         mpfr_printf(sf, X, R);
306     }
307     else
308         printf("Usage: %s g <X=value for exp(X)> <N> "
309             "<D=Digits to display> "
310             "<p=bits of precision>\n", argv[0]);
311 }
312 else
313     printf("Usage: %s g <X=value for exp(X)> <N> "
314         "<D=Digits to display> "
315         "<p=bits of precision>\n", argv[0]);
316 break;
317
318 case 'h':
319     if(argc == 6 &&
320         sscanf(argv[4], "%u", &D) == 1 &&
321         sscanf(argv[5], "%u", &p) == 1)
322     {
323         mpfr_set_default_prec(p);
324         INIT_CONSTANTS
325
326         if(mpfr_init_set_str(X, argv[2], 10, MPFR_RNDN) == 0 &&
327             mpz_init_set_str(N, argv[3], 10) == 0)
328         {
329             mpfr_init(R);
330
331             sprintf(sf, "ln(%%.%%uRNf) =~"
332                 "\n\t%.%%uRNf\n", d(D), D);
333
334             mpfr_taylor_nat_log(R, X, N);
335             mpfr_printf(sf, X, R);
336         }
337         else
338             printf("Usage: %s h <X=value for ln(X)> <N> "
339                 "<D=Digits to display> "
340                 "<p=bits of precision>\n", argv[0]);
341     }
342     else
343         printf("Usage: %s h <X=value for ln(X)> <N> "
344             "<D=Digits to display> "
345             "<p=bits of precision>\n", argv[0]);
346 break;
347

```

```

348 case 'i':
349     if (argc == 7 &&
350         sscanf(argv[5], "%u", &D) == 1 &&
351         sscanf(argv[6], "%u", &p) == 1)
352     {
353         mpfr_set_default_prec(p);
354         INIT_CONSTANTS
355
356         if (mpfr_init_set_str(X, argv[2], 10, MPFR_RNDN) == 0 &&
357             mpfr_init_set_str(Y, argv[3], 10, MPFR_RNDN) == 0 &&
358             mpz_init_set_str(N, argv[4], 10) == 0)
359         {
360             mpfr_init(R);
361
362             sprintf(sf, "pow(%%.%%uRNf, %.%%uRNf) =~"
363                     "\n\t%.%%uRNf\n", d(D), d(D), D);
364
365             mpfr_taylor_pow(R, X, Y, N);
366             mpfr_printf(sf, X, Y, R);
367         }
368     else
369         printf("Usage: %s i <X=value for pow(X,Y)> "
370              "<Y=value for pow(X,Y)> <N> "
371              "<D=Digits to display> "
372              "<p=bits of precision>\n", argv[0]);
373 }
374 else
375     printf("Usage: %s i <X=value for pow(X,Y)> "
376          "<Y=value for pow(X,Y)> <N> "
377          "<D=Digits to display> "
378          "<p=bits of precision>\n", argv[0]);
379 break;
380
381 case 'j':
382     if (argc == 7 &&
383         sscanf(argv[5], "%u", &D) == 1 &&
384         sscanf(argv[6], "%u", &p) == 1)
385     {
386         mpfr_set_default_prec(p);
387         INIT_CONSTANTS
388
389         if (mpfr_init_set_str(X, argv[2], 10, MPFR_RNDN) == 0 &&
390             mpfr_init_set_str(Y, argv[3], 10, MPFR_RNDN) == 0 &&
391             mpz_init_set_str(N, argv[4], 10) == 0)
392         {
393             mpfr_init(R);
394
395             sprintf(sf, "log(%%.%%uRNf, %.%%uRNf) =~"
396                     "\n\t%.%%uRNf\n", d(D), d(D), D);
397
398             mpfr_taylor_log(R, X, Y, N);
399             mpfr_printf(sf, X, Y, R);
400         }
401     else
402         printf("Usage: %s j <X=value for log(X,Y)> "
403              "<Y=value for log(X,Y)> <N> "
404              "<D=Digits to display> "
405              "<p=bits of precision>\n", argv[0]);

```



```

406     }
407     else
408         printf("Usage: %s j <X=value for log(X,Y)> "
409             "<Y=value for log(X,Y)> <N> "
410             "<D=Digits to display> "
411             "<p=bits of precision>\n", argv[0]);
412     break;
413
414     default:
415         printf("Usage: %s <a/b/c/d/e/f/g/h/i/j> <arguments>\n",
416             argv[0]);
417     }
418 }
419 else
420     printf("Usage: %s <a/b/c/d/e/f/g/h/i/j> <arguments>\n", argv[0]);
421 }
422 #endif

```

Code for Hyperbolic Logarithms:

File : hyperbolic_log.c

```

1  #include <stdio.h>
2  #include <assert.h>
3  #include <gmp.h>
4  #include <mpfr.h>
5  #include <math.h>
6
7  #include "log_exp_utilities.h"
8  #include "utilities.h"
9  #include "hyperbolic_log.h"
10
11 #define INIT_CONSTANTS in = fopen(NAT_LOG_2_INFILE, "r"); \
12     mpfr_init(MPFR_NAT_LOG_2); \
13     mpfr_inp_str(MPFR_NAT_LOG_2, in, 10, MPFR_RNDN); \
14     fclose(in);
15
16 mpfr_t MPFR_NAT_LOG_2;
17
18 double hyperbolic_nat_log(double x, unsigned int n)
19 {
20     assert(x > 0);
21
22     double t, y, z, a;
23     unsigned int d;
24     int b;
25
26     a = frexp(x, &b);
27
28     t = (a-1)/(a+1);
29     d = 1;
30     y = t*t;
31     z = t;
32
33     for(unsigned int k = 1; k <= n; ++k)
34     {
35         t *= y;
36         d += 2;
37         z += t/d;

```

```

38     }
39
40     return b*NAT_LOG_2 + 2*z;
41 }
42
43 void mpfr_hyperbolic_nat_log(mpfr_t R, mpfr_t x, mpz_t n)
44 {
45     assert(mpfr_cmp_ui(x, 0) > 0);
46     assert(mpz_cmp_ui(n, 0) >= 0);
47
48     mpfr_t t, y, z, tmp, a;
49     mpfr_exp_t b;
50     mpz_t d, k;
51
52     mpfr_init(a);
53     mpfr_init(t);
54     mpfr_init(tmp);
55     mpfr_init(y);
56
57     mpfr_frexp(&b, a, x, MPFR_RNDN);
58
59     mpfr_sub_ui(t, a, 1, MPFR_RNDN);
60     mpfr_add_ui(tmp, a, 1, MPFR_RNDN);
61     mpfr_div(t, t, tmp, MPFR_RNDN);
62     mpfr_mul(y, t, t, MPFR_RNDN);
63
64     mpfr_init_set(z, t, MPFR_RNDN);
65     mpz_init_set_ui(d, 1);
66
67     for(mpz_init_set_ui(k, 1); mpz_cmp(k, n) <= 0; mpz_add_ui(k, k, 1))
68     {
69         mpfr_mul(t, t, y, MPFR_RNDN);
70         mpz_add_ui(d, d, 2);
71         mpfr_div_z(tmp, t, d, MPFR_RNDN);
72         mpfr_add(z, z, tmp, MPFR_RNDN);
73     }
74
75     mpfr_mul_ui(R, z, 2, MPFR_RNDN);
76     mpfr_mul_si(tmp, MPFR_NAT_LOG_2, b, MPFR_RNDN);
77     mpfr_add(R, R, tmp, MPFR_RNDN);
78 }
79
80 double hyperbolic_log(double x, double y, unsigned int n)
81 {
82     assert(x > 0);
83     assert(y > 0);
84
85     return hyperbolic_nat_log(y, n)/hyperbolic_nat_log(x, n);
86 }
87
88 void mpfr_hyperbolic_log(mpfr_t R, mpfr_t x, mpfr_t y, mpz_t n)
89 {
90     assert(mpfr_cmp_ui(x, 0) > 0);
91     assert(mpfr_cmp_ui(y, 0) > 0);
92     assert(mpz_cmp_ui(n, 0) >= 0);
93
94     mpfr_t A;
95     mpfr_init(A);

```

```

96
97     mpfr_hyperbolic_nat_log(R, y, n);
98     mpfr_hyperbolic_nat_log(A, x, n);
99
100     mpfr_div(R, R, A, MPFR_RNDN);
101 }
102
103 #ifdef COMPILE_MAIN
104 int main(int argc, char **argv)
105 {
106     double x, y;
107     unsigned int n, D, p;
108     mpfr_t X, Y, R;
109     mpz_t N;
110     char sf[50];
111     FILE *in;
112
113     if(argc > 1)
114     {
115         switch(argv[1][0])
116         {
117             case 'a':
118                 if(argc == 5 &&
119                     sscanf(argv[2], "%lf", &x) == 1 &&
120                     sscanf(argv[3], "%u", &n) == 1 &&
121                     sscanf(argv[4], "%u", &D) == 1)
122                     printf("ln(%.1f) ~= %.1f\n",
123                         d(D), x, D, hyperbolic_nat_log(x, n));
124                 else
125                     printf("Usage: %s a <x=Value for ln(x)> <n> "
126                         "<D=digits to display>\n", argv[0]);
127                 break;
128
129             case 'b':
130                 if(argc == 6 &&
131                     sscanf(argv[2], "%lf", &x) == 1 &&
132                     sscanf(argv[3], "%lf", &y) == 1 &&
133                     sscanf(argv[4], "%u", &n) == 1 &&
134                     sscanf(argv[5], "%u", &D) == 1)
135                     printf("log(%.1f, %.1f) ~= %.1f\n",
136                         d(D), x, d(D), y, D, hyperbolic_log(x, y, n));
137                 else
138                     printf("Usage: %s b <x=Value for log(x,y)> "
139                         "<y=Value for log(x,y)> <n> "
140                         "<D=digits to display>\n", argv[0]);
141                 break;
142
143             case 'c':
144                 if(argc == 6 &&
145                     sscanf(argv[4], "%u", &D) == 1 &&
146                     sscanf(argv[5], "%u", &p) == 1)
147                 {
148                     mpfr_set_default_prec(p);
149                     INIT_CONSTANTS
150
151                     if(mpfr_init_set_str(X, argv[2], 10, MPFR_RNDN) == 0 &&
152                         mpz_init_set_str(N, argv[3], 10) == 0)
153                 {

```

```

154     mpfr_init(R);
155
156     sprintf(sf, " ln(%%.%%uRNf) =~\n\t%%.%%uRNf\n",
157             d(D), D);
158
159     mpfr_hyperbolic_nat_log(R, X, N);
160     mpfr_printf(sf, X, R);
161 }
162 else
163     printf(" Usage: %s c <X=value for ln(X)> <N> "
164            "<D=Digits to display> "
165            "<p=bits of precision>\n", argv[0]);
166 }
167 else
168     printf(" Usage: %s c <X=value for ln(X)> <N> "
169            "<D=Digits to display> "
170            "<p=bits of precision>\n", argv[0]);
171 break;
172
173 case 'd':
174     if(argc == 7 &&
175        sscanf(argv[5], "%u", &D) == 1 &&
176        sscanf(argv[6], "%u", &p) == 1)
177     {
178         mpfr_set_default_prec(p);
179         INIT_CONSTANTS
180
181         if(mpfr_init_set_str(X, argv[2], 10, MPFR_RNDN) == 0 &&
182            mpfr_init_set_str(Y, argv[3], 10, MPFR_RNDN) == 0 &&
183            mpz_init_set_str(N, argv[4], 10) == 0)
184         {
185             mpfr_init(R);
186
187             sprintf(sf, " ln(%%.%%uRNf, %%.%%uRNf) =~\n\t"
188                    "%%.%%uRNf\n", d(D), d(D), D);
189
190             mpfr_hyperbolic_log(R, X, Y, N);
191             mpfr_printf(sf, X, Y, R);
192         }
193     else
194         printf(" Usage: %s d <X=Value for ln(X,Y)> "
195                "<Y=Value for ln(X,Y)> "
196                "<D=Digits to display> "
197                "<p=bits of precision>\n", argv[0]);
198     }
199 else
200     printf(" Usage: %s d <X=Value for ln(X, Y)> "
201            "<Y=Value for ln(X,Y)> "
202            "<D=Digits to display> "
203            "<p=bits of precision>\n", argv[0]);
204 break;
205
206 default:
207     printf(" Usage: %s <a/b/c/d> <arguments>\n", argv[0]);
208 }
209 }
210 else
211     printf(" Usage: %s <a/b/c/d> <arguments>\n", argv[0]);

```

```

212 }
213 #endif

```

Code for Continued Fraction Exponentials:

File : cont_frac_exp.c

```

1 #include <stdio.h>
2 #include <assert.h>
3 #include <gmp.h>
4 #include <mpfr.h>
5 #include <math.h>
6
7 #include "log_exp_utilities.h"
8 #include "utilities.h"
9 #include "hyperbolic_log.h"
10 #include "int_exp.h"
11 #include "cont_frac_exp.h"
12
13 #define INIT_CONSTANTS in = fopen(NAT_LOG_2_INFILE, "r"); \
14     mpfr_init(MPFR_NAT_LOG_2); \
15     mpfr_inp_str(MPFR_NAT_LOG_2, in, 10, MPFR_RNDN); \
16     fclose(in); \
17     in = fopen(E_CONST_INFILE, "r"); \
18     mpfr_init(MPFR_E_CONST); \
19     mpfr_inp_str(MPFR_E_CONST, in, 10, MPFR_RNDN); \
20     fclose(in); \
21     mpfr_init_set_ui(MPFR_TWO, 2, MPFR_RNDN);
22
23 mpfr_t MPFR_NAT_LOG_2, MPFR_TWO, MPFR_E_CONST;
24
25 double cont_frac_exp_v1(double x, unsigned int n)
26 {
27     double a = x, pA, A, nA, ca, pB, B, nB, cb;
28     unsigned int b = 0;
29
30     if (x < 0)
31         return 1/cont_frac_exp_v1(-x, n);
32     else if (x == 0)
33         return 1;
34     else if (x > 1)
35         a = frexp(x, &b);
36
37     pA = a + 1;
38     pB = 1;
39     A = a*a + 2*a + 2;
40     B = 2;
41
42     ca = -a;
43     cb = 2 + a;
44
45     for(unsigned int k = 2; k <= n; ++k)
46     {
47         ca -= a;
48         ++cb;
49
50         nA = cb*A + ca*pA;
51         nB = cb*B + ca*pB;
52

```

```

53     pA = A;
54     pB = B;
55     A = nA;
56     B = nB;
57 }
58
59     return b ? squaring_int_exp(A/B, (unsigned int)squaring_int_exp(2, b))
60             : A/B;
61 }
62
63 double cont_frac_exp_v2(double x, unsigned int n)
64 {
65     double a = x, pA, A, nA, ca, pB, B, nB;
66     unsigned int b = 0, cb;
67
68     if (x < 0)
69         return 1/cont_frac_exp_v2(-x, n);
70     else if (x == 0)
71         return 1;
72     else if (x > 1)
73         a = frexp(x, &b);
74
75     pA = 1;
76     pB = 1;
77     A = 1;
78     B = 1 - a;
79
80     ca = 0;
81     cb = 1;
82
83     for(unsigned int k = 3; k <= n; ++k)
84     {
85         ++cb;
86
87         if (k % 2)
88         {
89             ca += a;
90             nA = cb*A + ca*pA;
91             nB = cb*B + ca*pB;
92         }
93         else
94         {
95             nA = cb*A - ca*pA;
96             nB = cb*B - ca*pB;
97         }
98
99         pA = A;
100        pB = B;
101        A = nA;
102        B = nB;
103    }
104
105    return b ? squaring_int_exp(A/B, (unsigned int)squaring_int_exp(2, b))
106            : A/B;
107 }
108
109 double cont_frac_exp_v3(double x, unsigned int n)
110 {

```

```

111 double a = x, pA, A, nA, ca, pB, B, nB;
112 unsigned int b = 0, cb;
113
114 if (x < 0)
115     return 1/cont_frac_exp_v3(-x, n);
116 else if (x == 0)
117     return 1;
118 else if (x > 1)
119     a = frexp(x, &b);
120
121 pA = 1;
122 pB = 1;
123 A = 2 + a;
124 B = 2 - a;
125
126 ca = a*a;
127 cb = 2;
128
129 for(unsigned int k = 2; k <= n; ++k)
130 {
131     cb += 4;
132
133     nA = cb*A + ca*pA;
134     nB = cb*B + ca*pB;
135
136     pA = A;
137     pB = B;
138     A = nA;
139     B = nB;
140 }
141
142 return b ? squaring_int_exp(A/B, (unsigned int)squaring_int_exp(2, b))
143         : A/B;
144 }
145
146 void mpfr_cont_frac_exp_v3(mpfr_t R, mpfr_t x, mpz_t n)
147 {
148     assert(mpz_cmp_ui(n, 0) >= 0);
149
150     mpfr_t a, pA, A, nA, ca, pB, B, nB, C, tmp1, tmp2;
151     mpfr_exp_t b = 0;
152     mpz_t N, k, cb;
153
154     if (mpfr_cmp_ui(x, 0) < 0)
155     {
156         mpfr_neg(x, x, MPFR_RNDN);
157         mpfr_cont_frac_exp_v3(R, x, n);
158         mpfr_ui_div(R, 1, R, MPFR_RNDN);
159     }
160     else if (mpfr_cmp_ui(x, 0) == 0)
161     {
162         mpfr_set_ui(R, 1, MPFR_RNDN);
163     }
164     else
165     {
166         mpfr_init(a);
167         if (mpfr_cmp_ui(x, 1) > 0)
168         {

```

```

169     mpfr_frexp(&b, a, x, MPFR_RNDN);
170     mpz_init(N);
171     mpfr_init(C);
172     mpz_ui_pow_ui(N, 2, b);
173 }
174 else
175 {
176     mpfr_init_set_ui(C, 0, MPFR_RNDN);
177     mpfr_set(a, x, MPFR_RNDN);
178 }
179
180     mpfr_init_set_ui(pA, 1, MPFR_RNDN);
181     mpfr_init_set_ui(pB, 1, MPFR_RNDN);
182     mpfr_init_set_ui(A, 2, MPFR_RNDN);
183     mpfr_init_set_ui(B, 2, MPFR_RNDN);
184     mpfr_add(A, A, a, MPFR_RNDN);
185     mpfr_sub(B, B, a, MPFR_RNDN);
186
187     mpfr_init_set(ca, a, MPFR_RNDN);
188     mpfr_mul(ca, ca, a, MPFR_RNDN);
189     mpz_init_set_ui(cb, 2);
190
191     mpfr_init(nA);
192     mpfr_init(nB);
193     mpfr_init(tmp1);
194     mpfr_init(tmp2);
195
196     for(mpz_init_set_ui(k, 2); mpz_cmp(k, n) <= 0; mpz_add_ui(k, k, 1))
197     {
198         mpz_add_ui(cb, cb, 4);
199
200         mpfr_mul_z(tmp1, A, cb, MPFR_RNDN);
201         mpfr_mul(tmp2, ca, pA, MPFR_RNDN);
202         mpfr_add(nA, tmp1, tmp2, MPFR_RNDN);
203
204         mpfr_mul_z(tmp1, B, cb, MPFR_RNDN);
205         mpfr_mul(tmp2, ca, pB, MPFR_RNDN);
206         mpfr_add(nB, tmp1, tmp2, MPFR_RNDN);
207
208         mpfr_set(pA, A, MPFR_RNDN);
209         mpfr_set(pB, B, MPFR_RNDN);
210         mpfr_set(A, nA, MPFR_RNDN);
211         mpfr_set(B, nB, MPFR_RNDN);
212     }
213
214     mpfr_div(R, A, B, MPFR_RNDN);
215
216     if(b)
217     {
218         mpfr_set(A, R, MPFR_RNDN);
219         mpfr_squaring_int_exp(R, A, N);
220     }
221 }
222 }
223
224 double improved_pow(double x, double y, unsigned int n)
225 {
226     assert(x > 0);

```



```

227
228     return cont_frac_exp_v3(y * hyperbolic_nat_log(x, n), n);
229 }
230
231 void mpfr_improved_pow(mpfr_t R, mpfr_t x, mpfr_t y, mpz_t n)
232 {
233     assert(mpz_cmp_ui(n, 0) >= 0);
234     assert(mpfr_cmp_ui(x, 0) > 0);
235
236     mpfr_t A;
237     mpfr_init(A);
238
239     mpfr_hyperbolic_nat_log(A, x, n);
240     mpfr_mul(A, y, A, MPFR_RNDN);
241     mpfr_cont_frac_exp_v3(R, A, n);
242 }
243
244 #ifdef COMPILE_MAIN
245 int main(int argc, char **argv)
246 {
247     double x, y;
248     unsigned int n, D, p;
249     mpfr_t X, Y, R;
250     mpz_t N;
251     char sf[50];
252     FILE *in;
253
254     if(argc > 1)
255     {
256         switch(argv[1][0])
257         {
258             case 'a':
259                 if(argc == 5 &&
260                    sscanf(argv[2], "%lf", &x) == 1 &&
261                    sscanf(argv[3], "%u", &n) == 1 &&
262                    sscanf(argv[4], "%u", &D) == 1)
263                     printf("exp(%.1f) ~ = %.1f (v1)\n",
264                            d(D), x, D, cont_frac_exp_v1(x, n));
265                 else
266                     printf("Usage: %s a <x=Value for exp(x)> <n> "
267                            "<D=Digits to display>\n", argv[0]);
268                 break;
269
270             case 'b':
271                 if(argc == 5 &&
272                    sscanf(argv[2], "%lf", &x) == 1 &&
273                    sscanf(argv[3], "%u", &n) == 1 &&
274                    sscanf(argv[4], "%u", &D) == 1)
275                     printf("exp(%.1f) ~ = %.1f (v2)\n",
276                            d(D), x, D, cont_frac_exp_v2(x, n));
277                 else
278                     printf("Usage: %s b <x=Value for exp(x)> <n> "
279                            "<D=Digits to display>\n", argv[0]);
280                 break;
281
282             case 'c':
283                 if(argc == 5 &&
284                    sscanf(argv[2], "%lf", &x) == 1 &&

```

```

285     sscanf(argv[3], "%u", &n) == 1 &&
286     sscanf(argv[4], "%u", &D) == 1)
287     printf("exp(%.1f) ~= %.1f (v3)\n",
288           d(D), x, D, cont_frac_exp_v3(x, n));
289 else
290     printf("Usage: %s c <x=Value for exp(x)> <n> "
291           "<D=Digits to display>\n", argv[0]);
292 break;
293
294 case 'd':
295     if(argc == 6 &&
296        sscanf(argv[2], "%1f", &x) == 1 &&
297        sscanf(argv[3], "%1f", &y) == 1 &&
298        sscanf(argv[4], "%u", &n) == 1 &&
299        sscanf(argv[5], "%u", &D) == 1)
300        printf("pow(%.1f, %.1f) ~= %.1f\n",
301              d(D), x, d(D), y, D, improved_pow(x, y, n));
302 else
303     printf("Usage: %s d <x=Value for pow(x,y)> "
304           "<y=Value for pow(x,y)> <n> "
305           "<D=Digits to display>\n", argv[0]);
306 break;
307
308 case 'e':
309     if(argc == 6 &&
310        sscanf(argv[4], "%u", &D) == 1 &&
311        sscanf(argv[5], "%u", &p) == 1)
312     {
313         mpfr_set_default_prec(p);
314         INIT_CONSTANTS
315
316         if(mpfr_init_set_str(X, argv[2], 10, MPFR_RNDN) == 0 &&
317            mpz_init_set_str(N, argv[3], 10) == 0)
318         {
319             mpfr_init(R);
320
321             sprintf(sf, "exp(%%.1f) = ~\n\t%.1f\n",
322                   d(D), D);
323
324             mpfr_cont_frac_exp_v3(R, X, N);
325             mpfr_printf(sf, X, R);
326         }
327     else
328         printf("Usage: %s e <X=Value for exp(X)> <N> "
329               "<D=Digits to display> "
330               "<p=bits of precision>\n", argv[0]);
331 }
332 else
333     printf("Usage: %s e <X=Value for exp(X)> <N> "
334           "<D=Digits to display> "
335           "<p=bits of precision>\n", argv[0]);
336 break;
337
338 case 'f':
339     if(argc == 7 &&
340        sscanf(argv[5], "%u", &D) == 1 &&
341        sscanf(argv[6], "%u", &p) == 1)
342     {

```

```

343     mpfr_set_default_prec(p);
344     INIT_CONSTANTS
345
346     if (mpfr_init_set_str(X, argv[2], 10, MPFR_RNDN) == 0 &&
347         mpfr_init_set_str(Y, argv[3], 10, MPFR_RNDN) == 0 &&
348         mpz_init_set_str(N, argv[4], 10) == 0)
349     {
350         mpfr_init(R);
351
352         sprintf(sf, "pow(%%.%uRNf, %%%uRNf) =~"
353             "\n\t%%.%uRNf\n",
354             d(D), d(D), D);
355
356         mpfr_improved_pow(R, X, Y, N);
357         mpfr_printf(sf, X, Y, R);
358     }
359     else
360         printf("Usage: %s f <X=Value for pow(X, Y)> "
361             "<Y=Value for pow(X,Y)> <N> "
362             "<D=Digits to display> "
363             "<p=bits of precision>\n", argv[0]);
364 }
365 else
366     printf("Usage: %s f <X=Value for pow(X, Y)> "
367         "<Y=Value for pow(X,Y)> <N> "
368         "<D=Digits to display> "
369         "<p=bits of precision>\n", argv[0]);
370 break;
371
372 default:
373     printf("Usage: %s <a/b/c/d/e/f> <arguments>\n", argv[0]);
374 }
375 }
376 else
377     printf("Usage: %s <a/b/c/d/e/f> <arguments>\n", argv[0]);
378 }
379 #endif

```

Header Files for Exponential and Logarithmic Functions:

File : int_exp.h

```

1 #ifndef INT_EXP_HEADER
2 #define INT_EXP_HEADER
3
4 double naive_int_exp(const double, const int);
5 double squaring_int_exp(const double, const int);
6 void mpfr_naive_int_exp(mpfr_t, mpfr_t, mpz_t);
7 void mpfr_squaring_int_exp(mpfr_t, mpfr_t, mpz_t);
8
9 #endif

```

File : taylor_exp_log.h

```

1 #ifndef TAYLOR_EXP_LOG_HEADER
2 #define TAYLOR_EXP_LOG_HEARDR
3
4 double naive_exp(double, unsigned int);
5 void mpfr_naive_exp(mpfr_t, mpfr_t, mpz_t);
6

```

```

7 | double taylor_exp(double, unsigned int);
8 | double taylor_nat_log(double, unsigned int);
9 |
10 | double taylor_log(double, double, unsigned int);
11 | double taylor_pow(double, double, unsigned int);
12 |
13 | void mpfr_taylor_exp(mpfr_t, mpfr_t, mpz_t);
14 | void mpfr_taylor_nat_log(mpfr_t, mpfr_t, mpz_t);
15 |
16 | void mpfr_taylor_log(mpfr_t, mpfr_t, mpfr_t, mpz_t);
17 | void mpfr_taylor_pow(mpfr_t, mpfr_t, mpfr_t, mpz_t);
18 |
19 | #endif

```

File : hyperbolic_log.h

```

1 | #ifndef HYPERBOLIC_LOG_HEADER
2 | #define HYPERBOLIC_LOG_HEADER
3 |
4 | double hyperbolic_nat_log(double, unsigned int);
5 | double hyperbolic_log(double, double, unsigned int);
6 | void mpfr_hyperbolic_nat_log(mpfr_t, mpfr_t, mpz_t);
7 | void mpfr_hyperbolic_log(mpfr_t, mpfr_t, mpfr_t, mpz_t);
8 |
9 | #endif

```

File : cont_frac_exp.h

```

1 | #ifndef CONT_FRAC_EXP_HEADER
2 | #define CONT_FRAC_EXP_HEADER
3 |
4 | double cont_frac_exp_v1(double, unsigned int);
5 | double cont_frac_exp_v2(double, unsigned int);
6 | double cont_frac_exp_v3(double, unsigned int);
7 | void mpfr_cont_frac_v3(mpfr_t, mpfr_t, mpz_t);
8 |
9 | double improved_pow(double, double, unsigned int);
10 | void mpfr_improved_pow(mpfr_t, mpfr_t, mpfr_t, mpz_t);
11 |
12 | #endif

```